

Desenvolvimento de Aplicações com Arquitetura Baseada em Microservices

Prof. Vinicius Cardoso Garcia
vcg@cin.ufpe.br :: @vinicius3w :: assertlab.com

[IF1007] - Tópicos Avançados em SI 4
<https://github.com/vinicius3w/if1007-Microservices>

Licença do material

Este Trabalho foi licenciado com uma Licença
Creative Commons - Atribuição-NãoComercial-Compartilhalgual
3.0 Não Adaptada

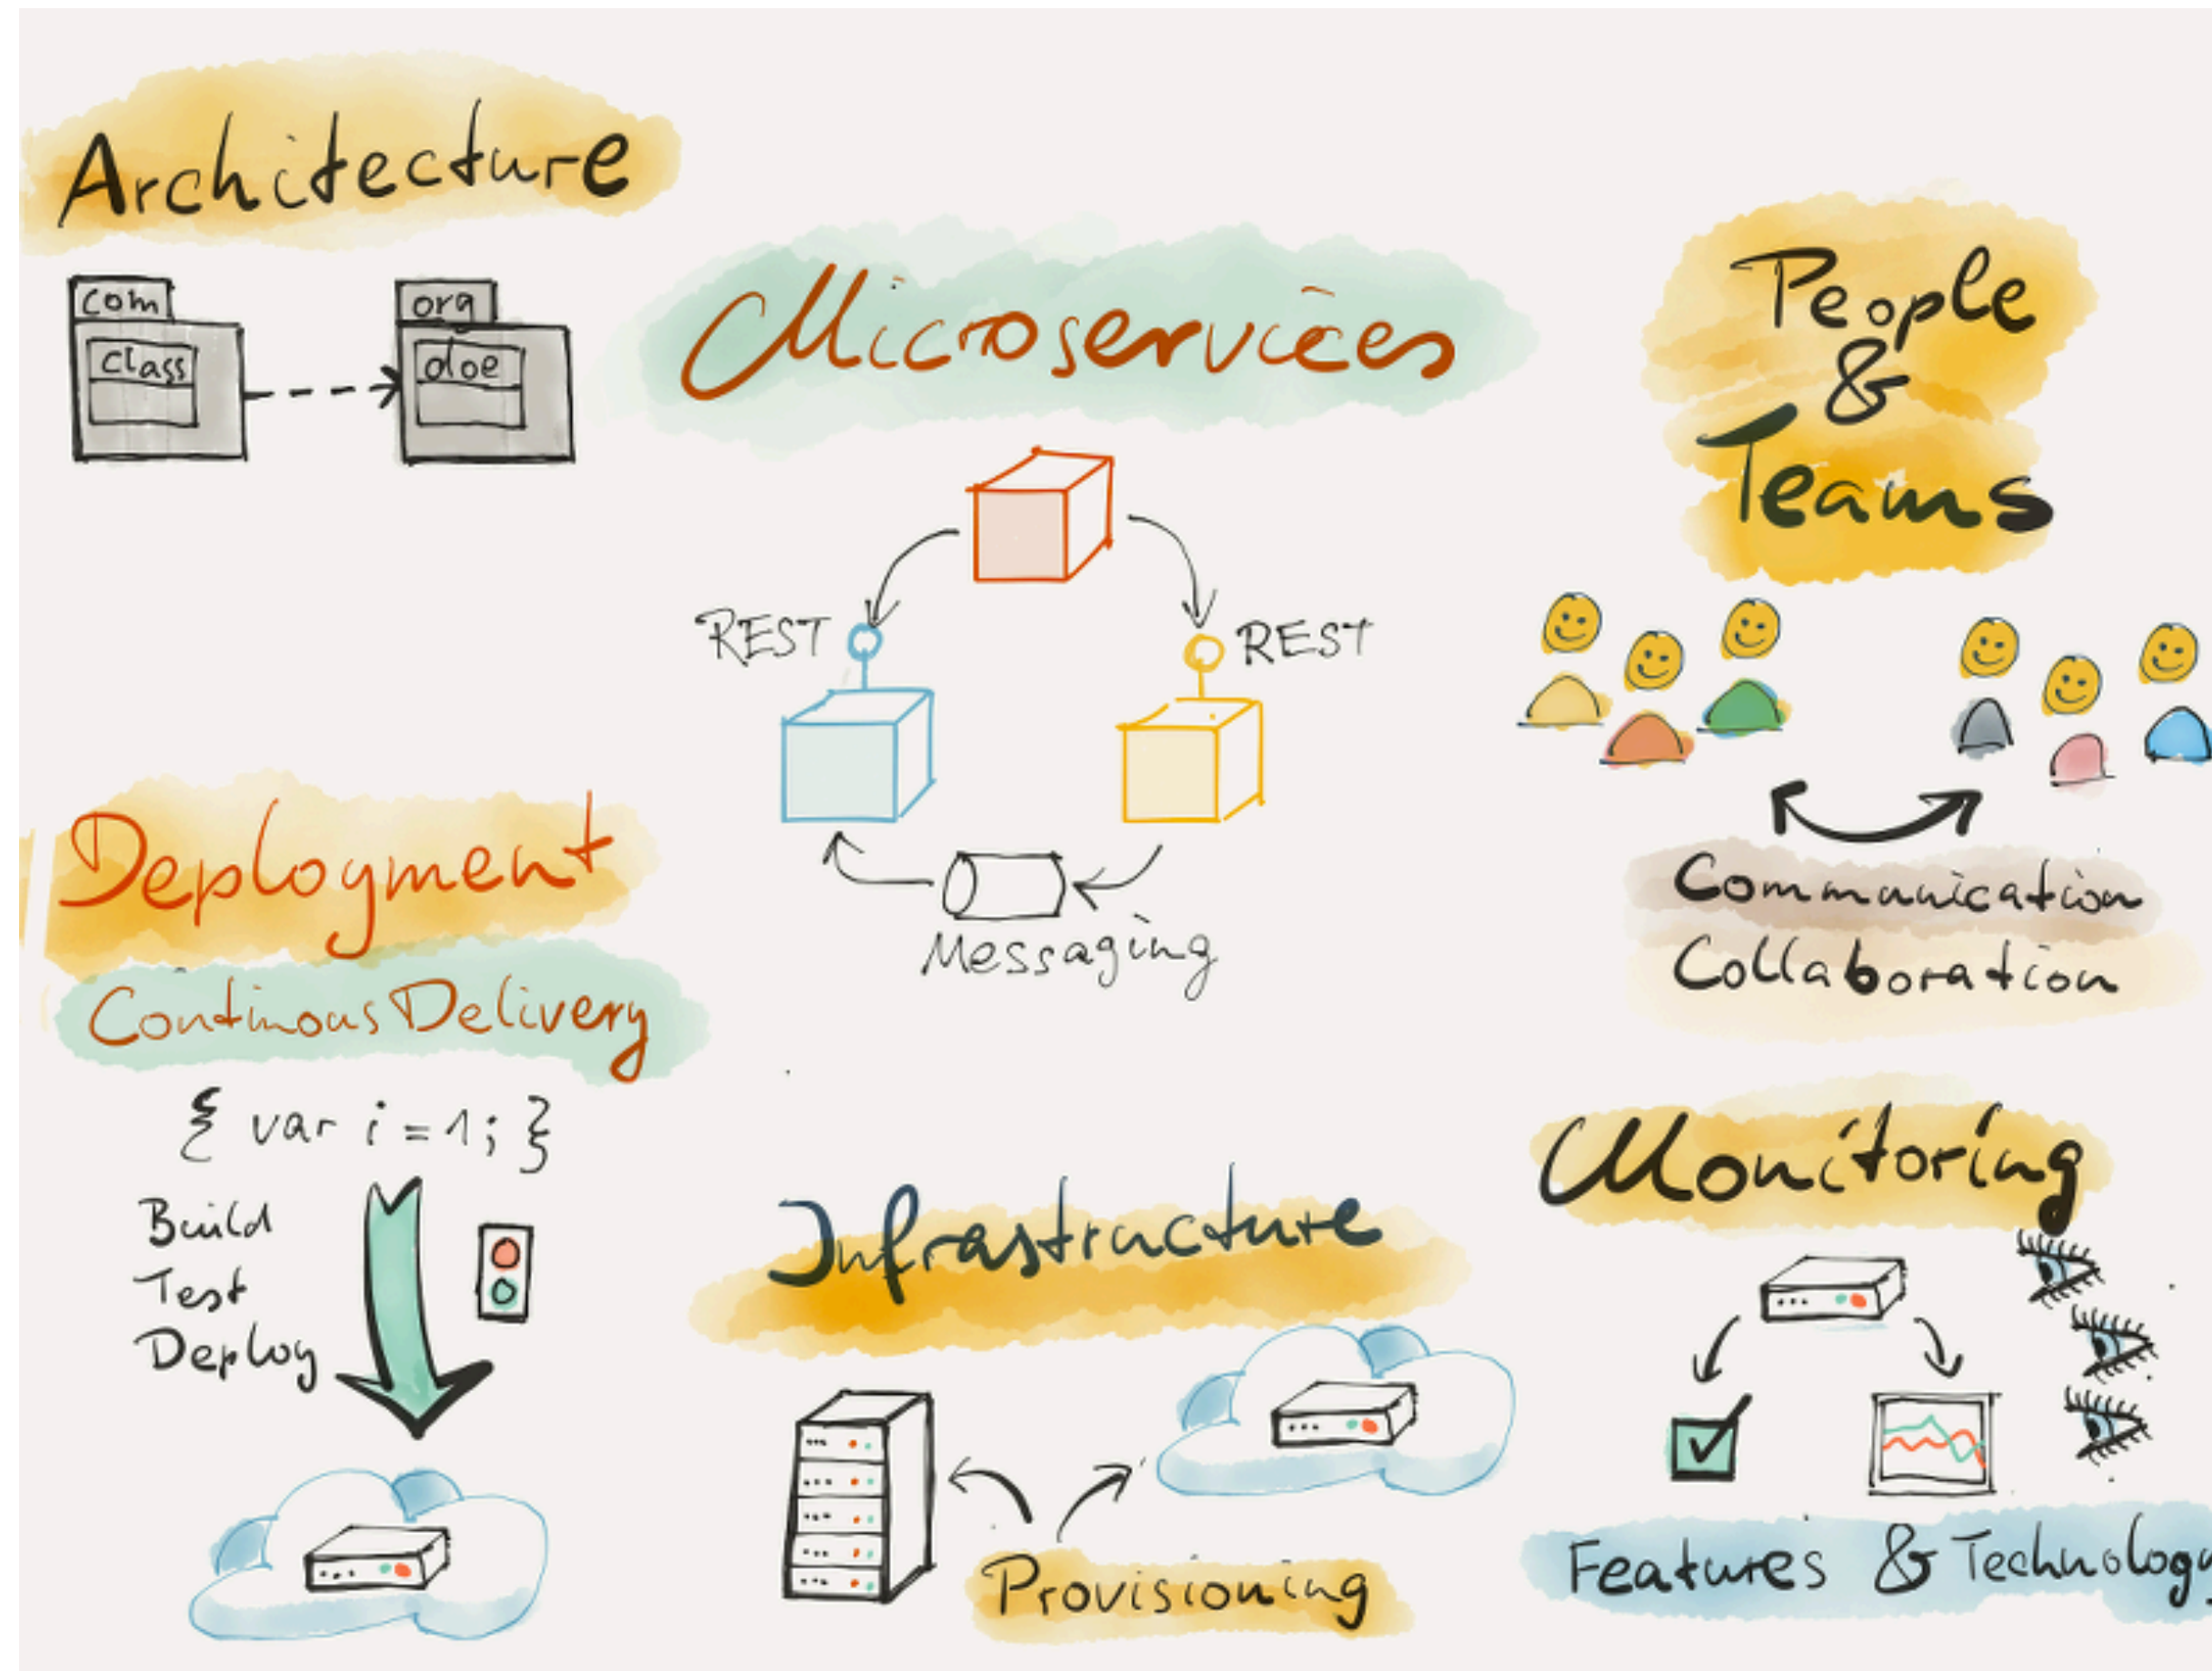


Mais informações visite

<http://creativecommons.org/licenses/by-nc-sa/3.0/deed.pt>

Resources

- There is no textbook required. However, the following are some books that may be recommended:
 - [Building Microservices: Designing Fine-Grained Systems](#)
 - [Spring Microservices](#)
 - [Spring Boot: Acelere o desenvolvimento de microserviços](#)
 - [Microservices for Java Developers A Hands-on Introduction to Frameworks and Containers](#)
 - [Migrating to Cloud-Native Application Architectures](#)
 - [Continuous Integration](#)
 - [Getting started guides from spring.io](#)



Applying Microservices Concepts

Context

- Microservices are **good**, but can also be an **evil** if they are not properly conceived.
- **Wrong** microservice interpretations could lead to **irrecoverable** failures
- What are the technical **challenges** around **practical** implementations of microservices?
 - design decisions, solutions and patterns?

Patterns and common design decisions

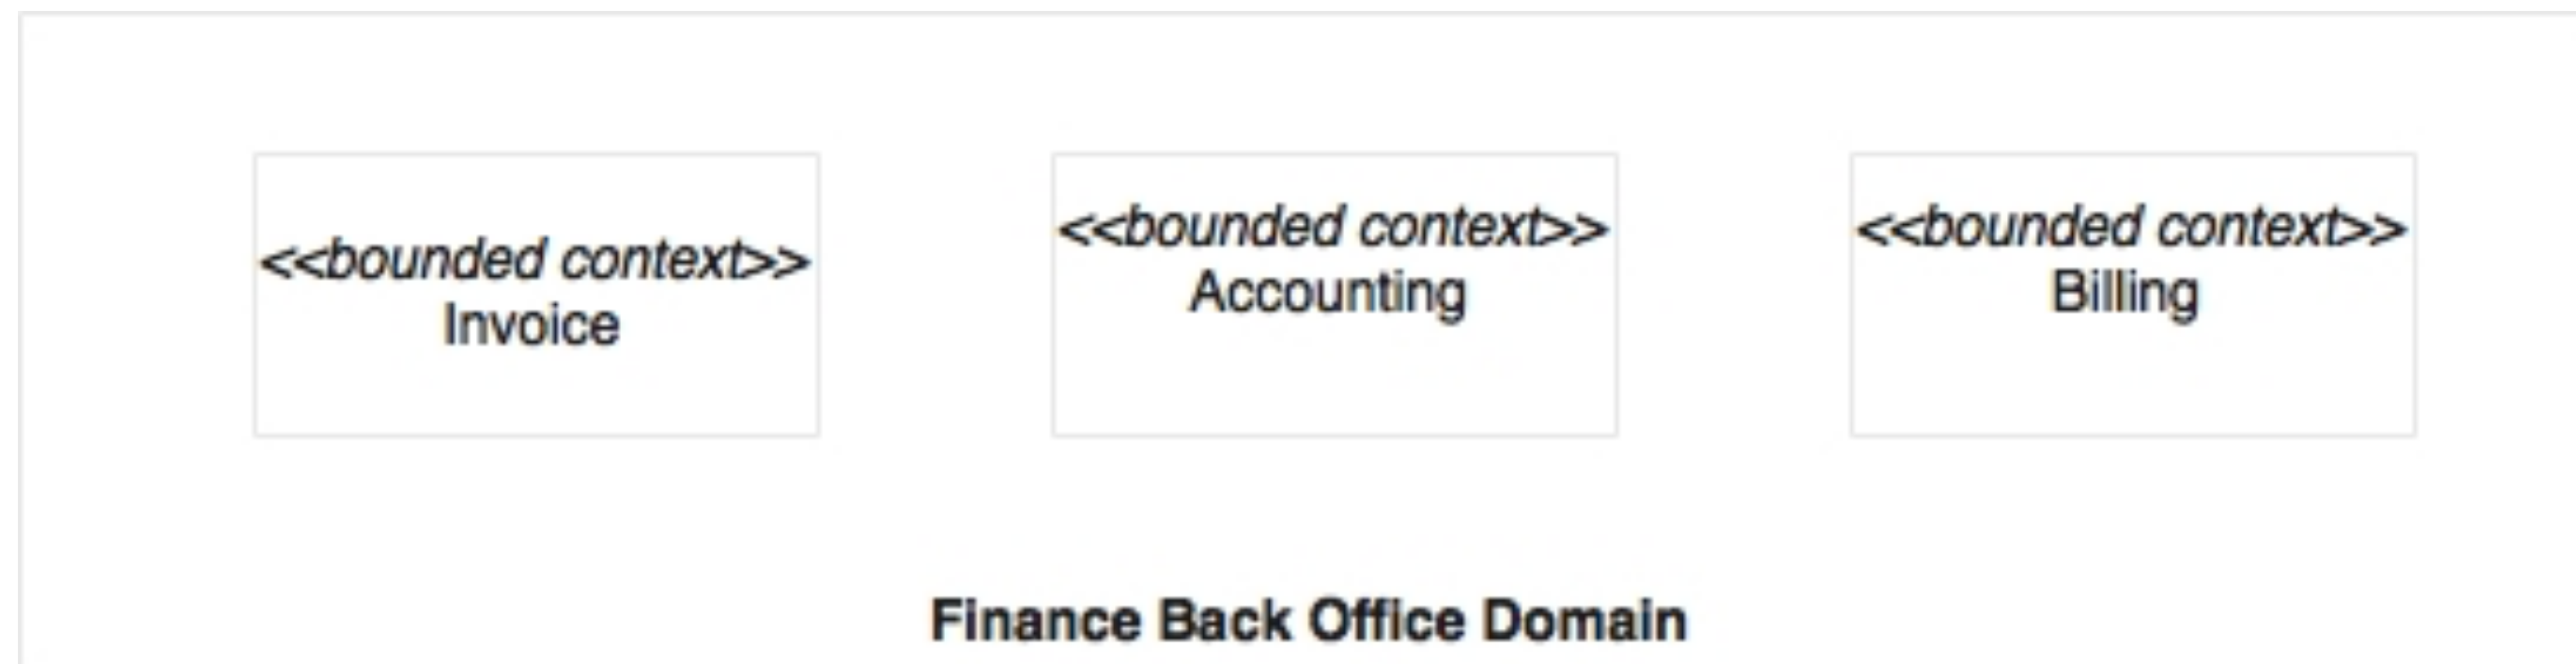
- Microservices are a vehicle for developing **scalable** cloud native systems, successful microservices need to be **carefully** designed to **avoid** catastrophes.
- Microservices **are not** the **one-size-fits-all**, universal solution for **all architecture** problems
- Generally speaking, microservices are a great choice for building a **lightweight**, **modular**, **scalable**, and **distributed** system of systems

Establishing appropriate microservice boundaries

- How **big** (mini-monolithic) or how **small** (nano service) can a microservice be, or is there anything like **right-sized services**?
- Does size really **matter**?
- A quick answer could be
 - "one REST endpoint per microservice", or
 - "less than 300 lines of code", or
 - "a component that performs a single responsibility"

Establishing appropriate microservice boundaries

- Domain-driven design (DDD) defines the concept of a **bounded context**. A bounded context is a subdomain or a subsystem of a larger domain or system that is responsible for **performing a particular function**



Establishing appropriate microservice boundaries

- A bounded context is a good way to determine the boundaries of microservices
 - Each bounded context could be mapped to a single microservice.
- In the real world, communication between bounded contexts are typically less coupled, and often, disconnected


Establishing appropriate microservice boundaries

- There is no **silver bullet** to establish microservices boundaries
 - Establishing boundaries is much easier in the scenario of **monolithic application to microservices migration**, as the service boundaries and dependencies are **known** from the existing system.
 - On the other hand, in a **green field microservices development**, the dependencies are hard to establish upfront.
- The most **pragmatic** way to design microservices boundaries is to run a lot of **scenarios**


Scenarios could help in defining the microservice boundaries

- Autonomous functions
- Size of a deployable unit
- Most appropriate function or subdomain
- Polyglot architecture
- Selective scaling
- Small, agile teams
- Single responsibility (business capability or a technical capability)
- Replicability or changeability
- Coupling and cohesion
- Think microservice as a product

Architecture




Microservices



People & Teams




Deployment Continuous Delivery



Infrastructure



Monitoring



Designing communication styles

- Synchronous style communication
 - there is no shared state or object
 - When a caller requests a service, it passes the required information and waits for a response
 - Advantages and downsides?

Designing communication styles

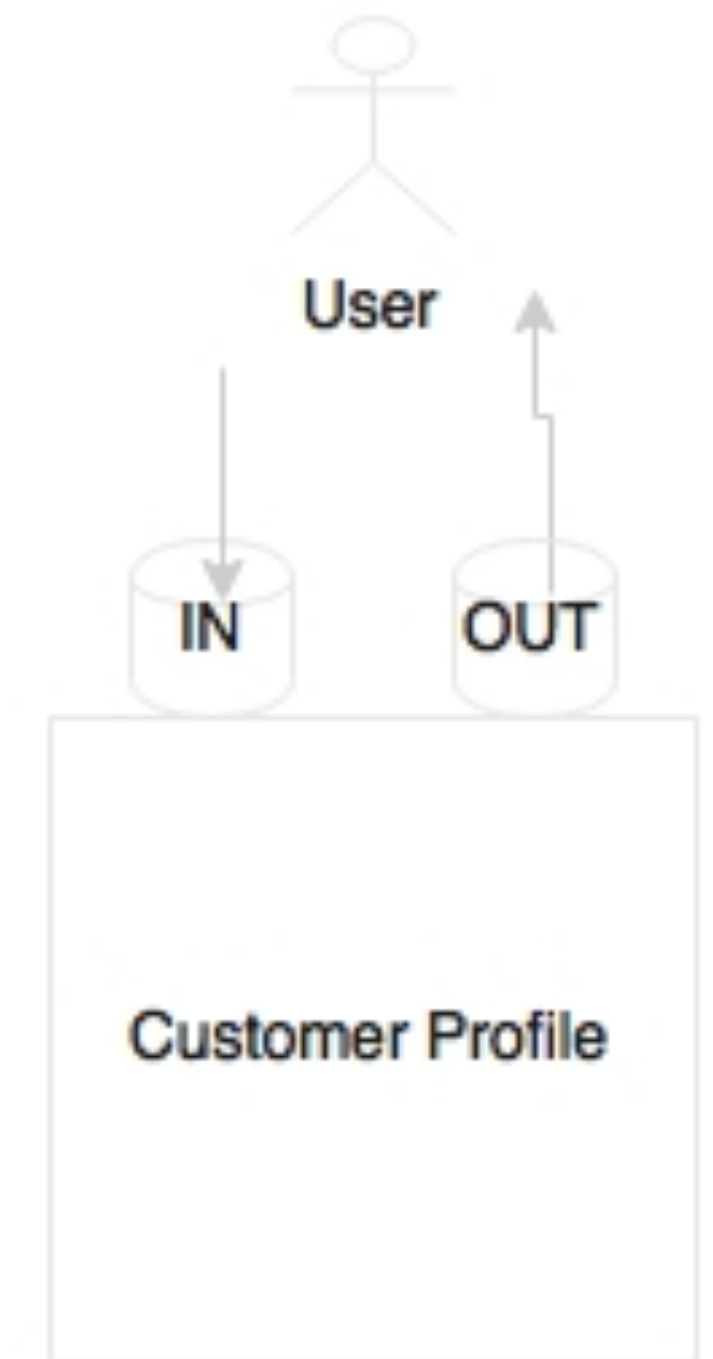
- Asynchronous style communication
 - The asynchronous style is based on **reactive event loop semantics** which decouple microservices.
 - This approach provides higher levels of **scalability**, because services are **independent**, and can internally spawn threads to handle an increase in load.
 - When overloaded, messages will be **queued** in a messaging server for **later processing**

How to decide which style to choose?

- Both approaches have their own merits and constraints
 - In principle, the asynchronous approach is great for building true, scalable microservice systems
 - However, attempting to model everything as asynchronous leads to complex system designs

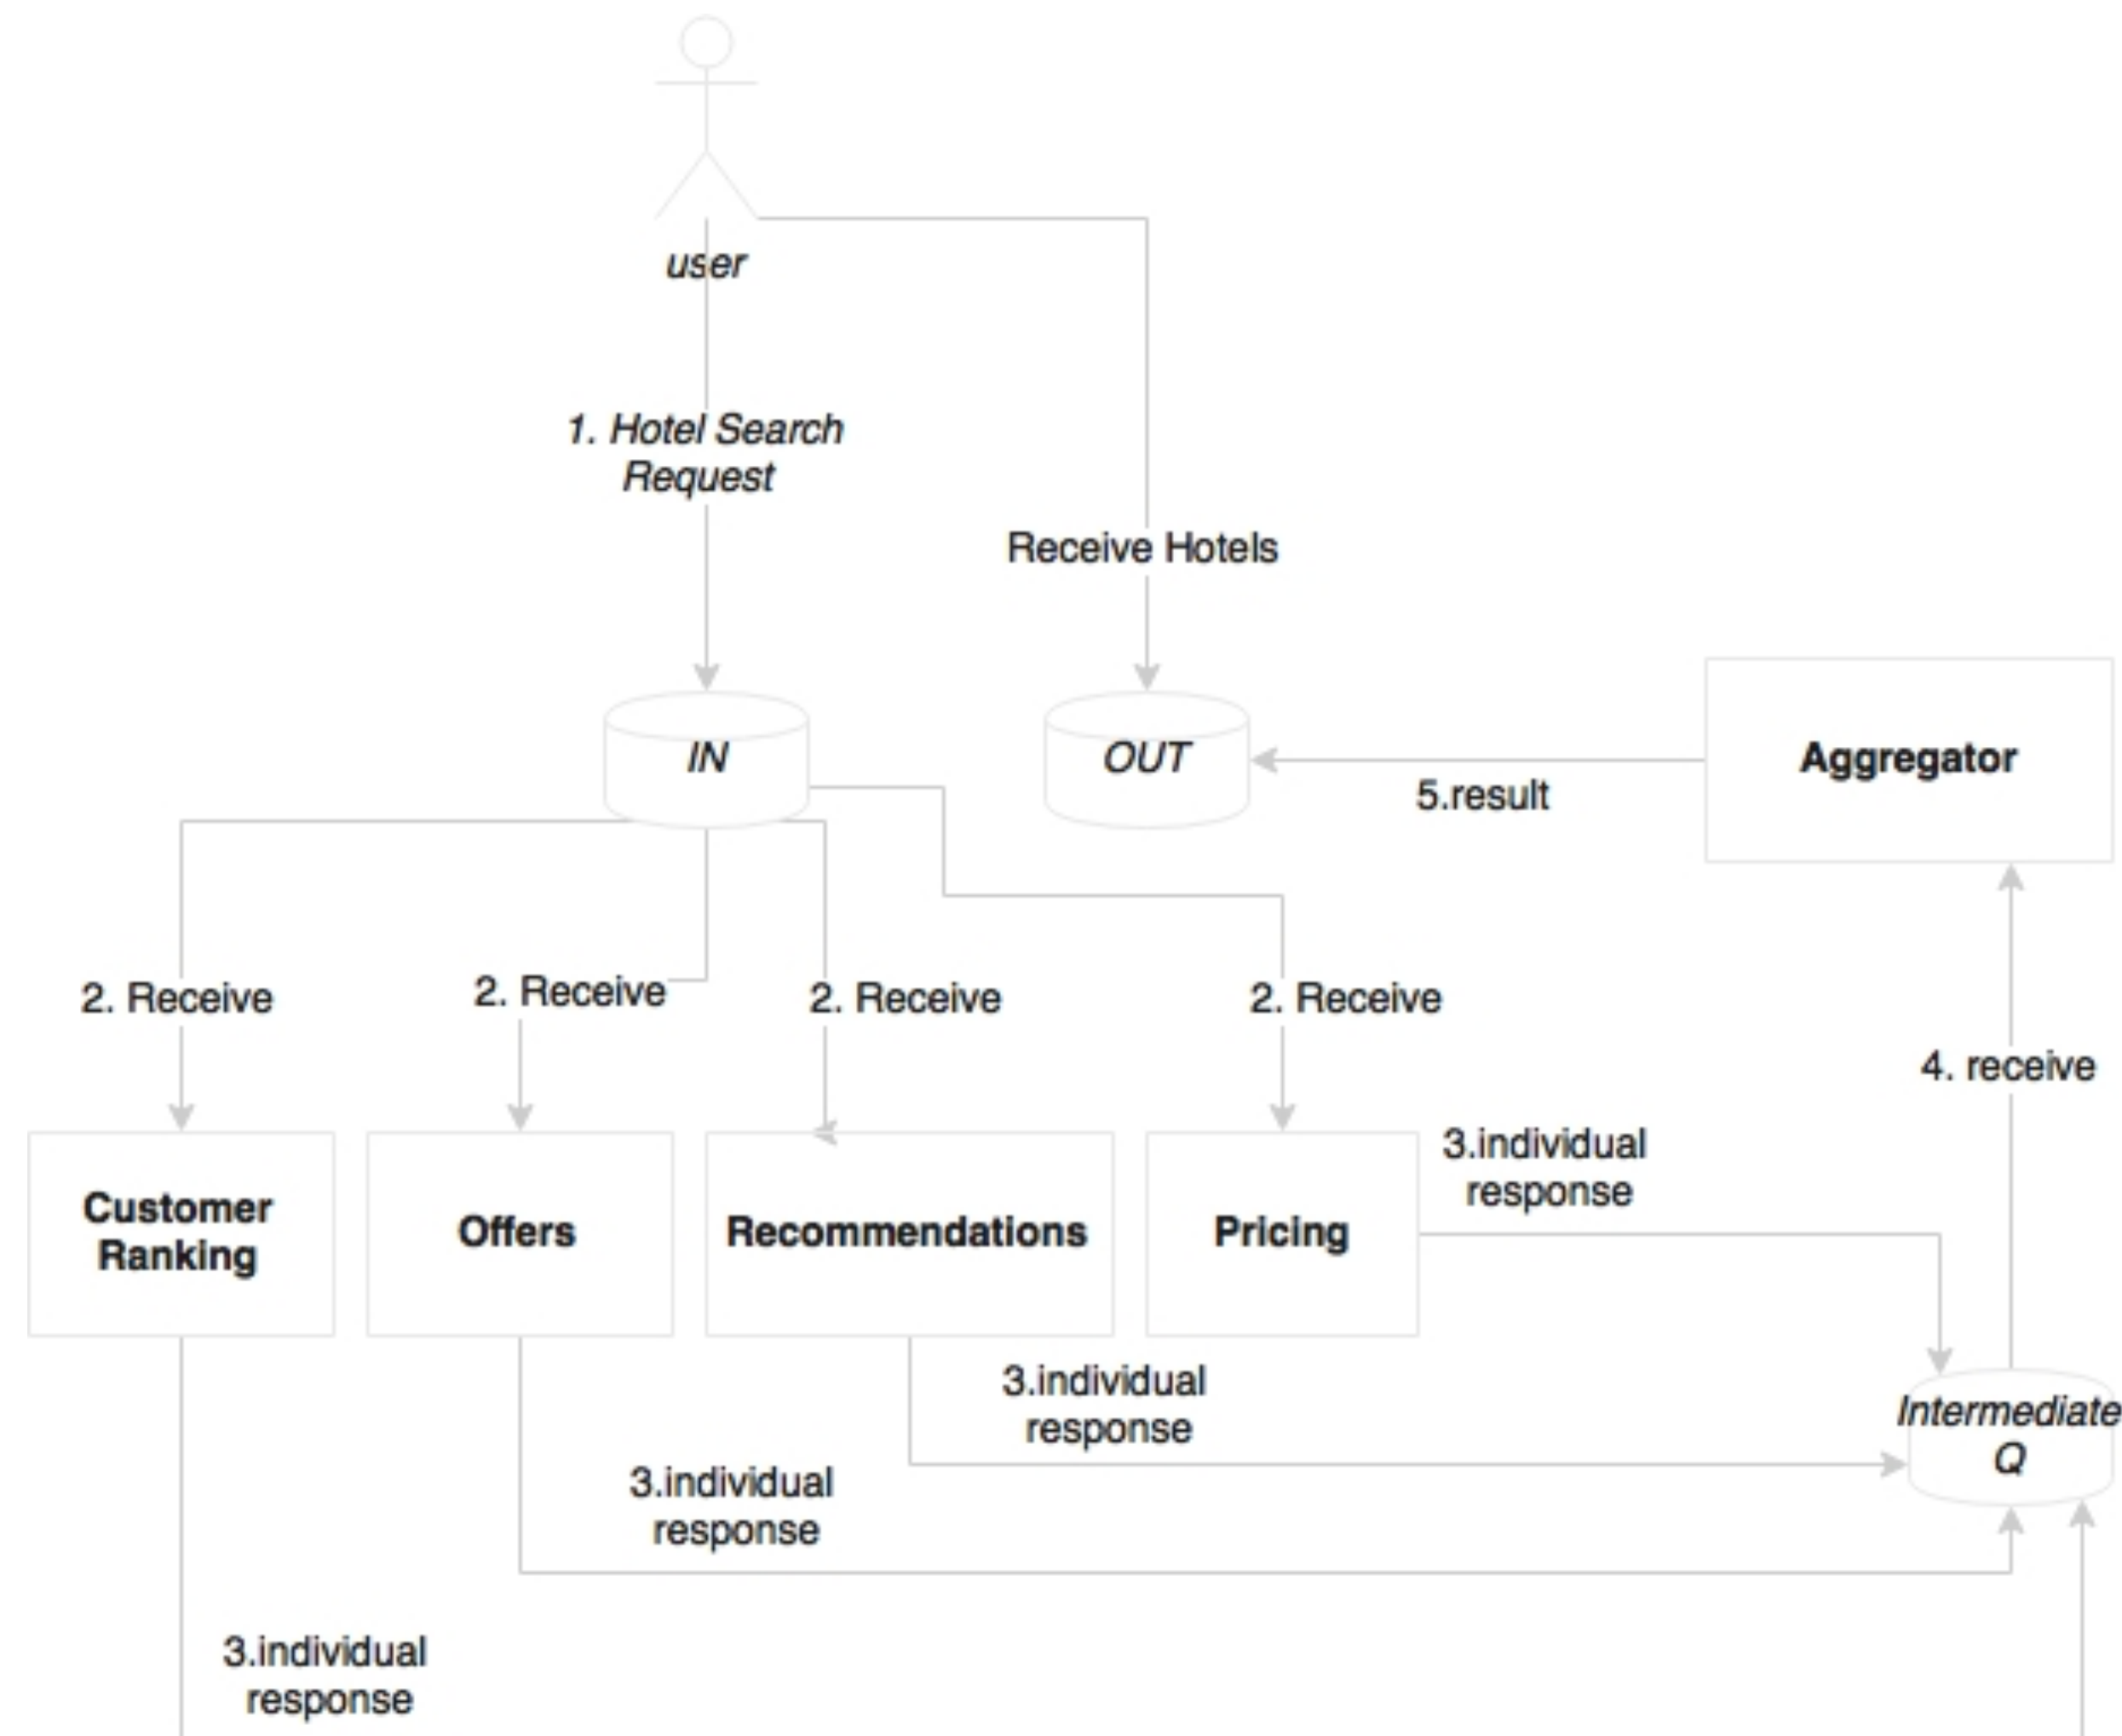


A) synchronous request - response



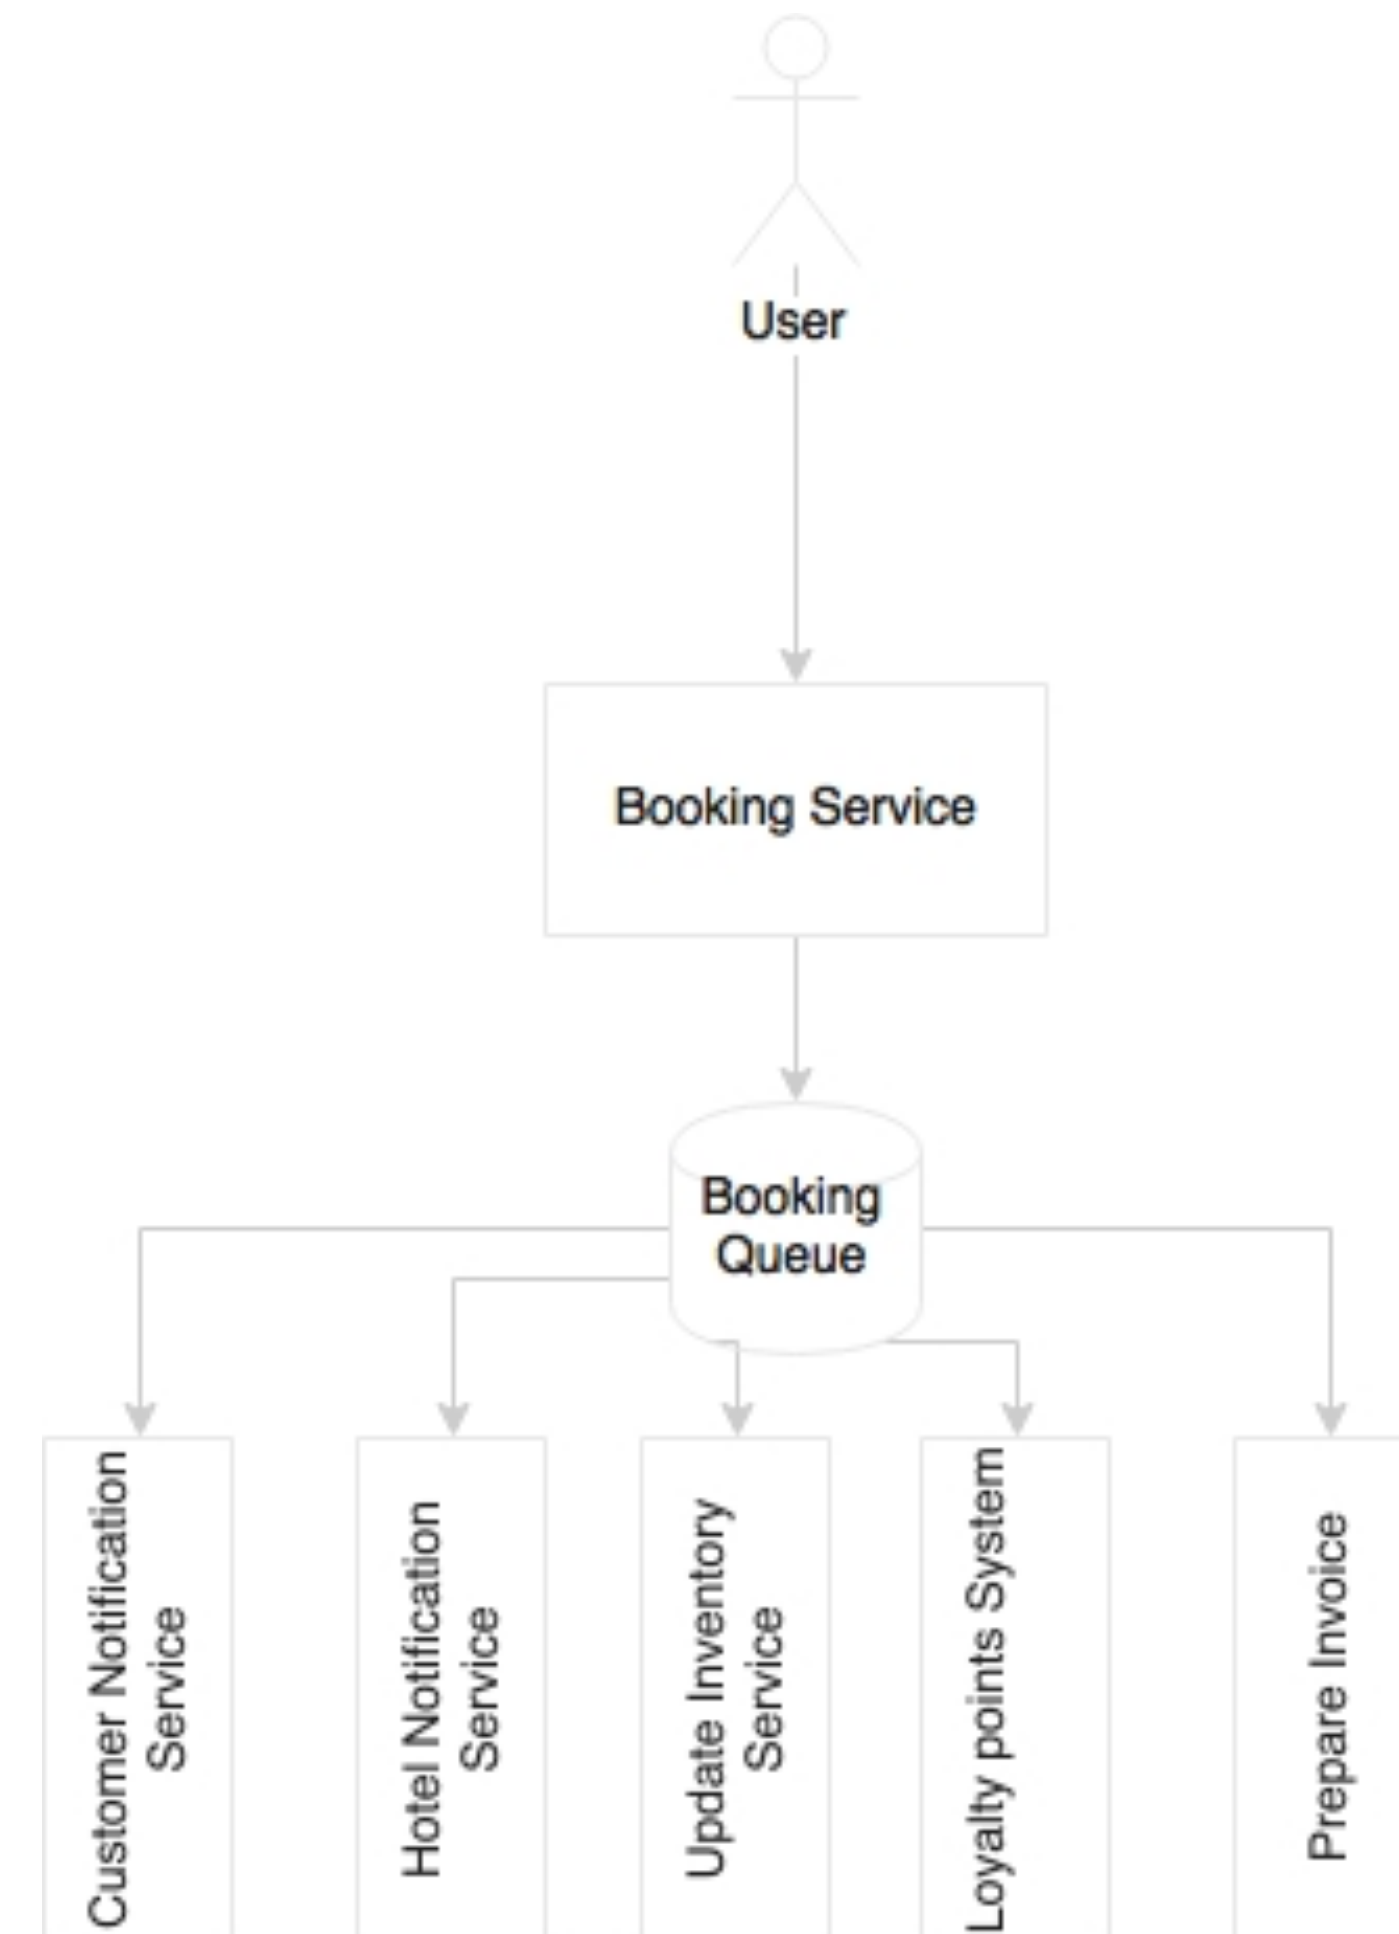
B) asynchronous request - response

How to decide which style to choose?



How to decide which style to choose?

- The service is triggered when the user clicks on the booking function.
- It is again, by nature, a synchronous style communication.
- When booking is successful, it sends a message to the customer's e-mail address, sends a message to the hotel's booking system, updates the cached inventory, updates the loyalty points system, prepares an invoice, and perhaps more

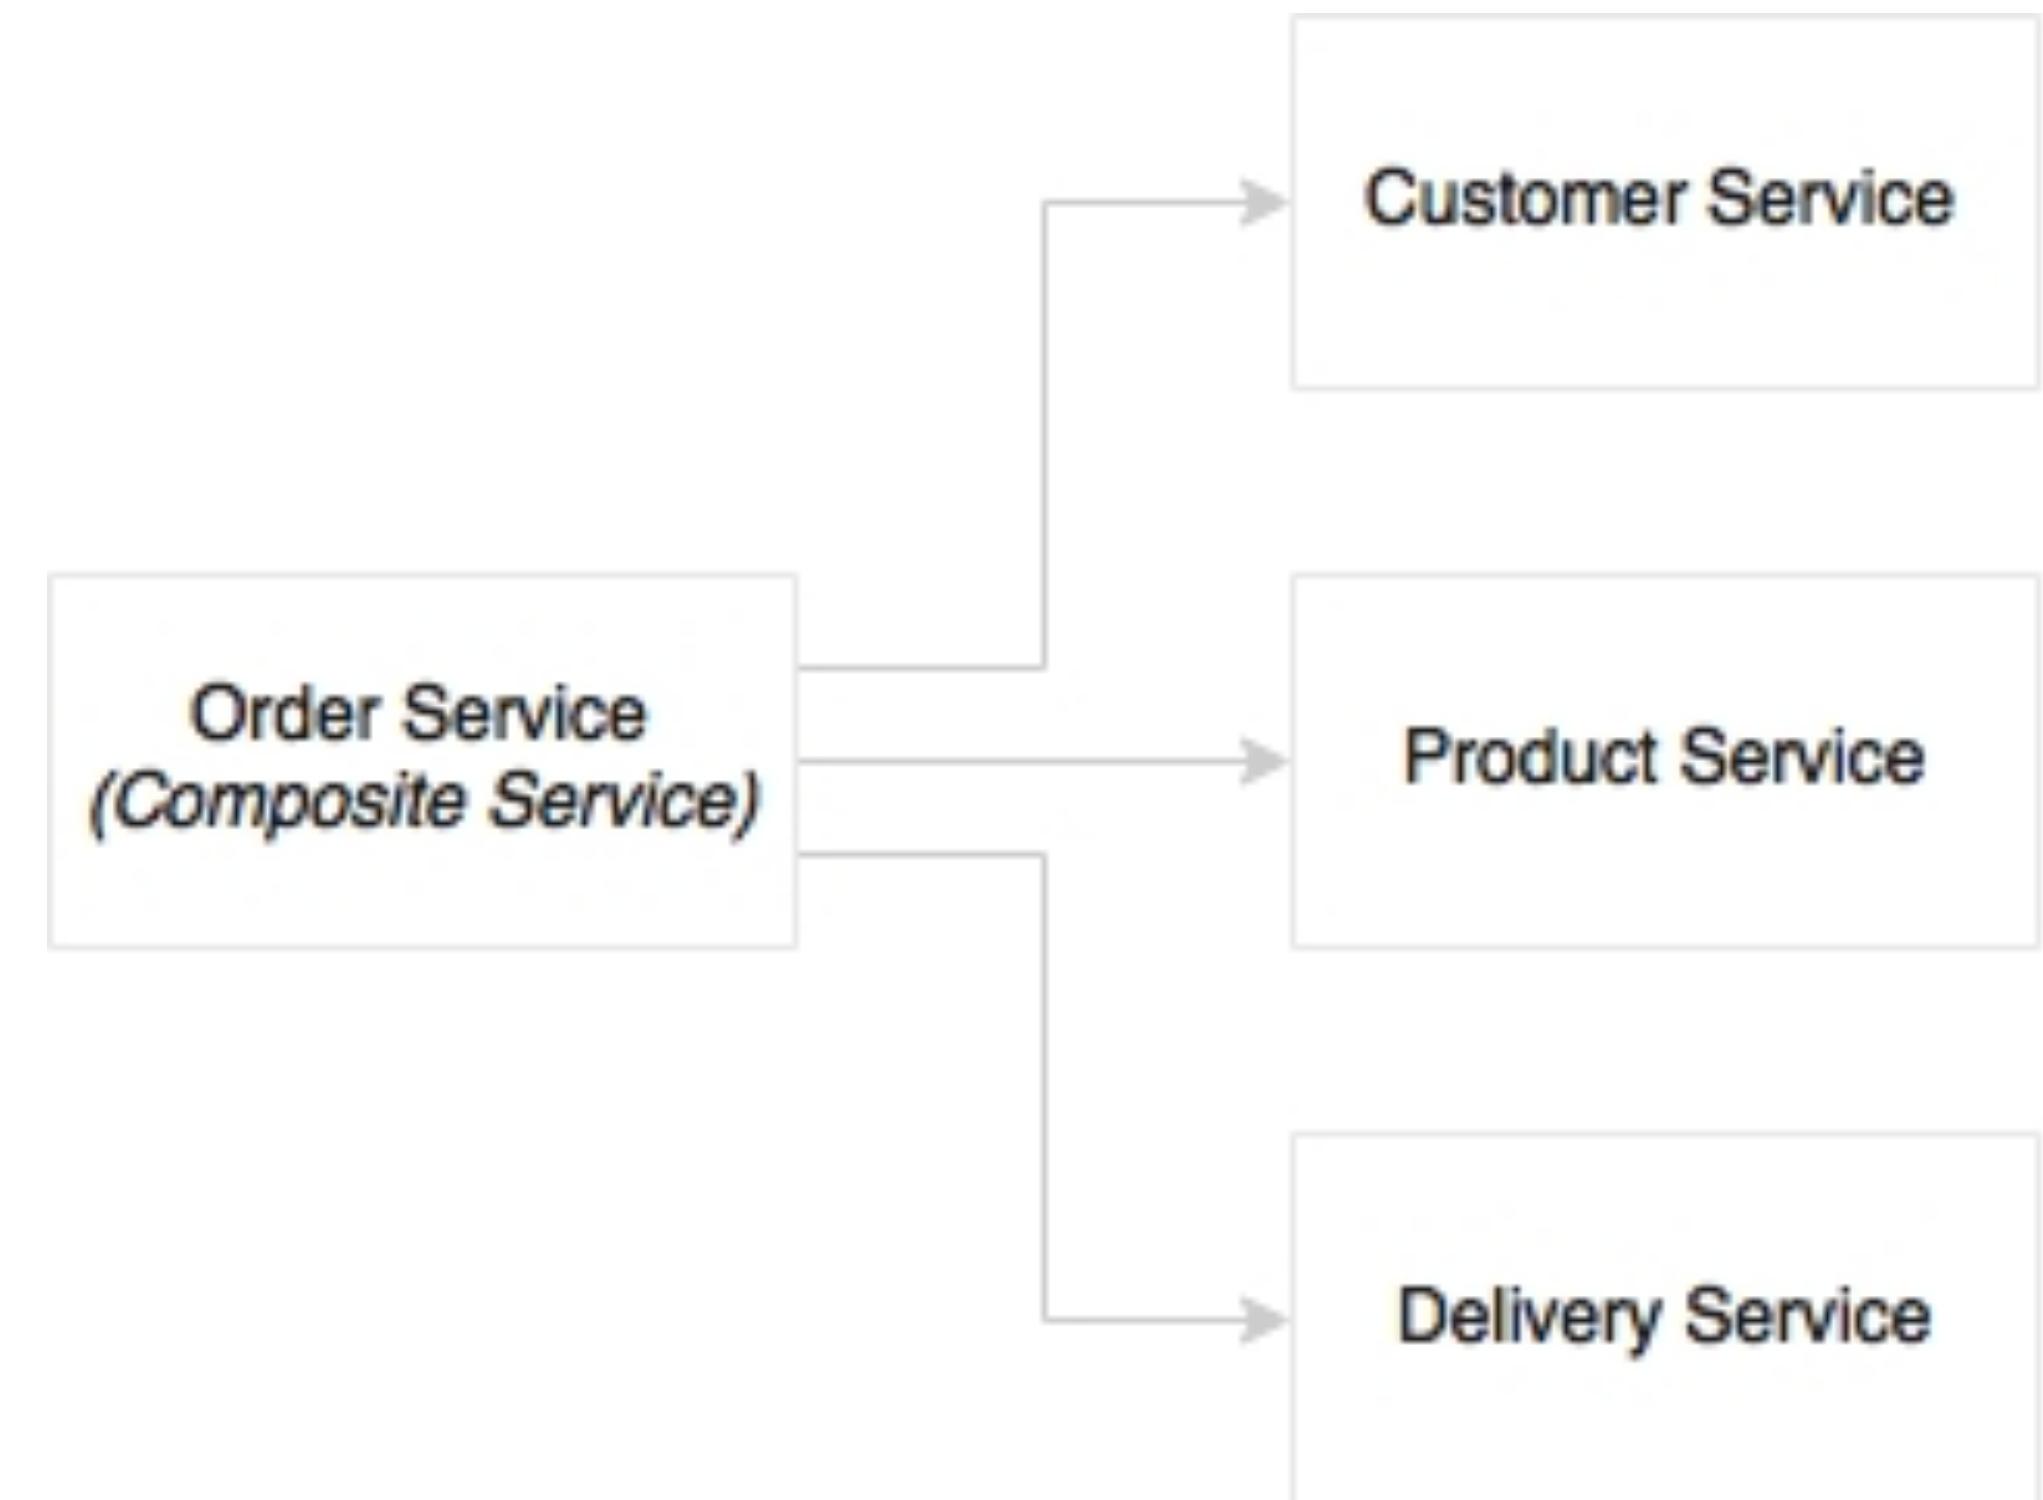


Orchestration of microservices

- Composability is one of the service design principles
 - Who is responsible for the composing services?
- SOA use ESBs (act as a proxy in some cases)
 - The first approach is orchestration

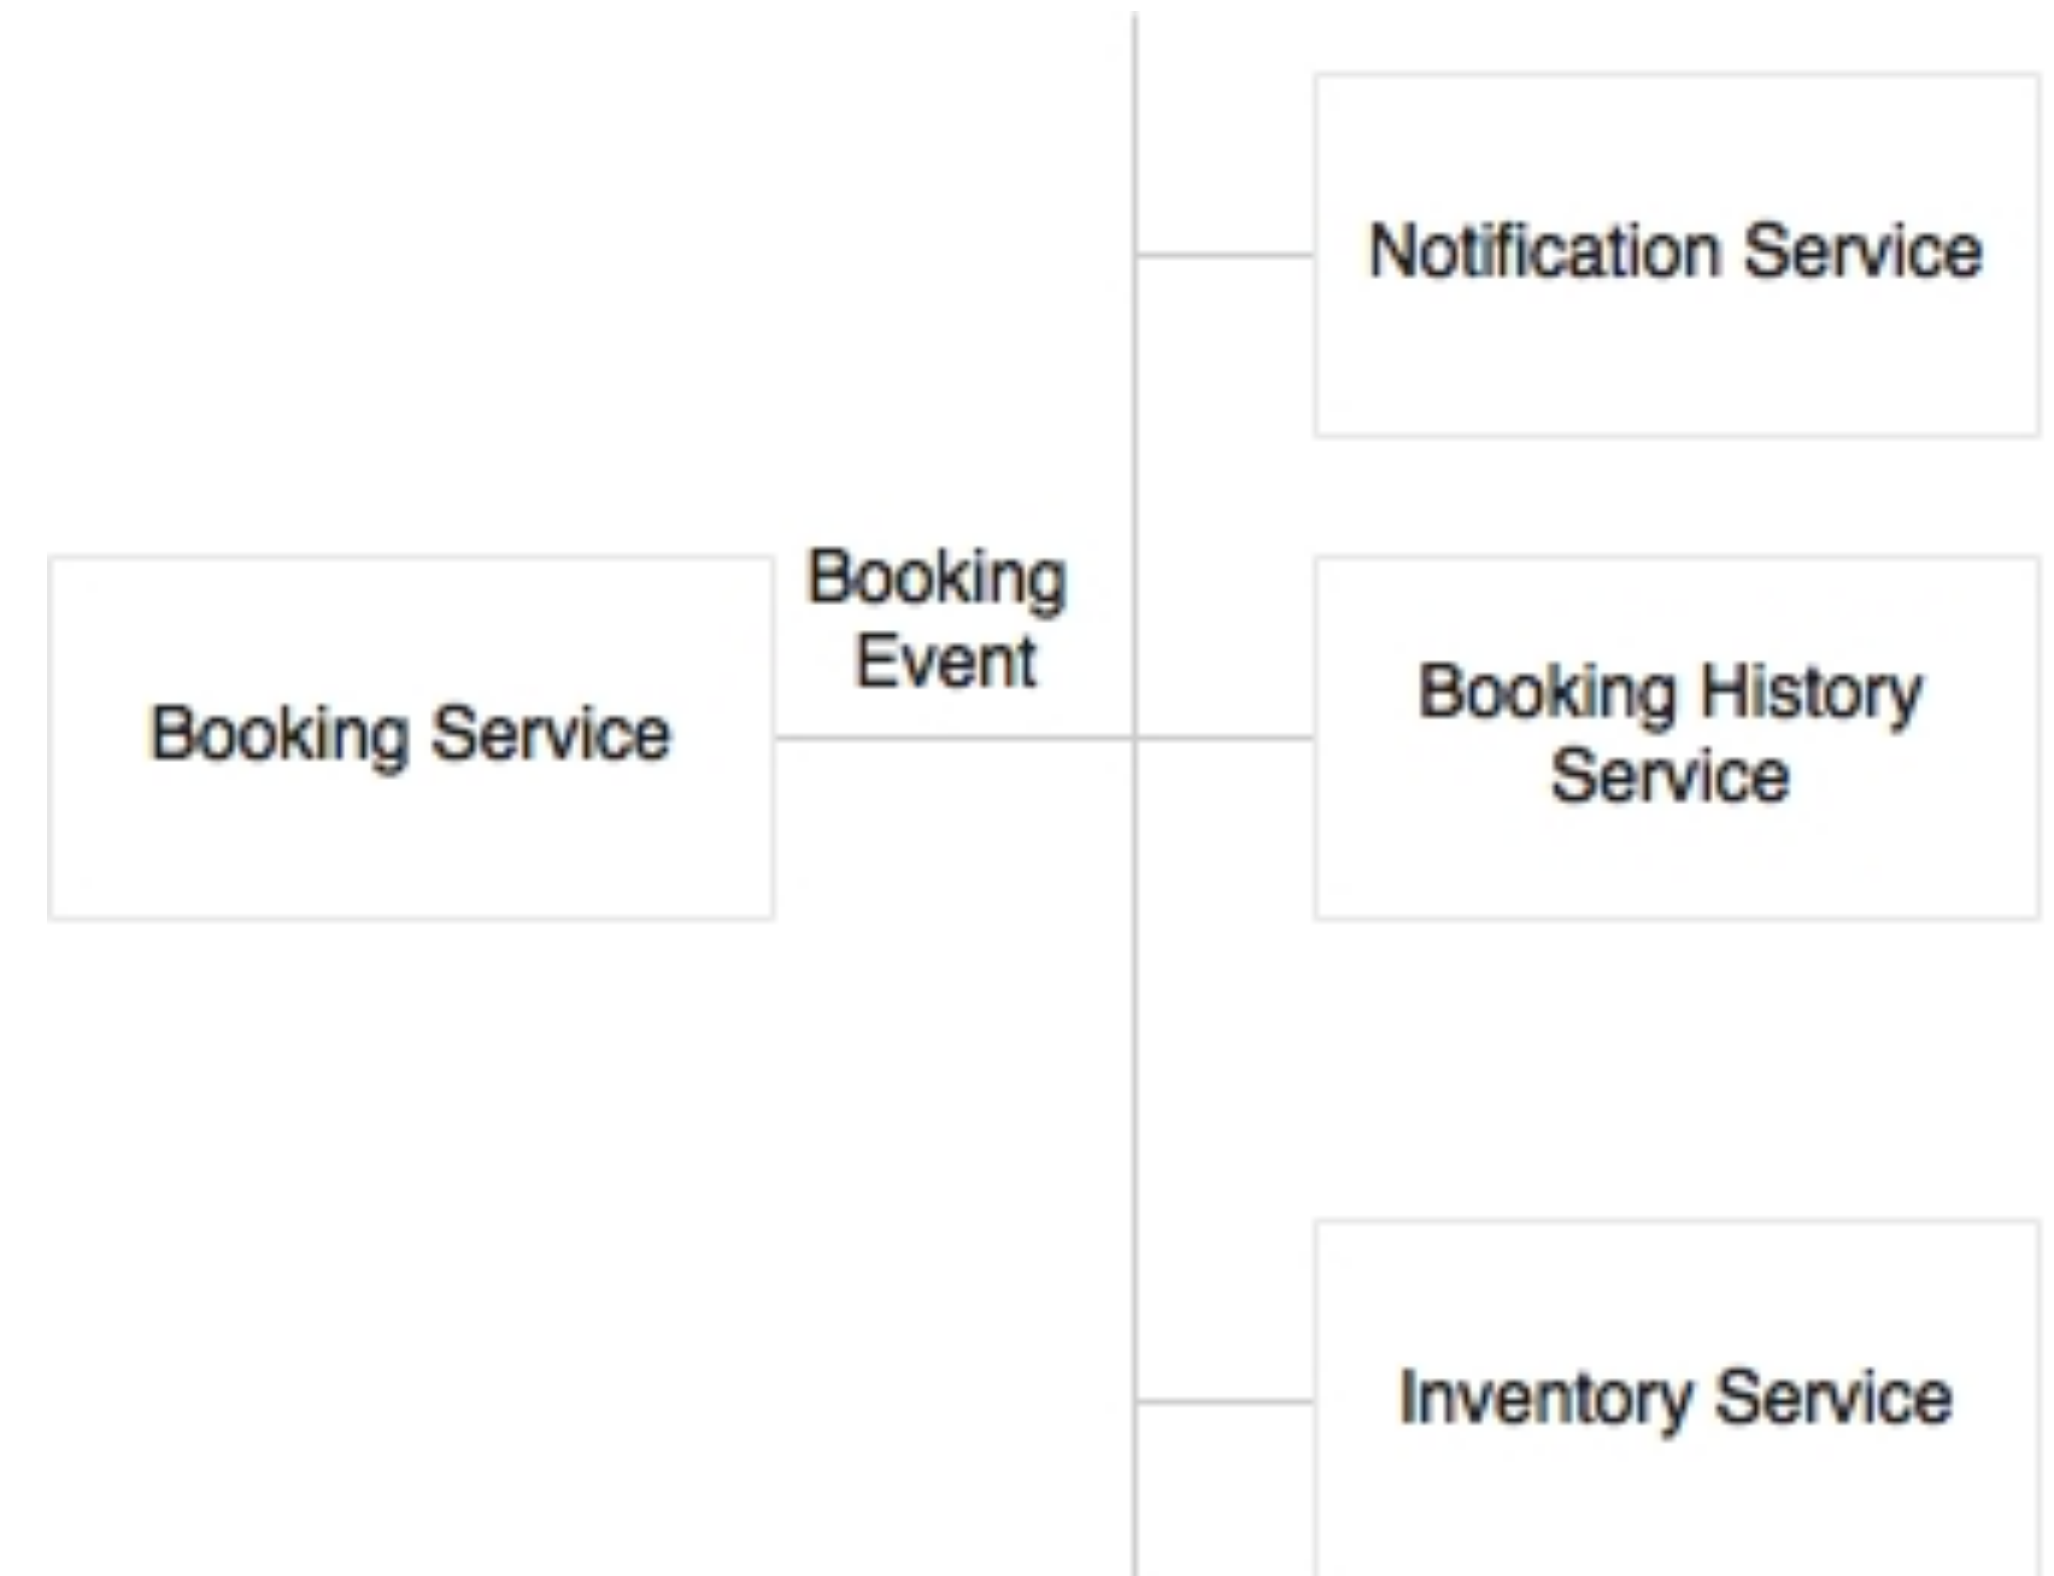
Orchestration of microservices

- Multiple services are **stitched together** to get a **complete function**. A central brain acts as the **orchestrator**
- In the SOA world, ESBs play the role of **orchestration**
- The orchestrated service will be exposed by ESBs as a composite service



Orchestration of microservices

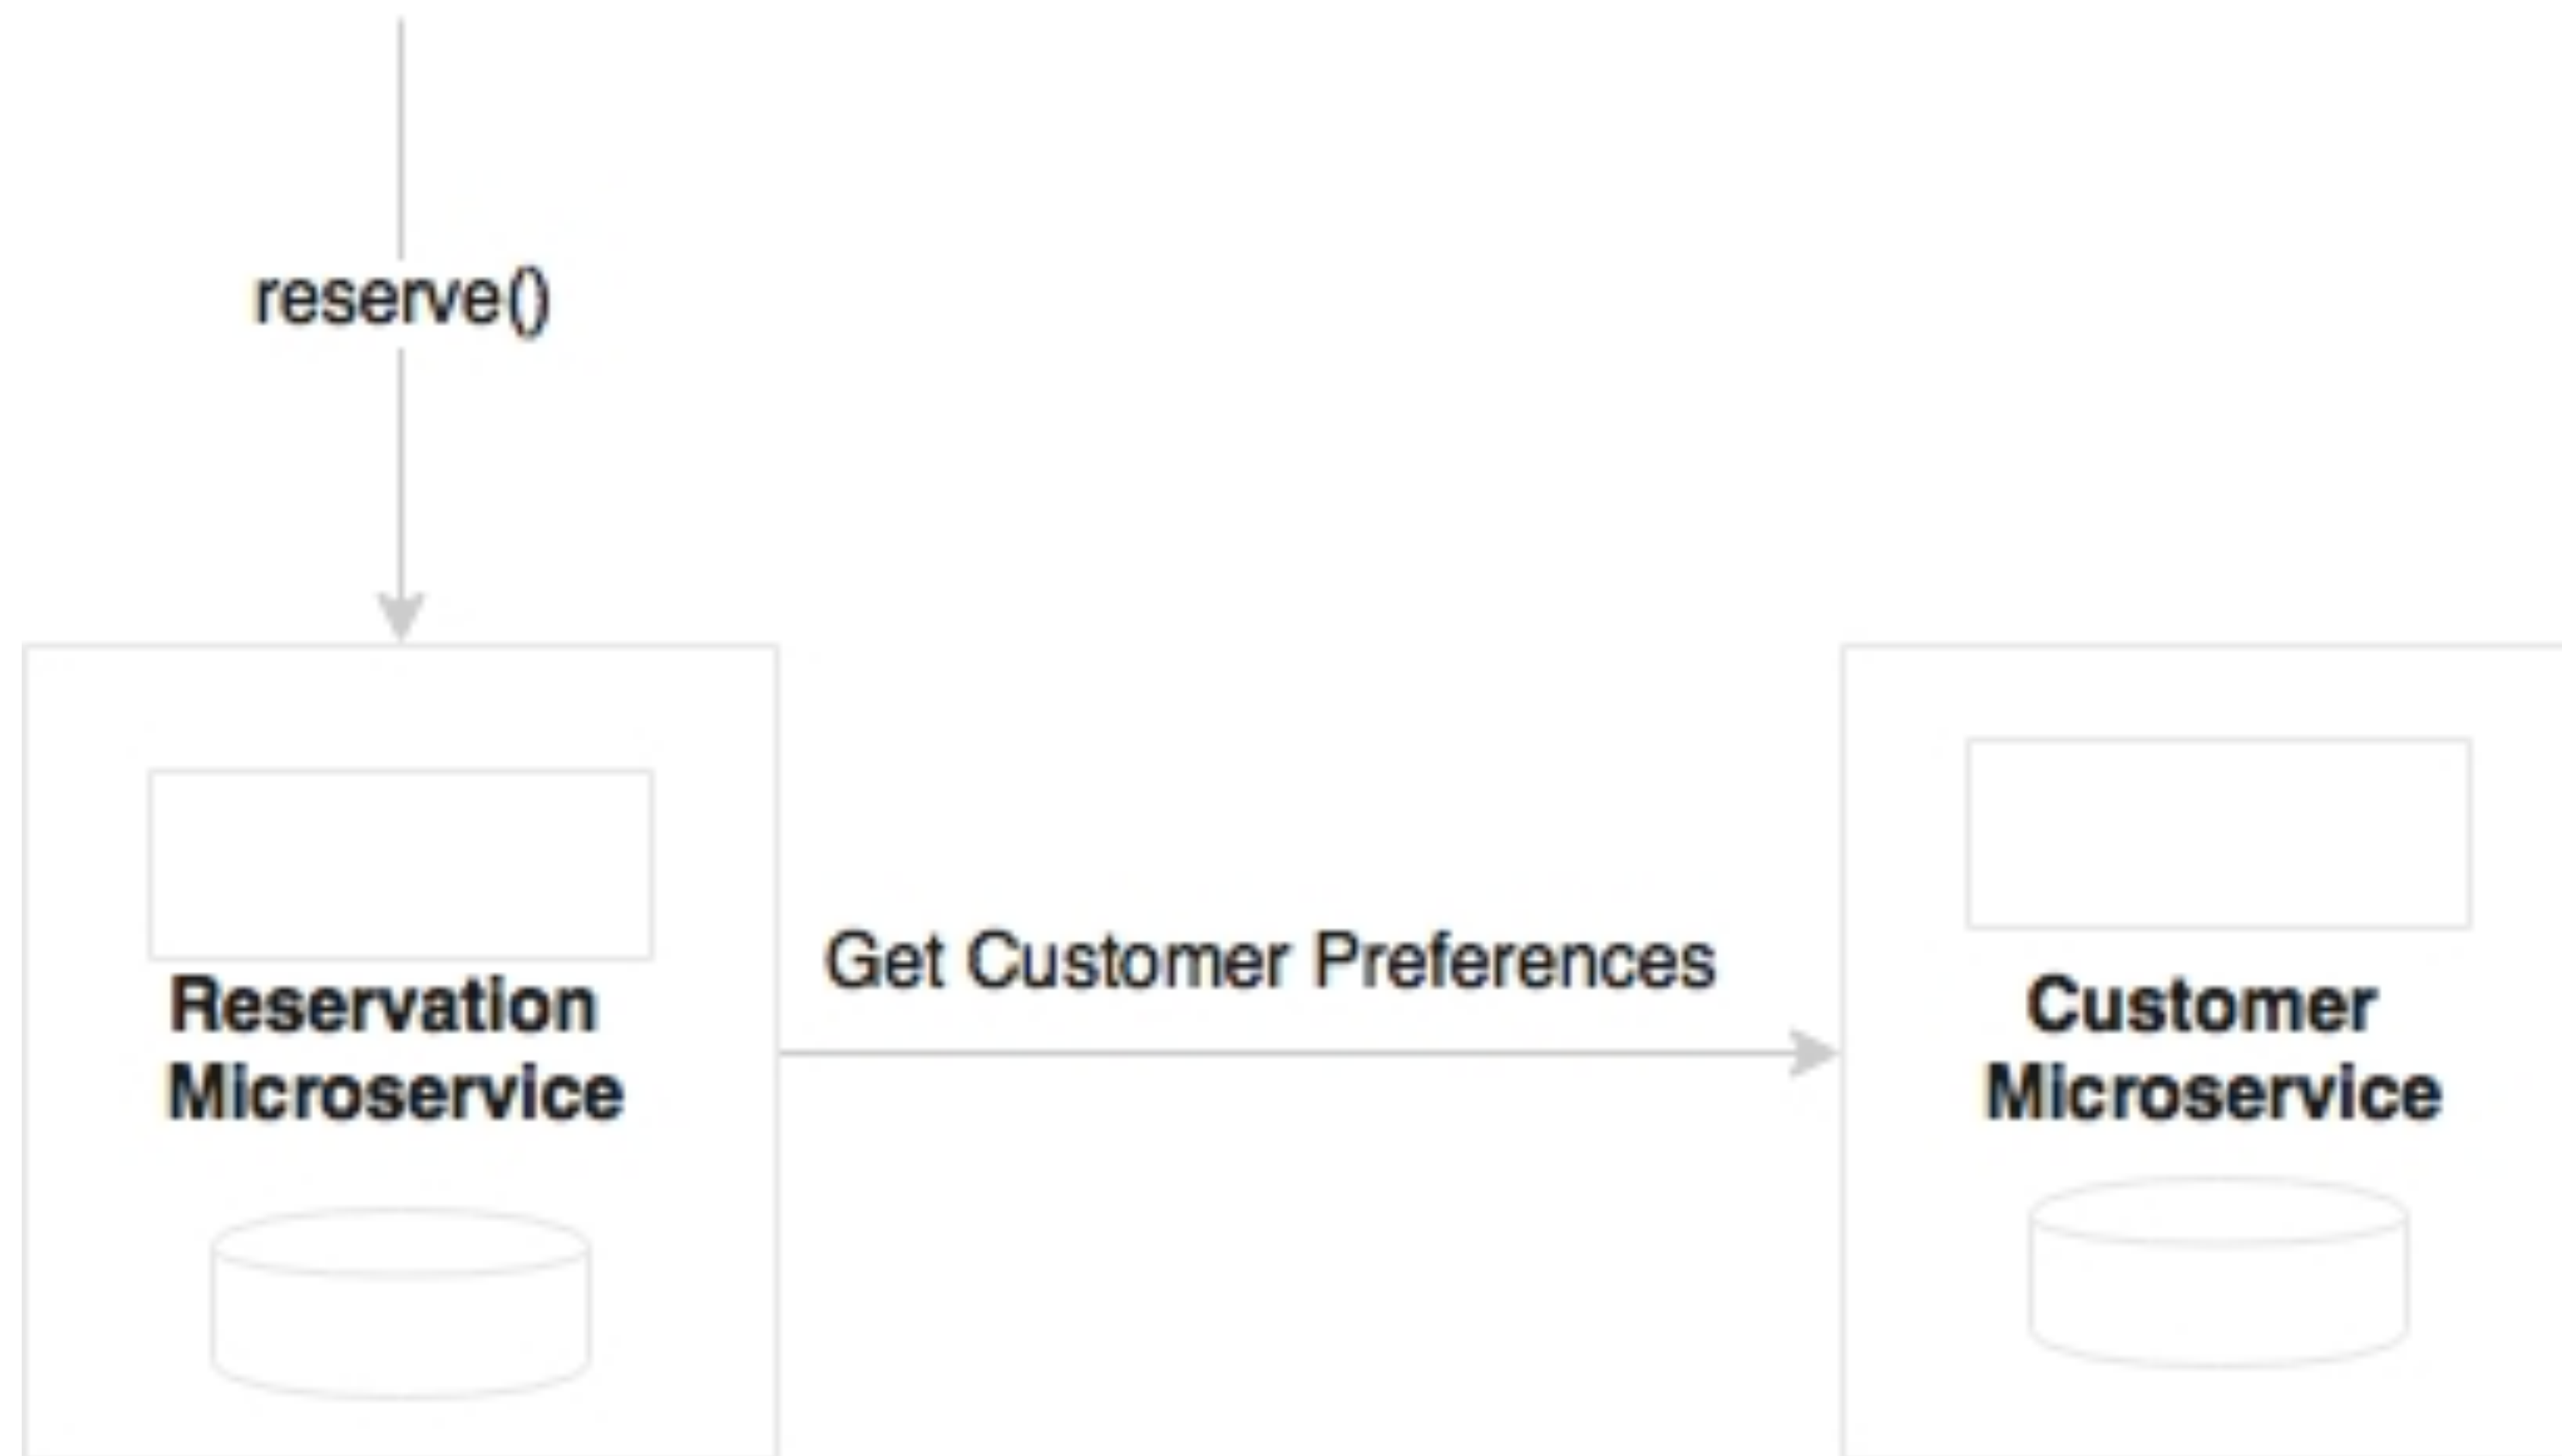
- The second approach is choreography, in where, there is no central brain
- In the SOA world, the caller pushes a message to the ESB, and the downstream flow will be automatically determined by the consuming services



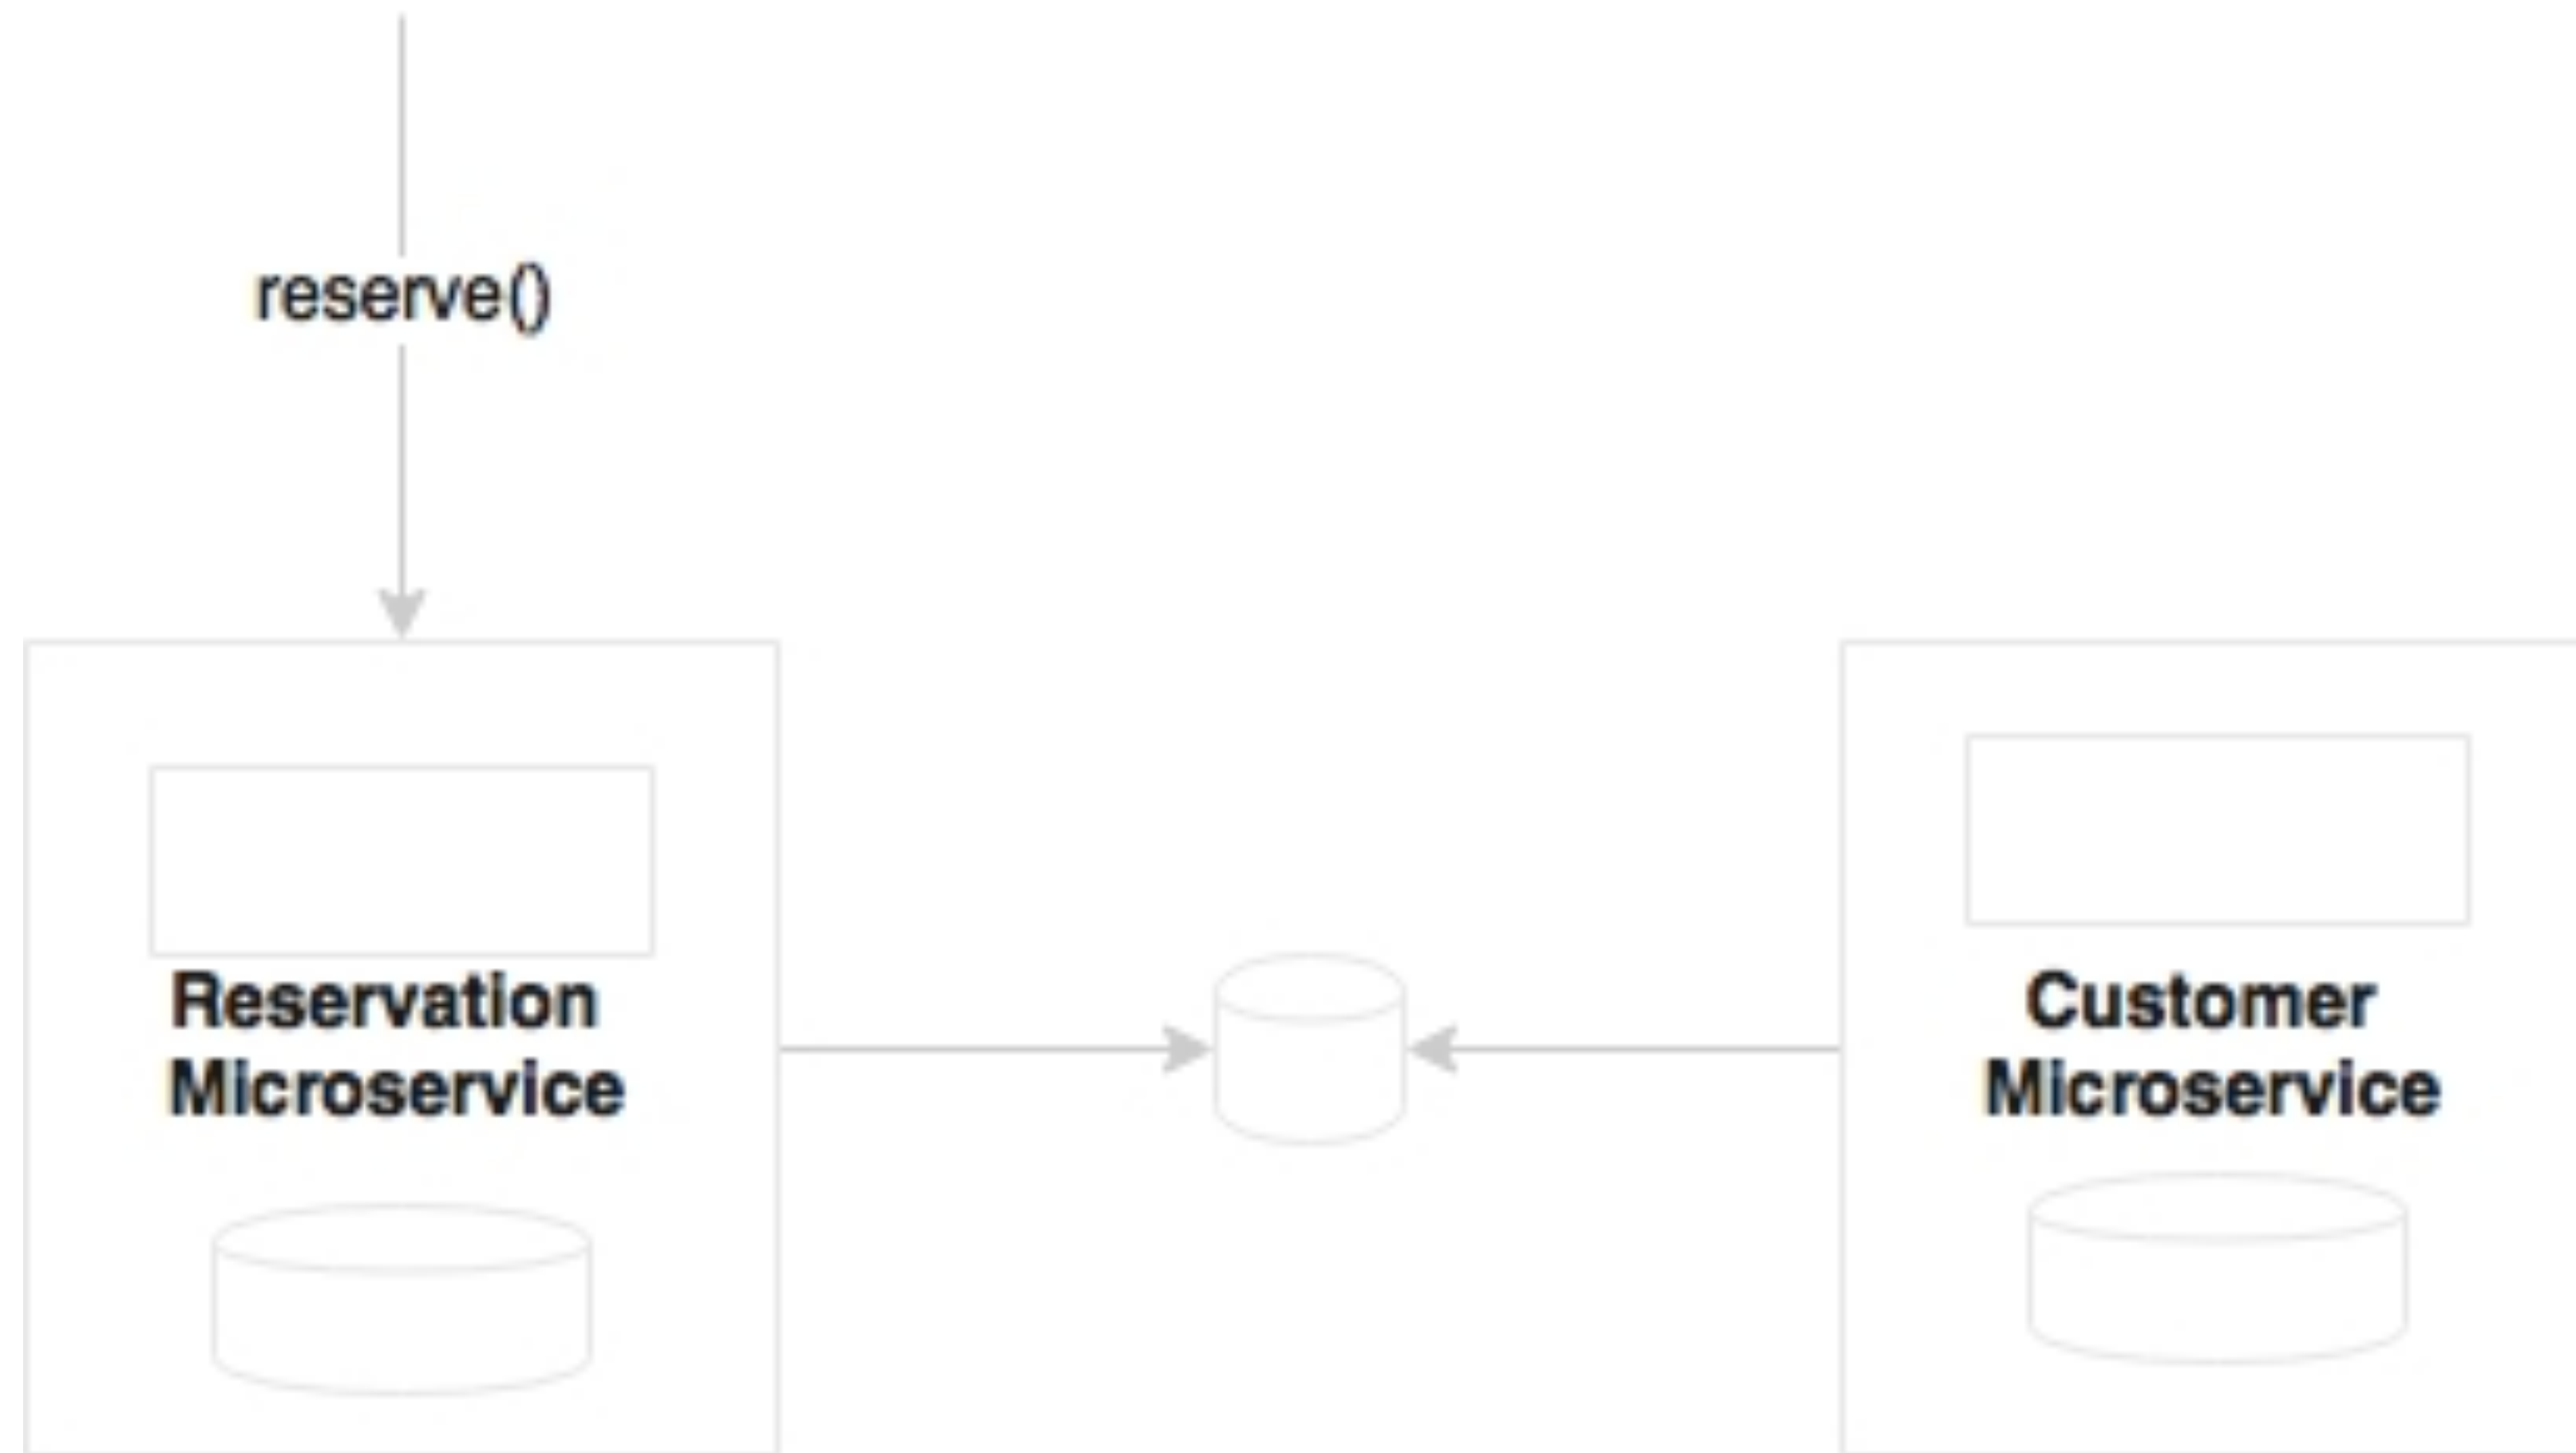
Orchestration of microservices

- Microservices are autonomous
 - all required components to complete their function should be within the service
- The service endpoints provide coarse-grained APIs
 - there are no external touch points required
 - microservices may need to talk to other microservices to fulfil their function

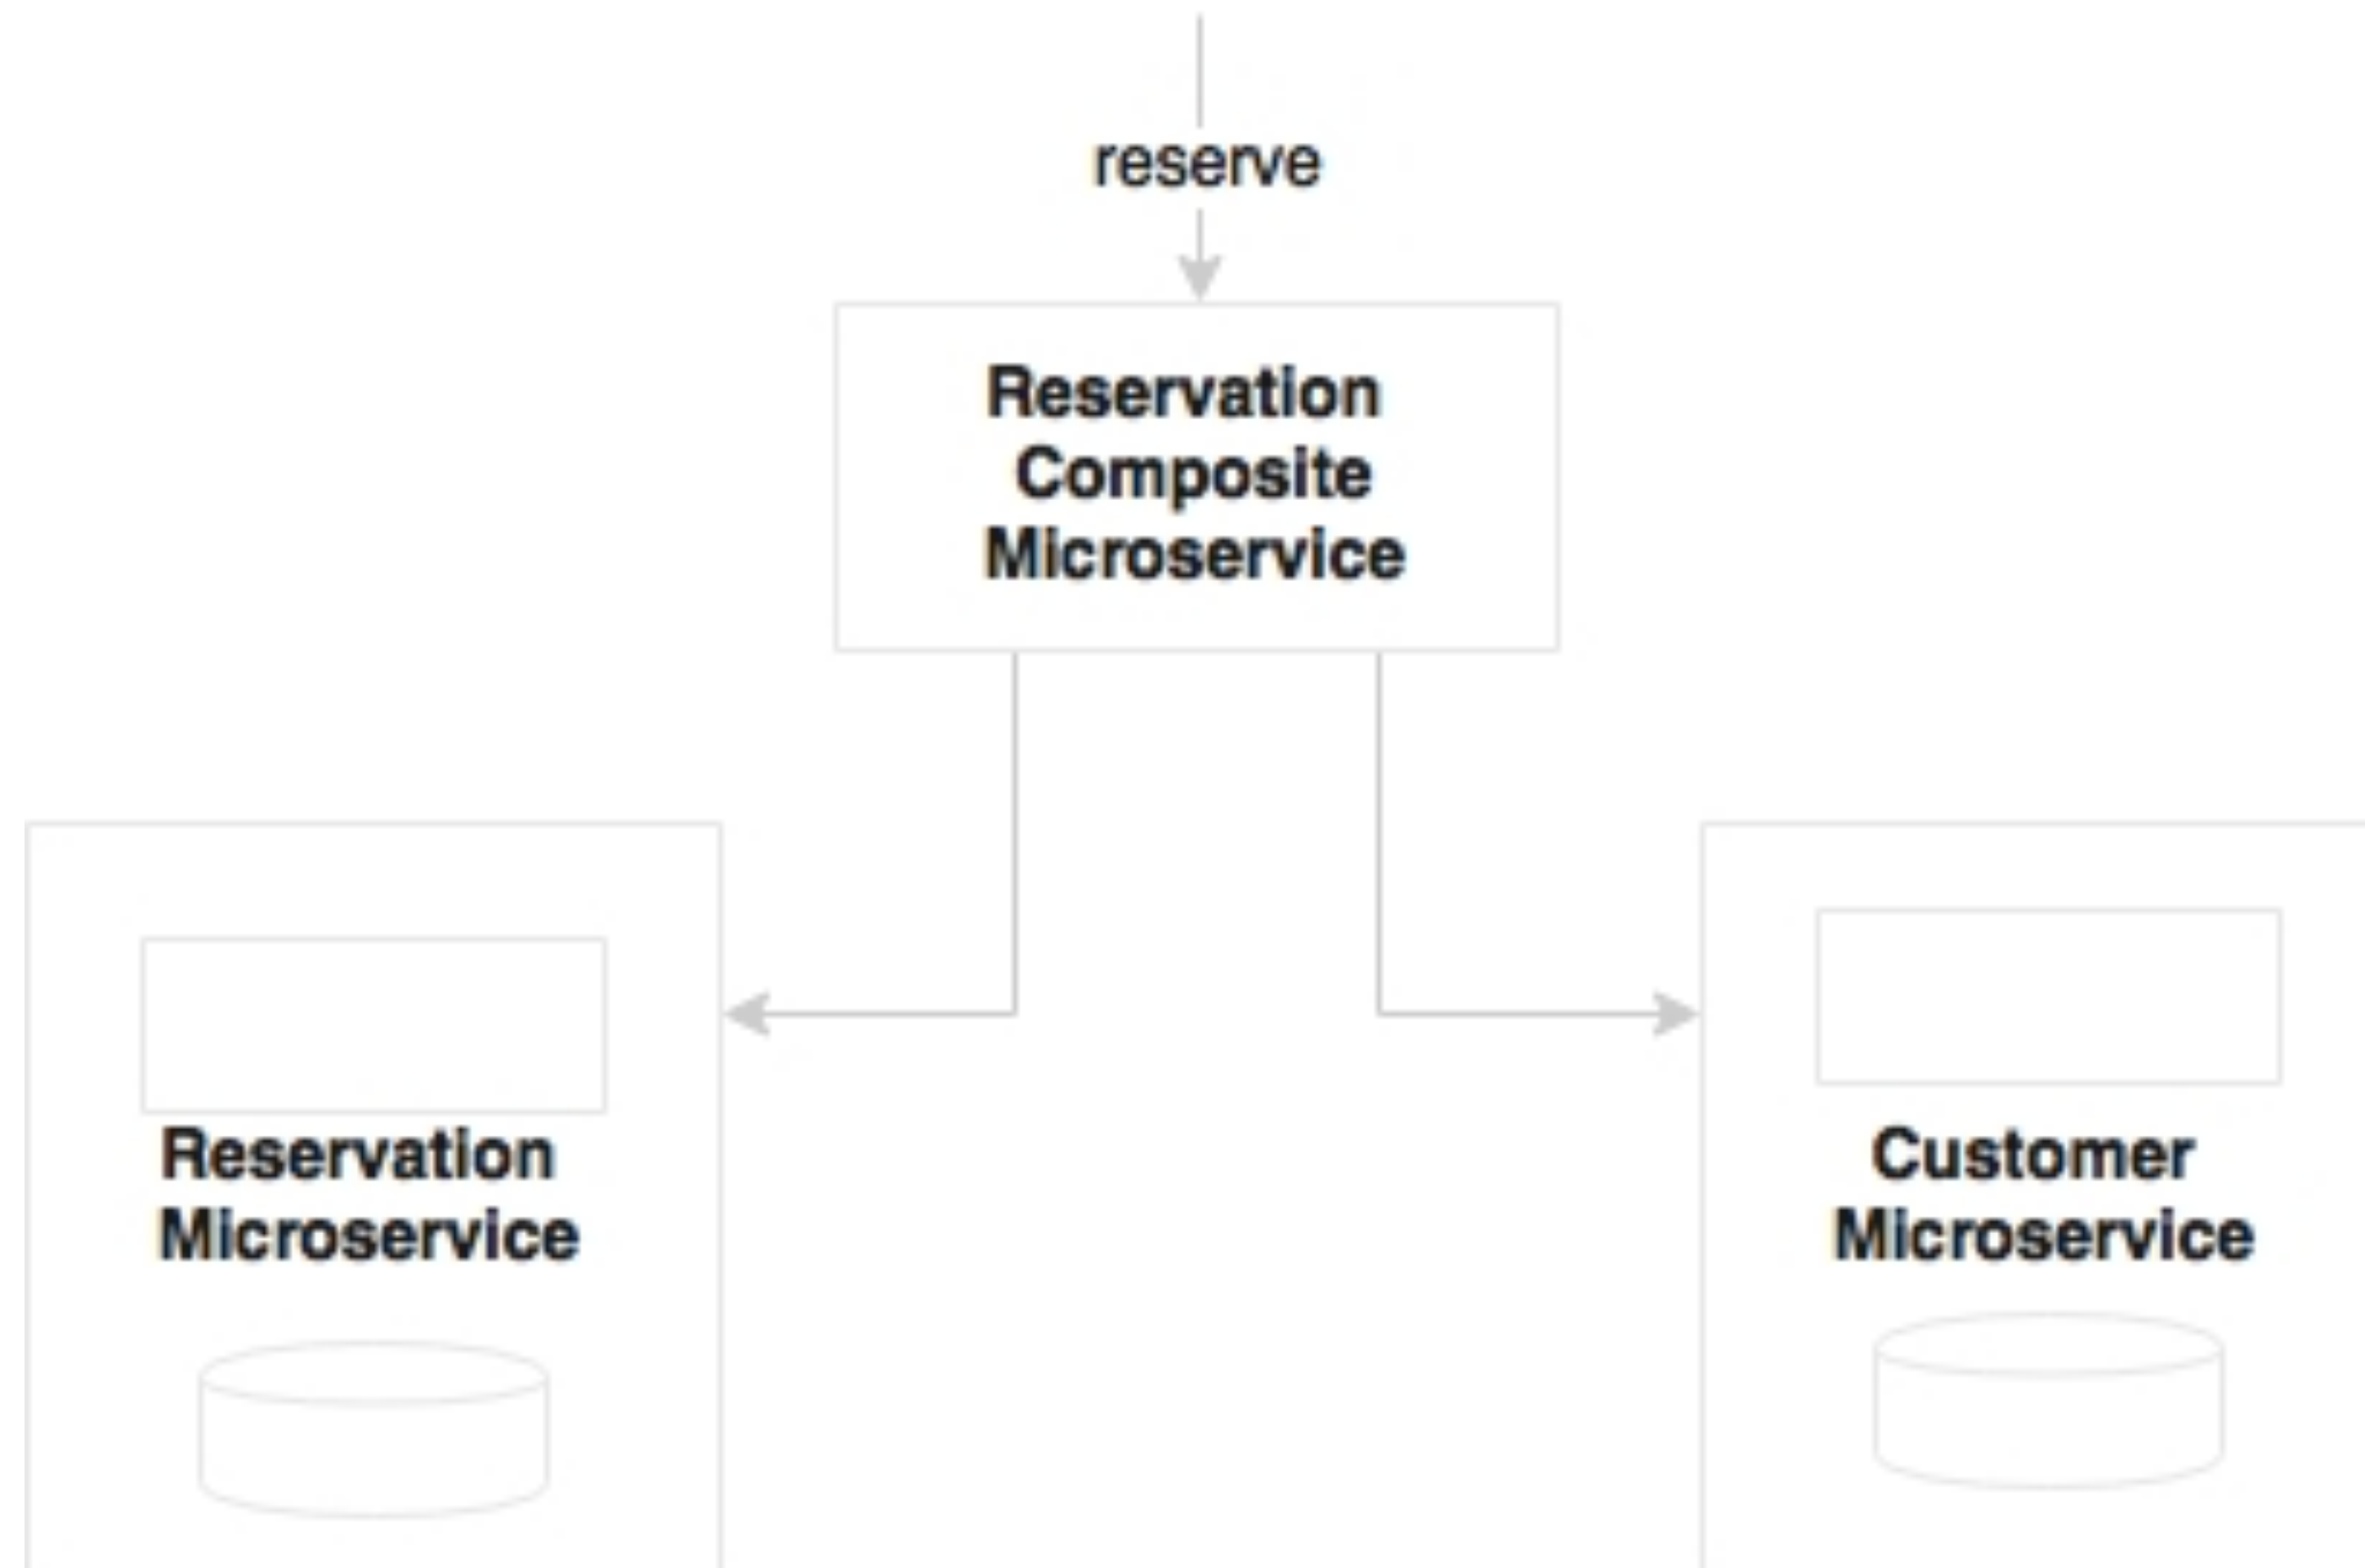
We can model choreography in all cases?



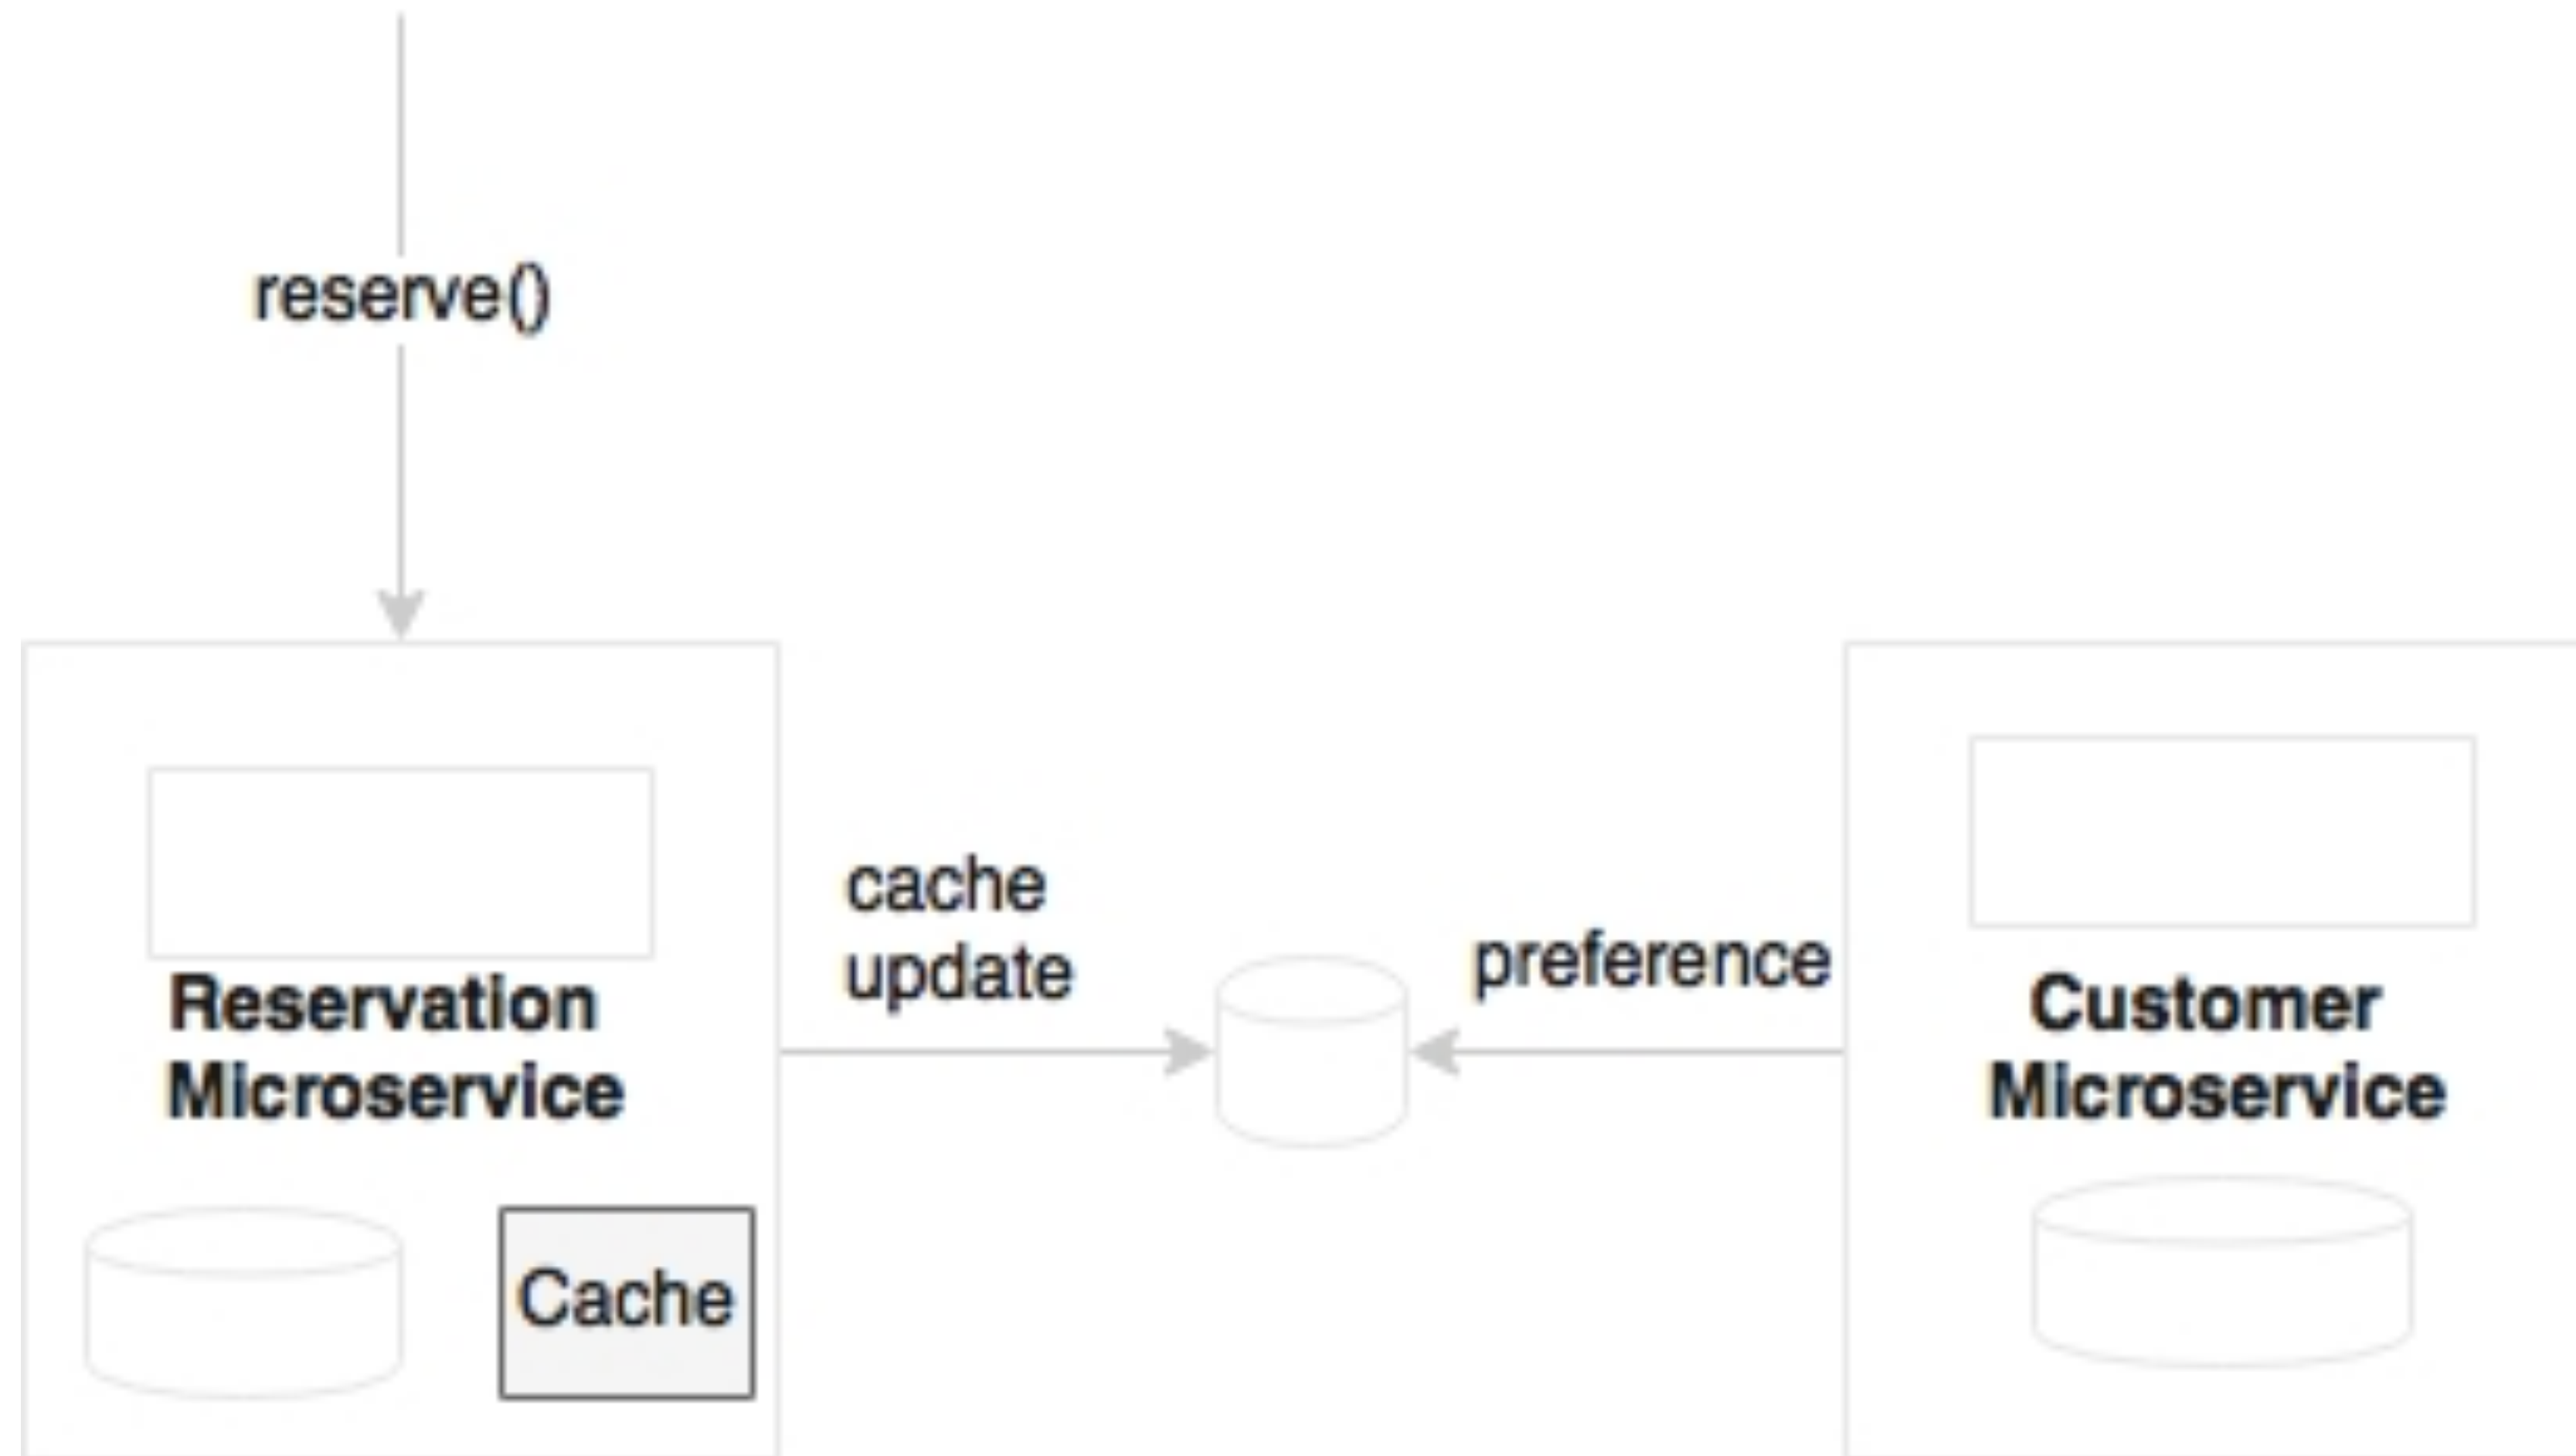
We can model choreography in all cases?



We can model choreography in all cases?

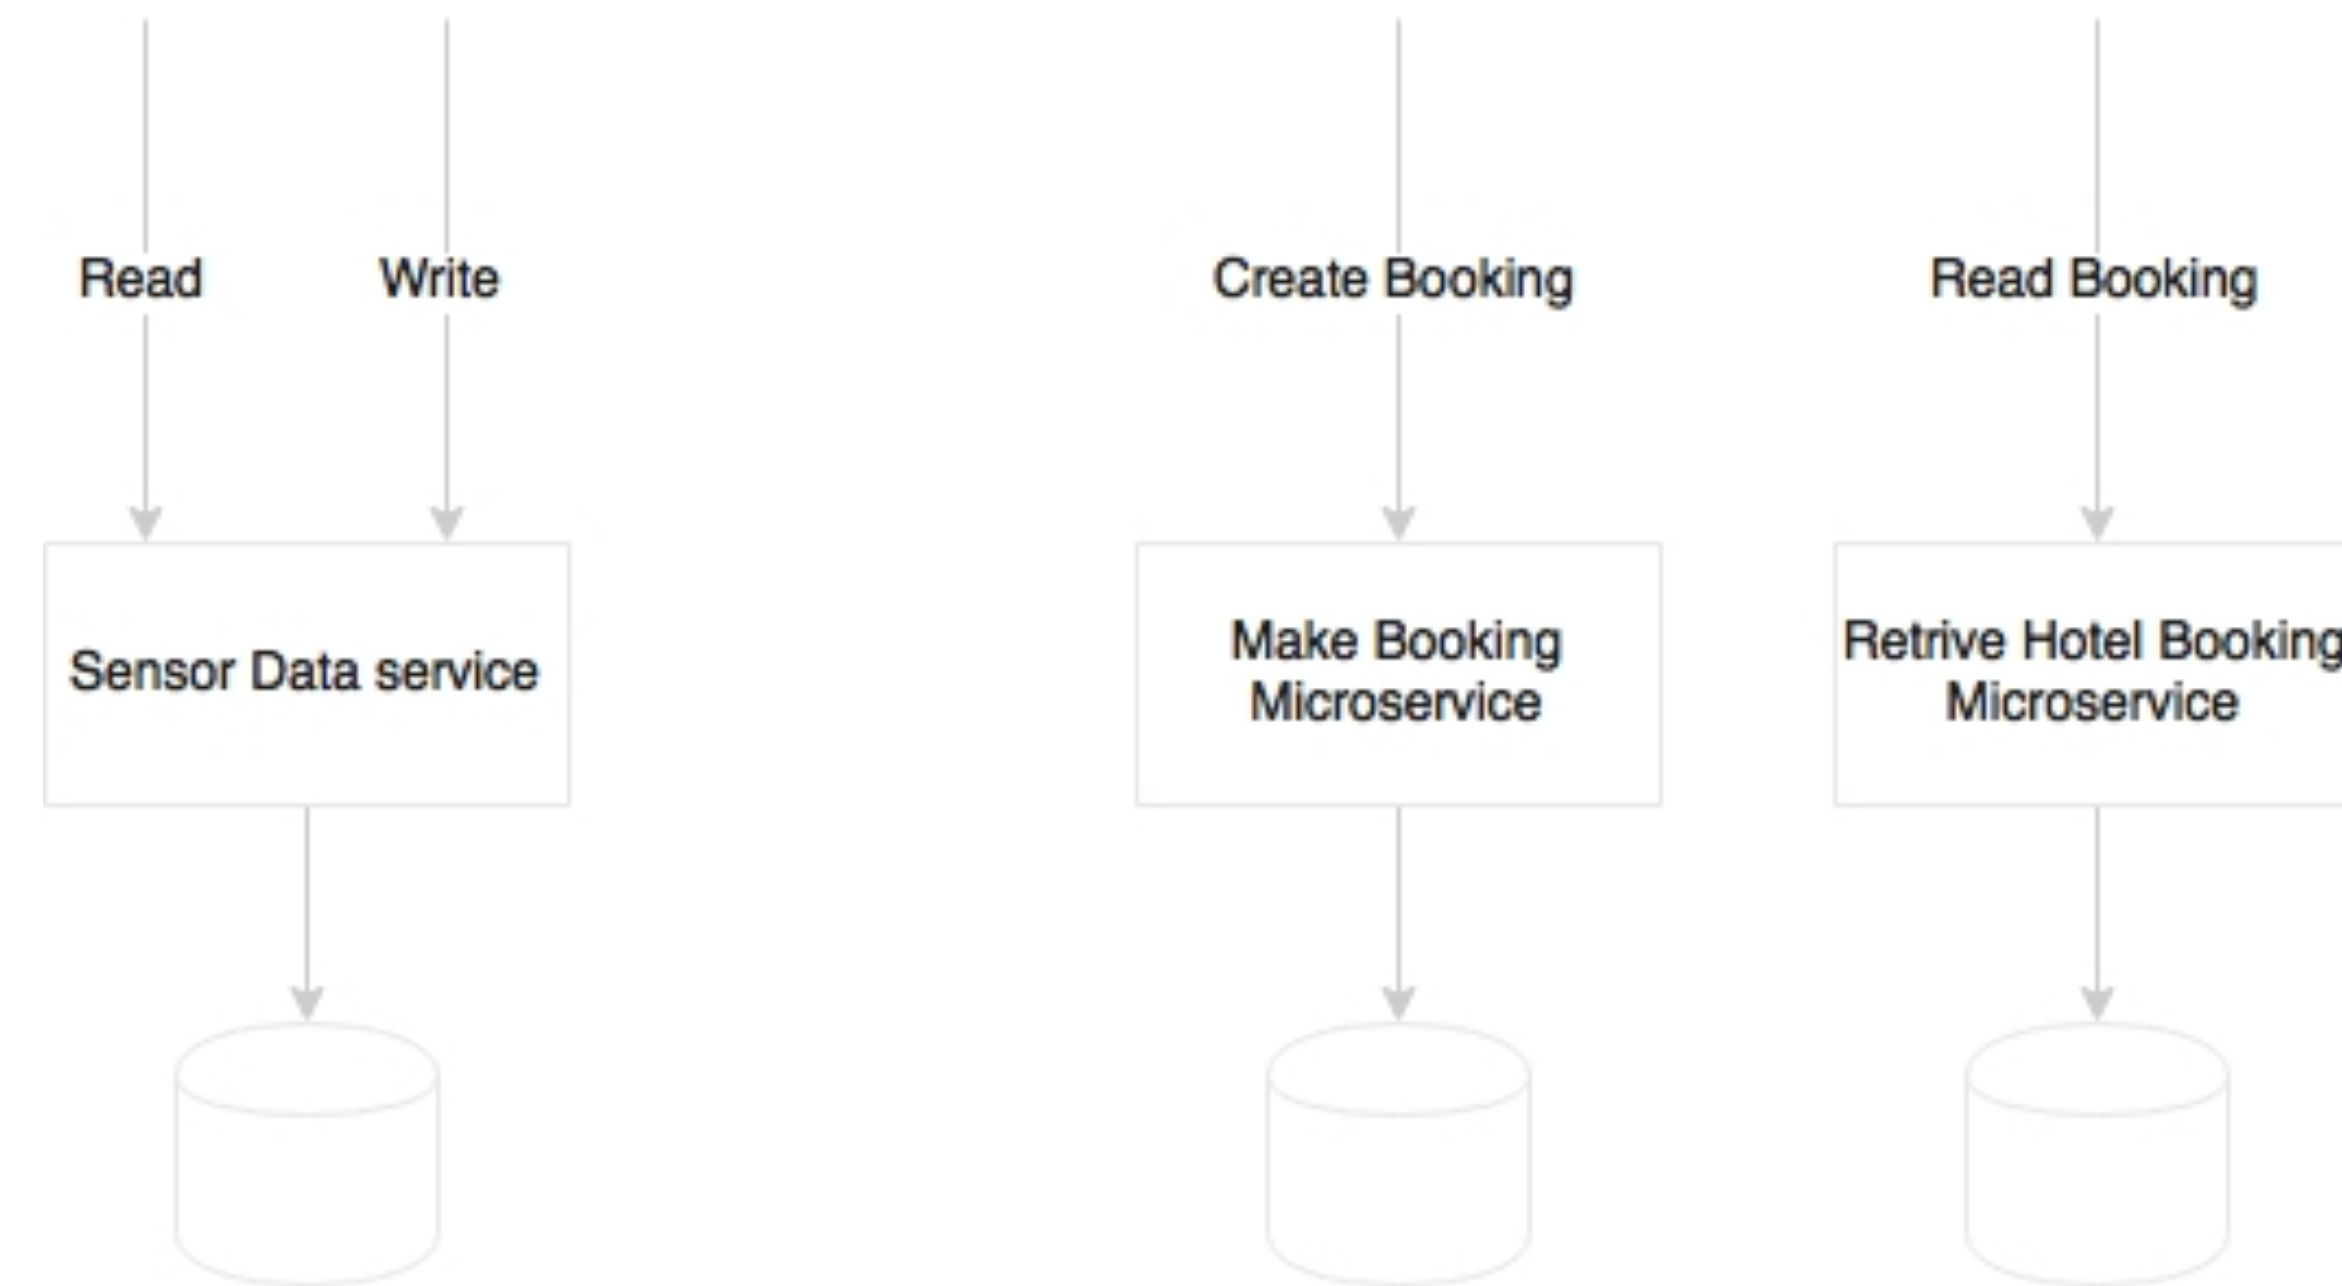


We can model choreography in all cases?



How many endpoints in a microservice?

- The question really is whether to limit each microservice with one endpoint or multiple endpoints
- Polyglot architecture could be another scenario where we may split endpoints into different microservices



One microservice per VM or multiple?

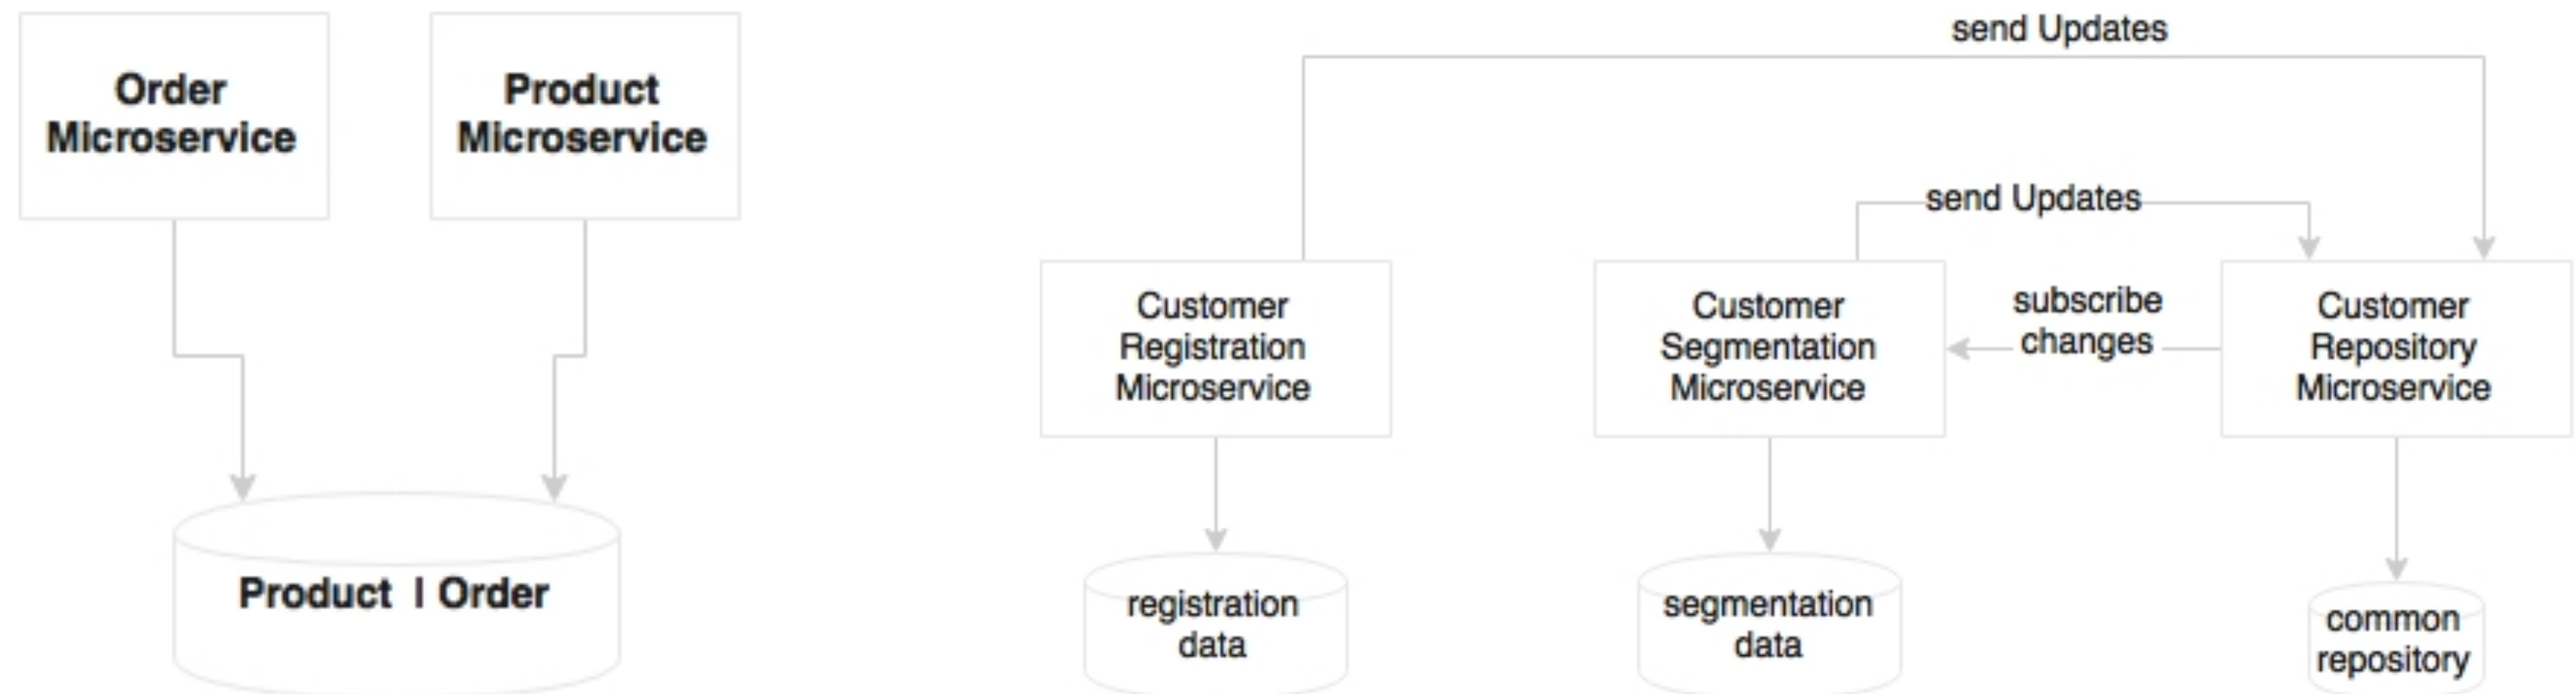
- Whether multiple microservices could be deployed in one virtual machine?
- The simplest way to approach this problem is to ask a few questions:
 - Does the VM have enough capacity to run both services under peak usage?
 - Do we want to treat these services differently to achieve SLAs (selective scaling)? For example, for scalability, if we have an all-in-one VM, we will have to replicate VMs which replicate all services.
 - Are there any conflicting resource requirements? For example, different OS versions, JDK versions, and others

Rules engine – shared or embedded?

- Either we hand code rules, or we may use a rules engine. Many enterprises manage rules centrally in a rules repository as well as execute them centrally
- **Drools** is one of the popular open source rules engines

Can microservices share data stores?

- In principle, microservices should abstract presentation, business logic, and data stores.
- If the services are broken as per the guidelines, each microservice logically **could** use an independent database



Setting up transaction boundaries

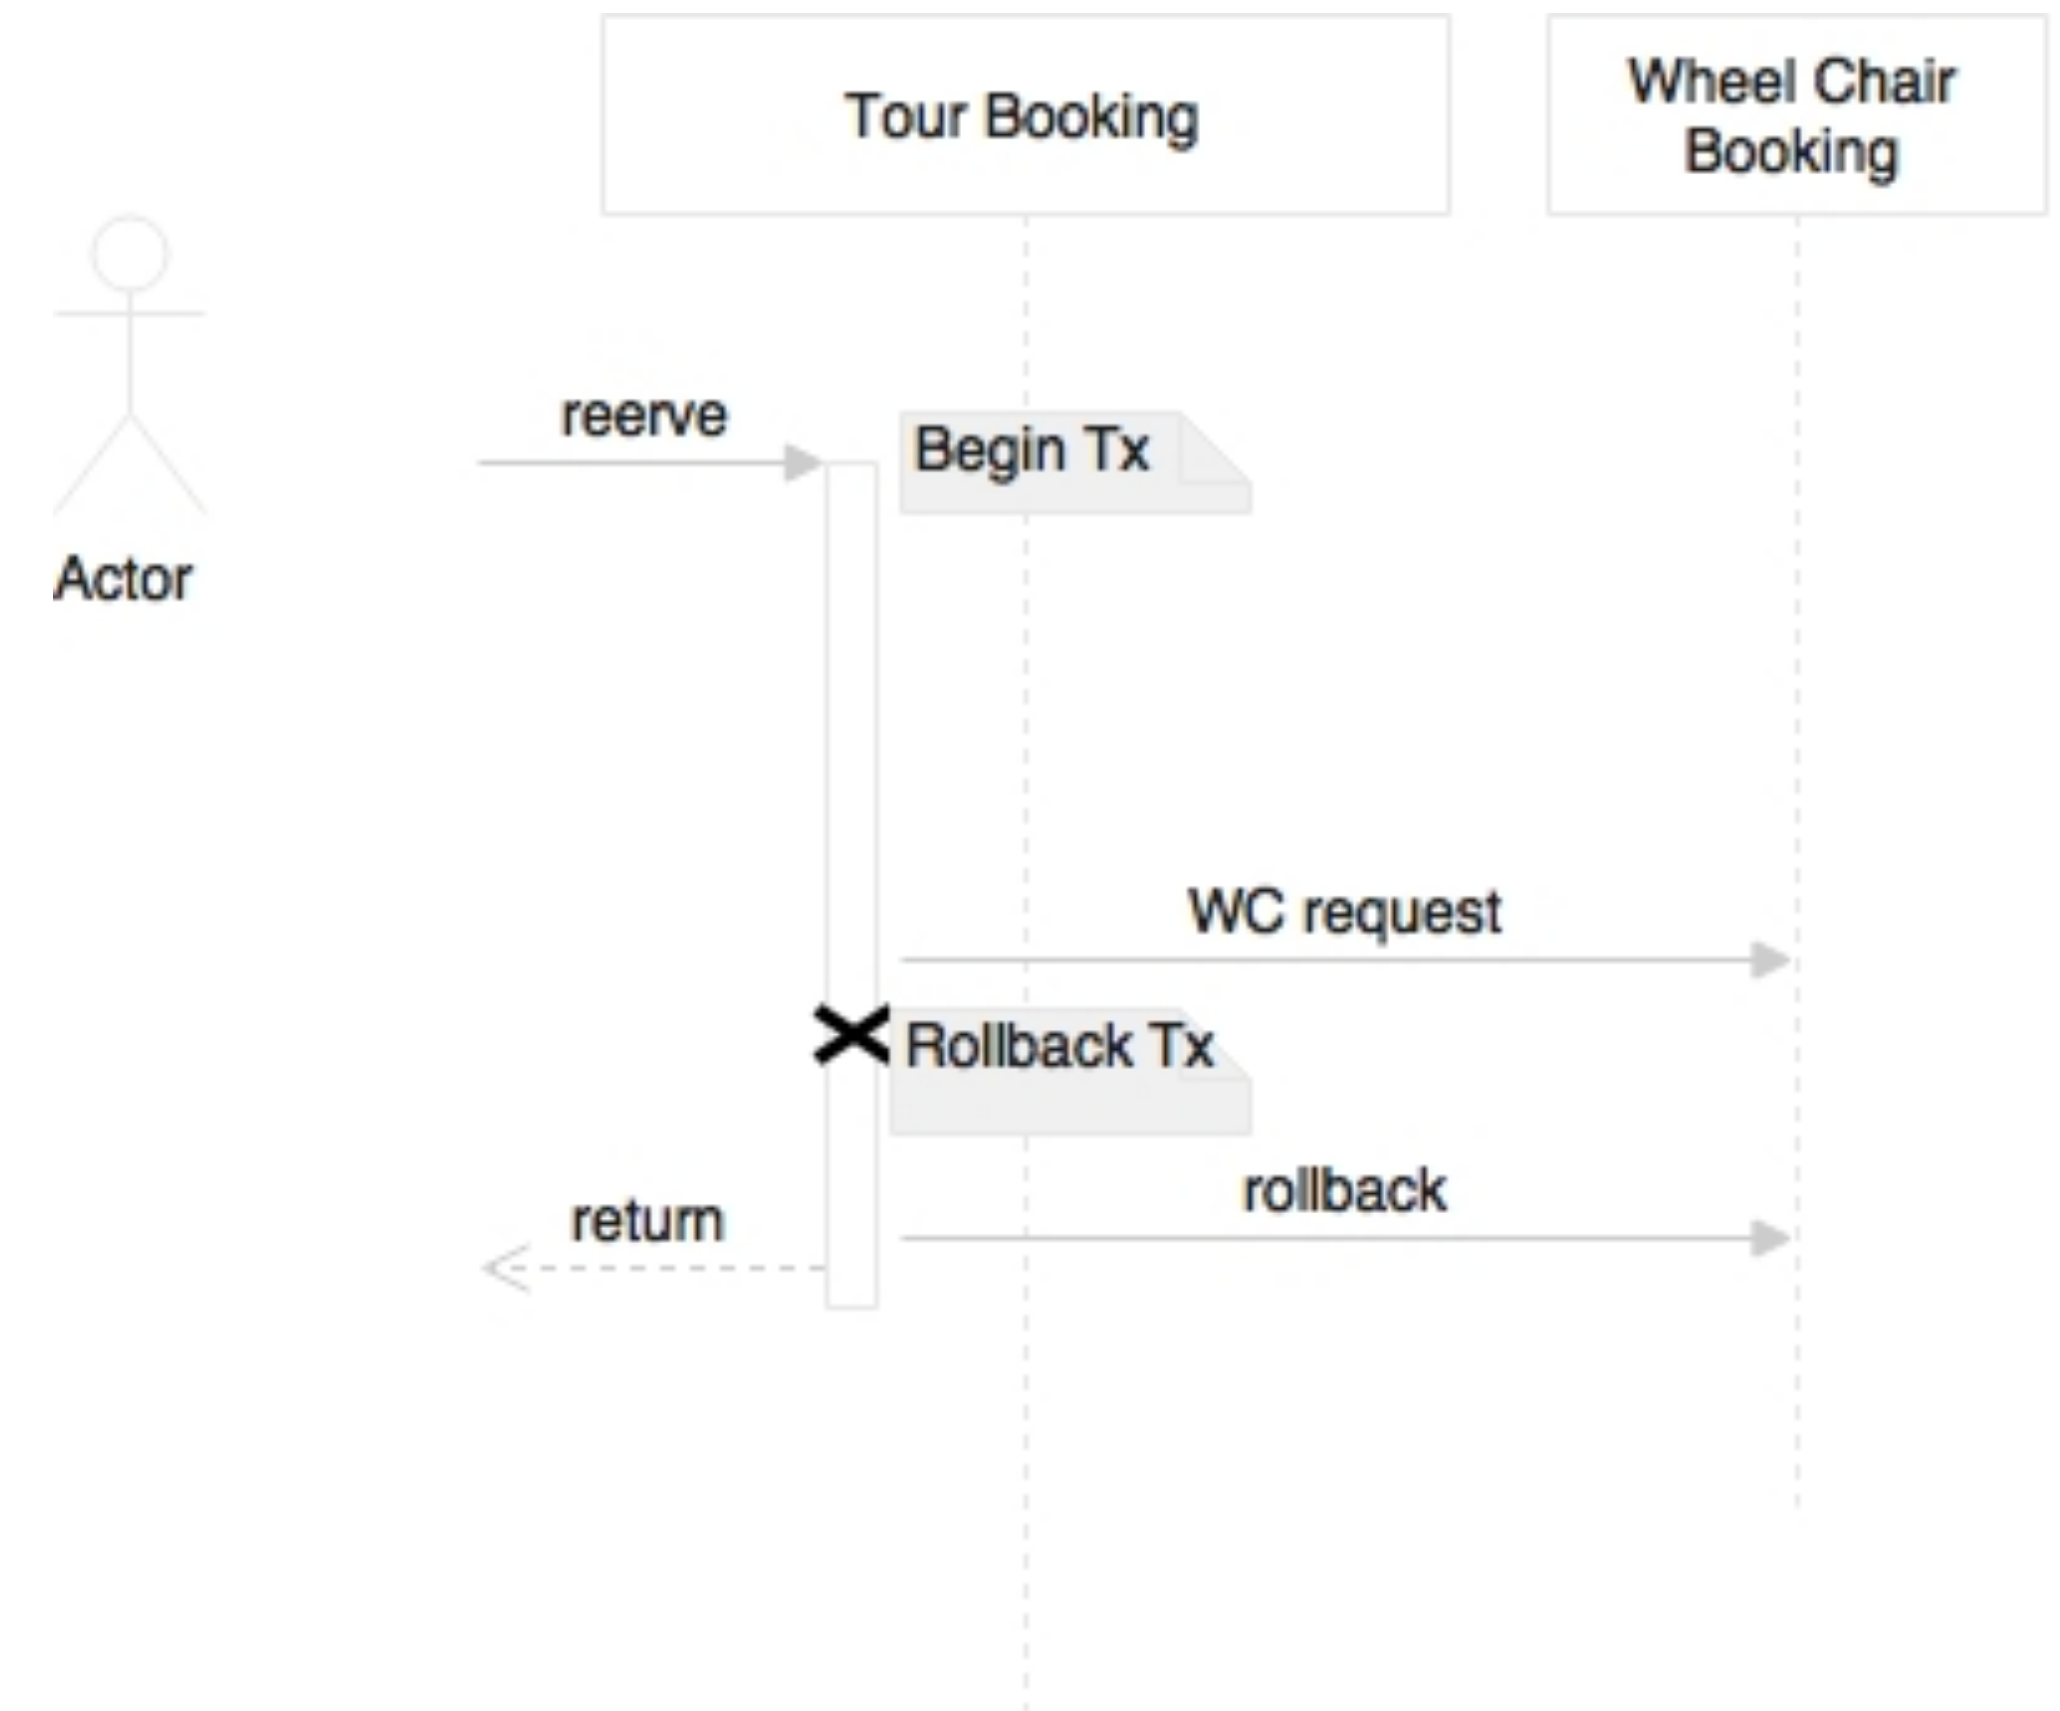
- Is there a place for transactions in microservices?
 - It is appropriate to define transaction boundaries within the microsystem using local transactions
 - However, distributed global transactions should be avoided in the microservices context

Altering use cases to simplify transactional requirements

- **Eventual consistency** is a better option than distributed transactions that span across multiple microservices
 - reduces a lot of overheads
 - but application developers may need to **re-think** the way they **write** application code

Distributed transaction scenarios

- The ideal scenario is to use local transactions within a microservice if required, and completely avoid distributed transactions
- There could be scenarios where at the end of the execution of one service, we may want to send a message to another microservice



Service endpoint design consideration

- Service design has two key elements: contract design and protocol selection
- **Contract design**
 - A **complex** service contract **reduces** the usability of the service
 - KISS (Keep It Simple Stupid)
 - YAGNI (You Ain't Gonna Need It)
 - Evolutionary design is a great concept
 - Consumer Driven Contracts (CDC) and Postel's law

Service endpoint design consideration

- Protocol selection

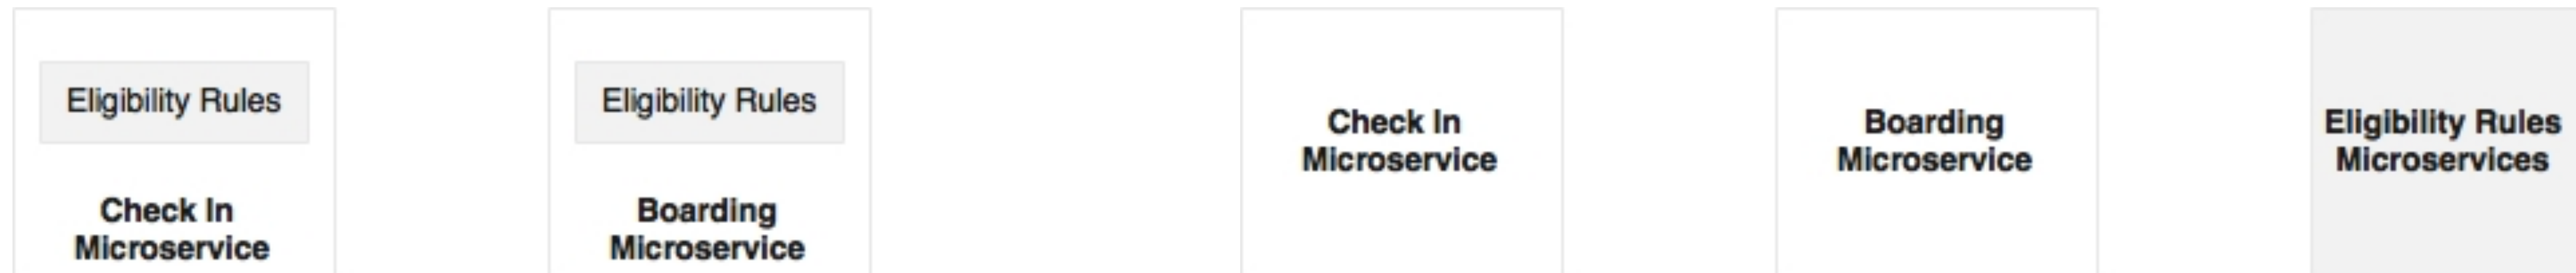
- In the SOA world, HTTP/SOAP, and messaging were kinds of default service protocols for service interactions
- increases the communication cost; susceptible to network failures; could result in poor performance of services
 - Message-oriented services (asynchronous style of communication)
 - HTTP and REST endpoints (interoperability, protocol handling, traffic routing, load balancing, security systems...)
 - Optimized communication protocols
 - API documentations

Homework 4.1

- What is the main concepts and principles of Consumer Driven Contracts (CDC)?
- How does it fit into the microservices approach?
- How Postel's law could be also relevant in this scenario?
- What are the main solutions (patterns, architectural styles, tools, etc.) to implement:
 - Message-oriented services;
 - HTTP and REST endpoints;
 - Optimized communication protocols;
 - API documentations

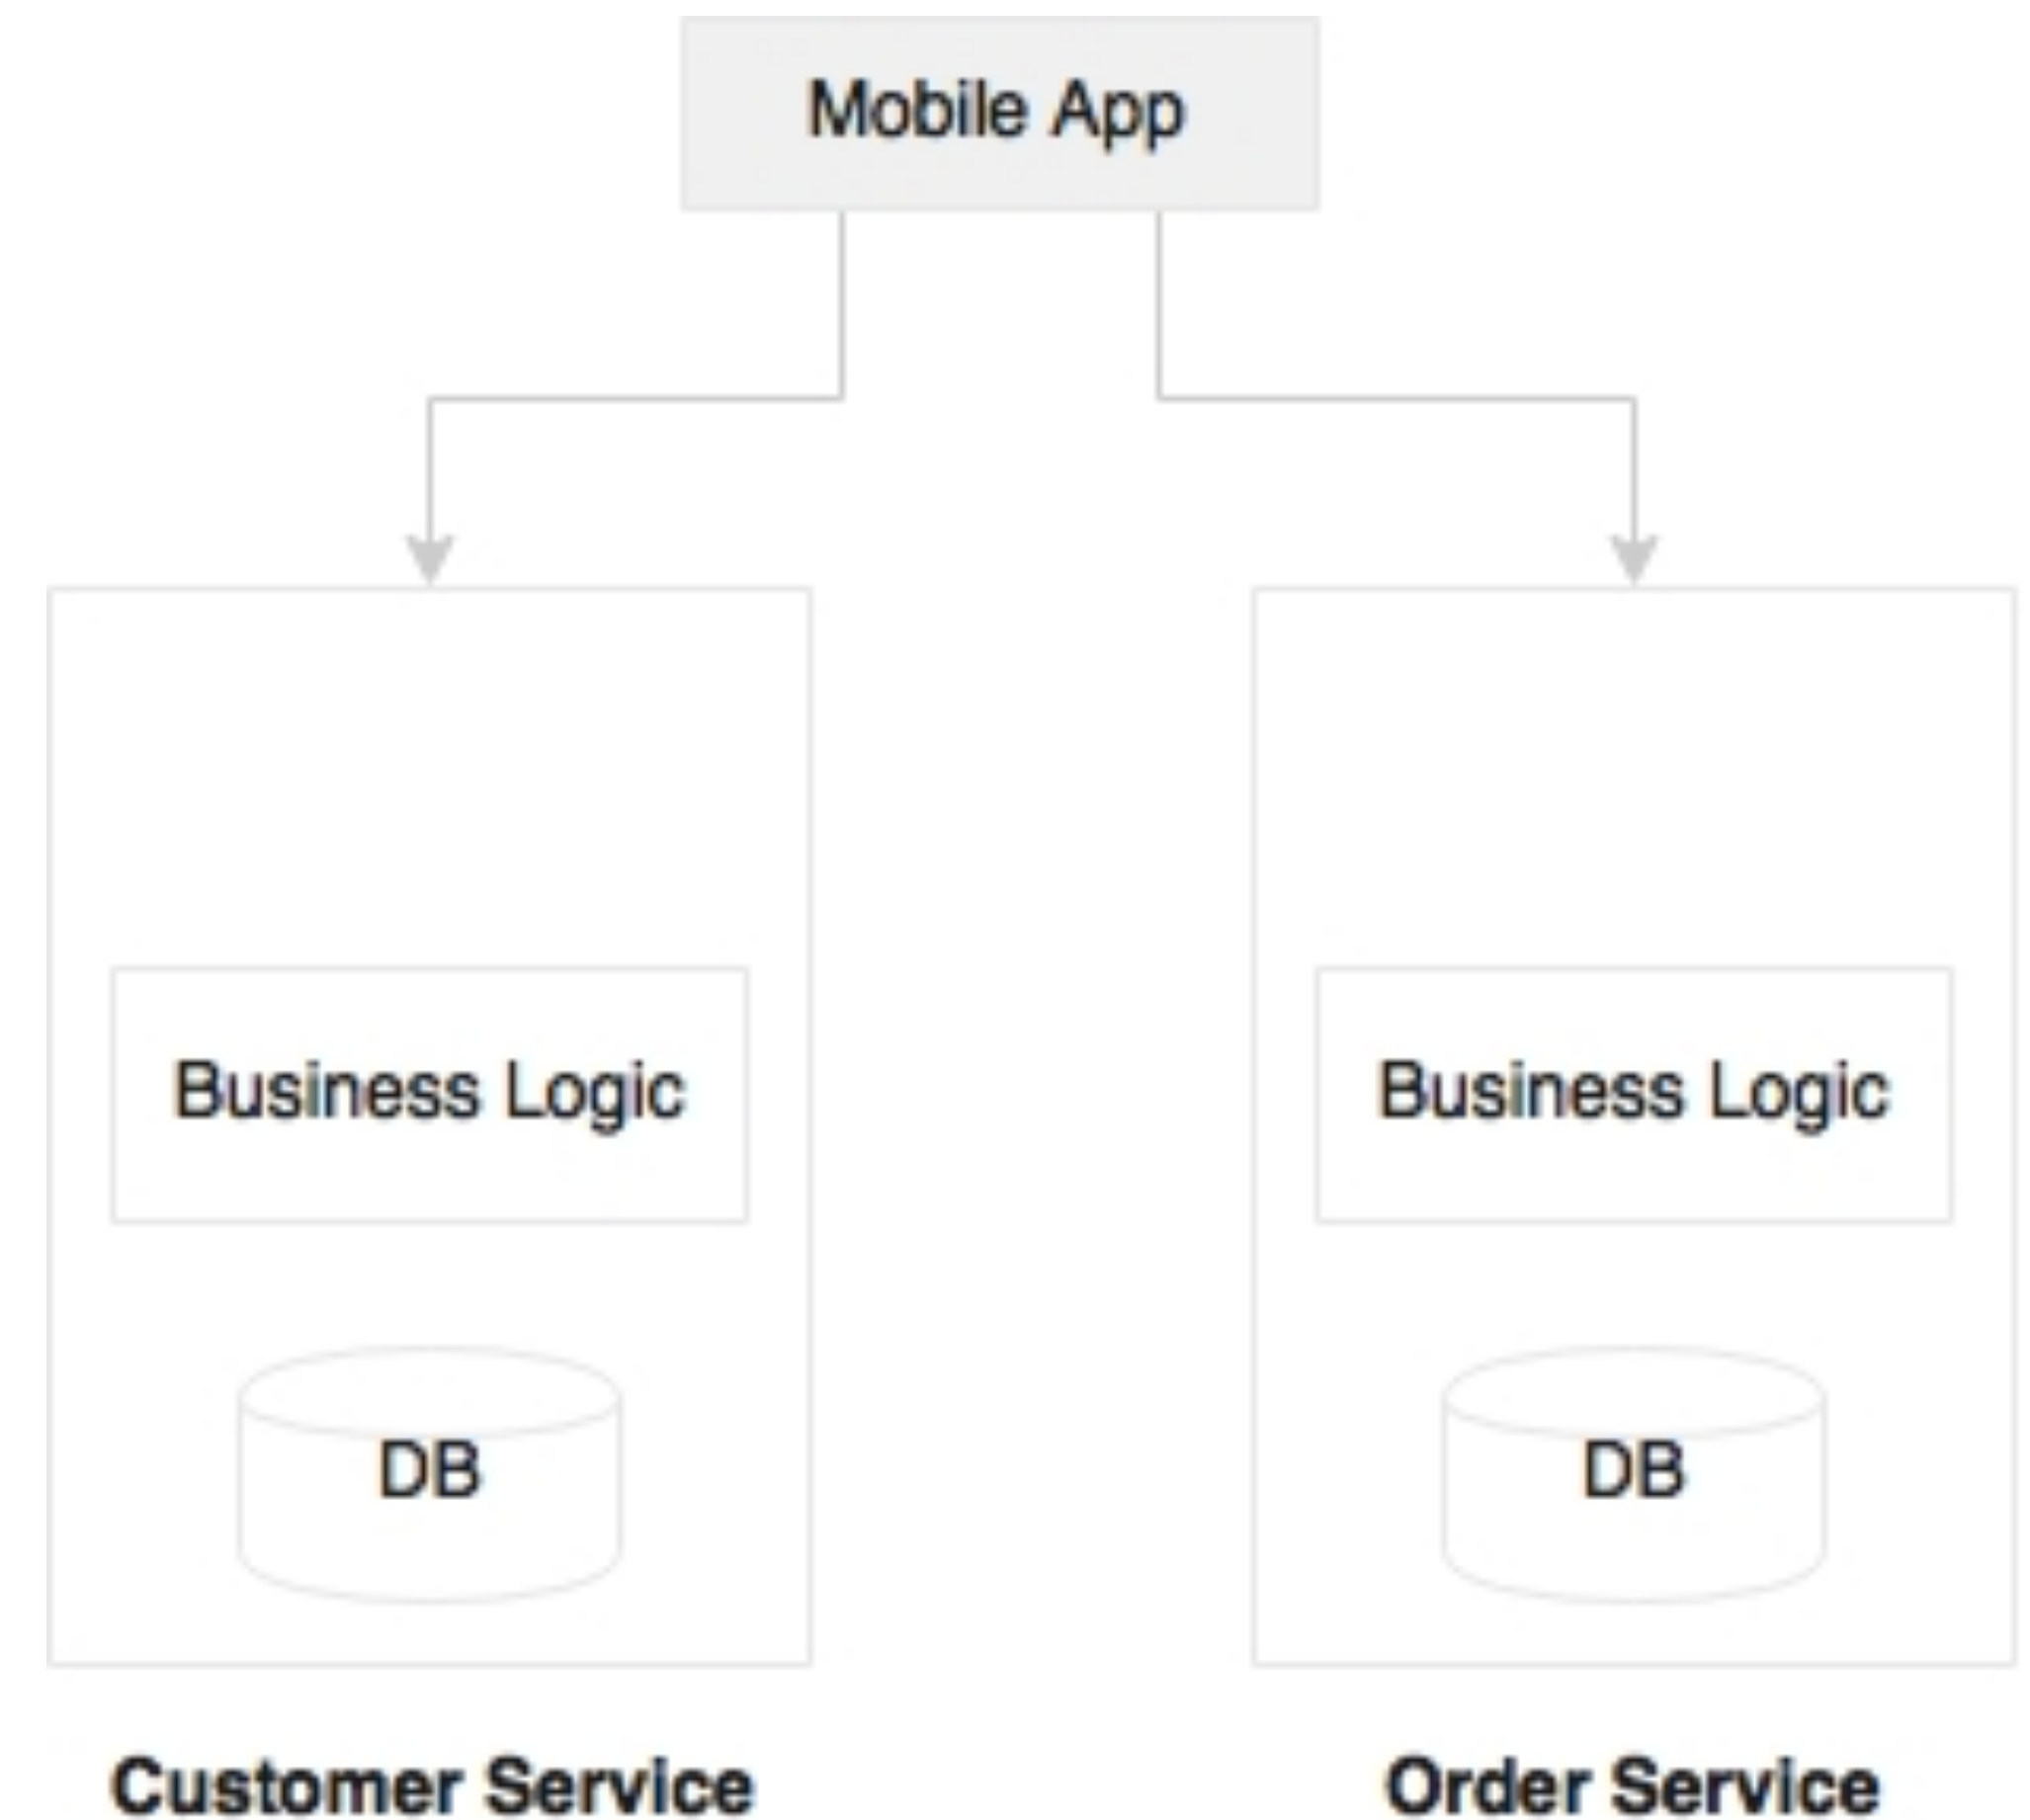
Handling shared libraries

- Autonomous and self-contained
 - duplicate code and libraries



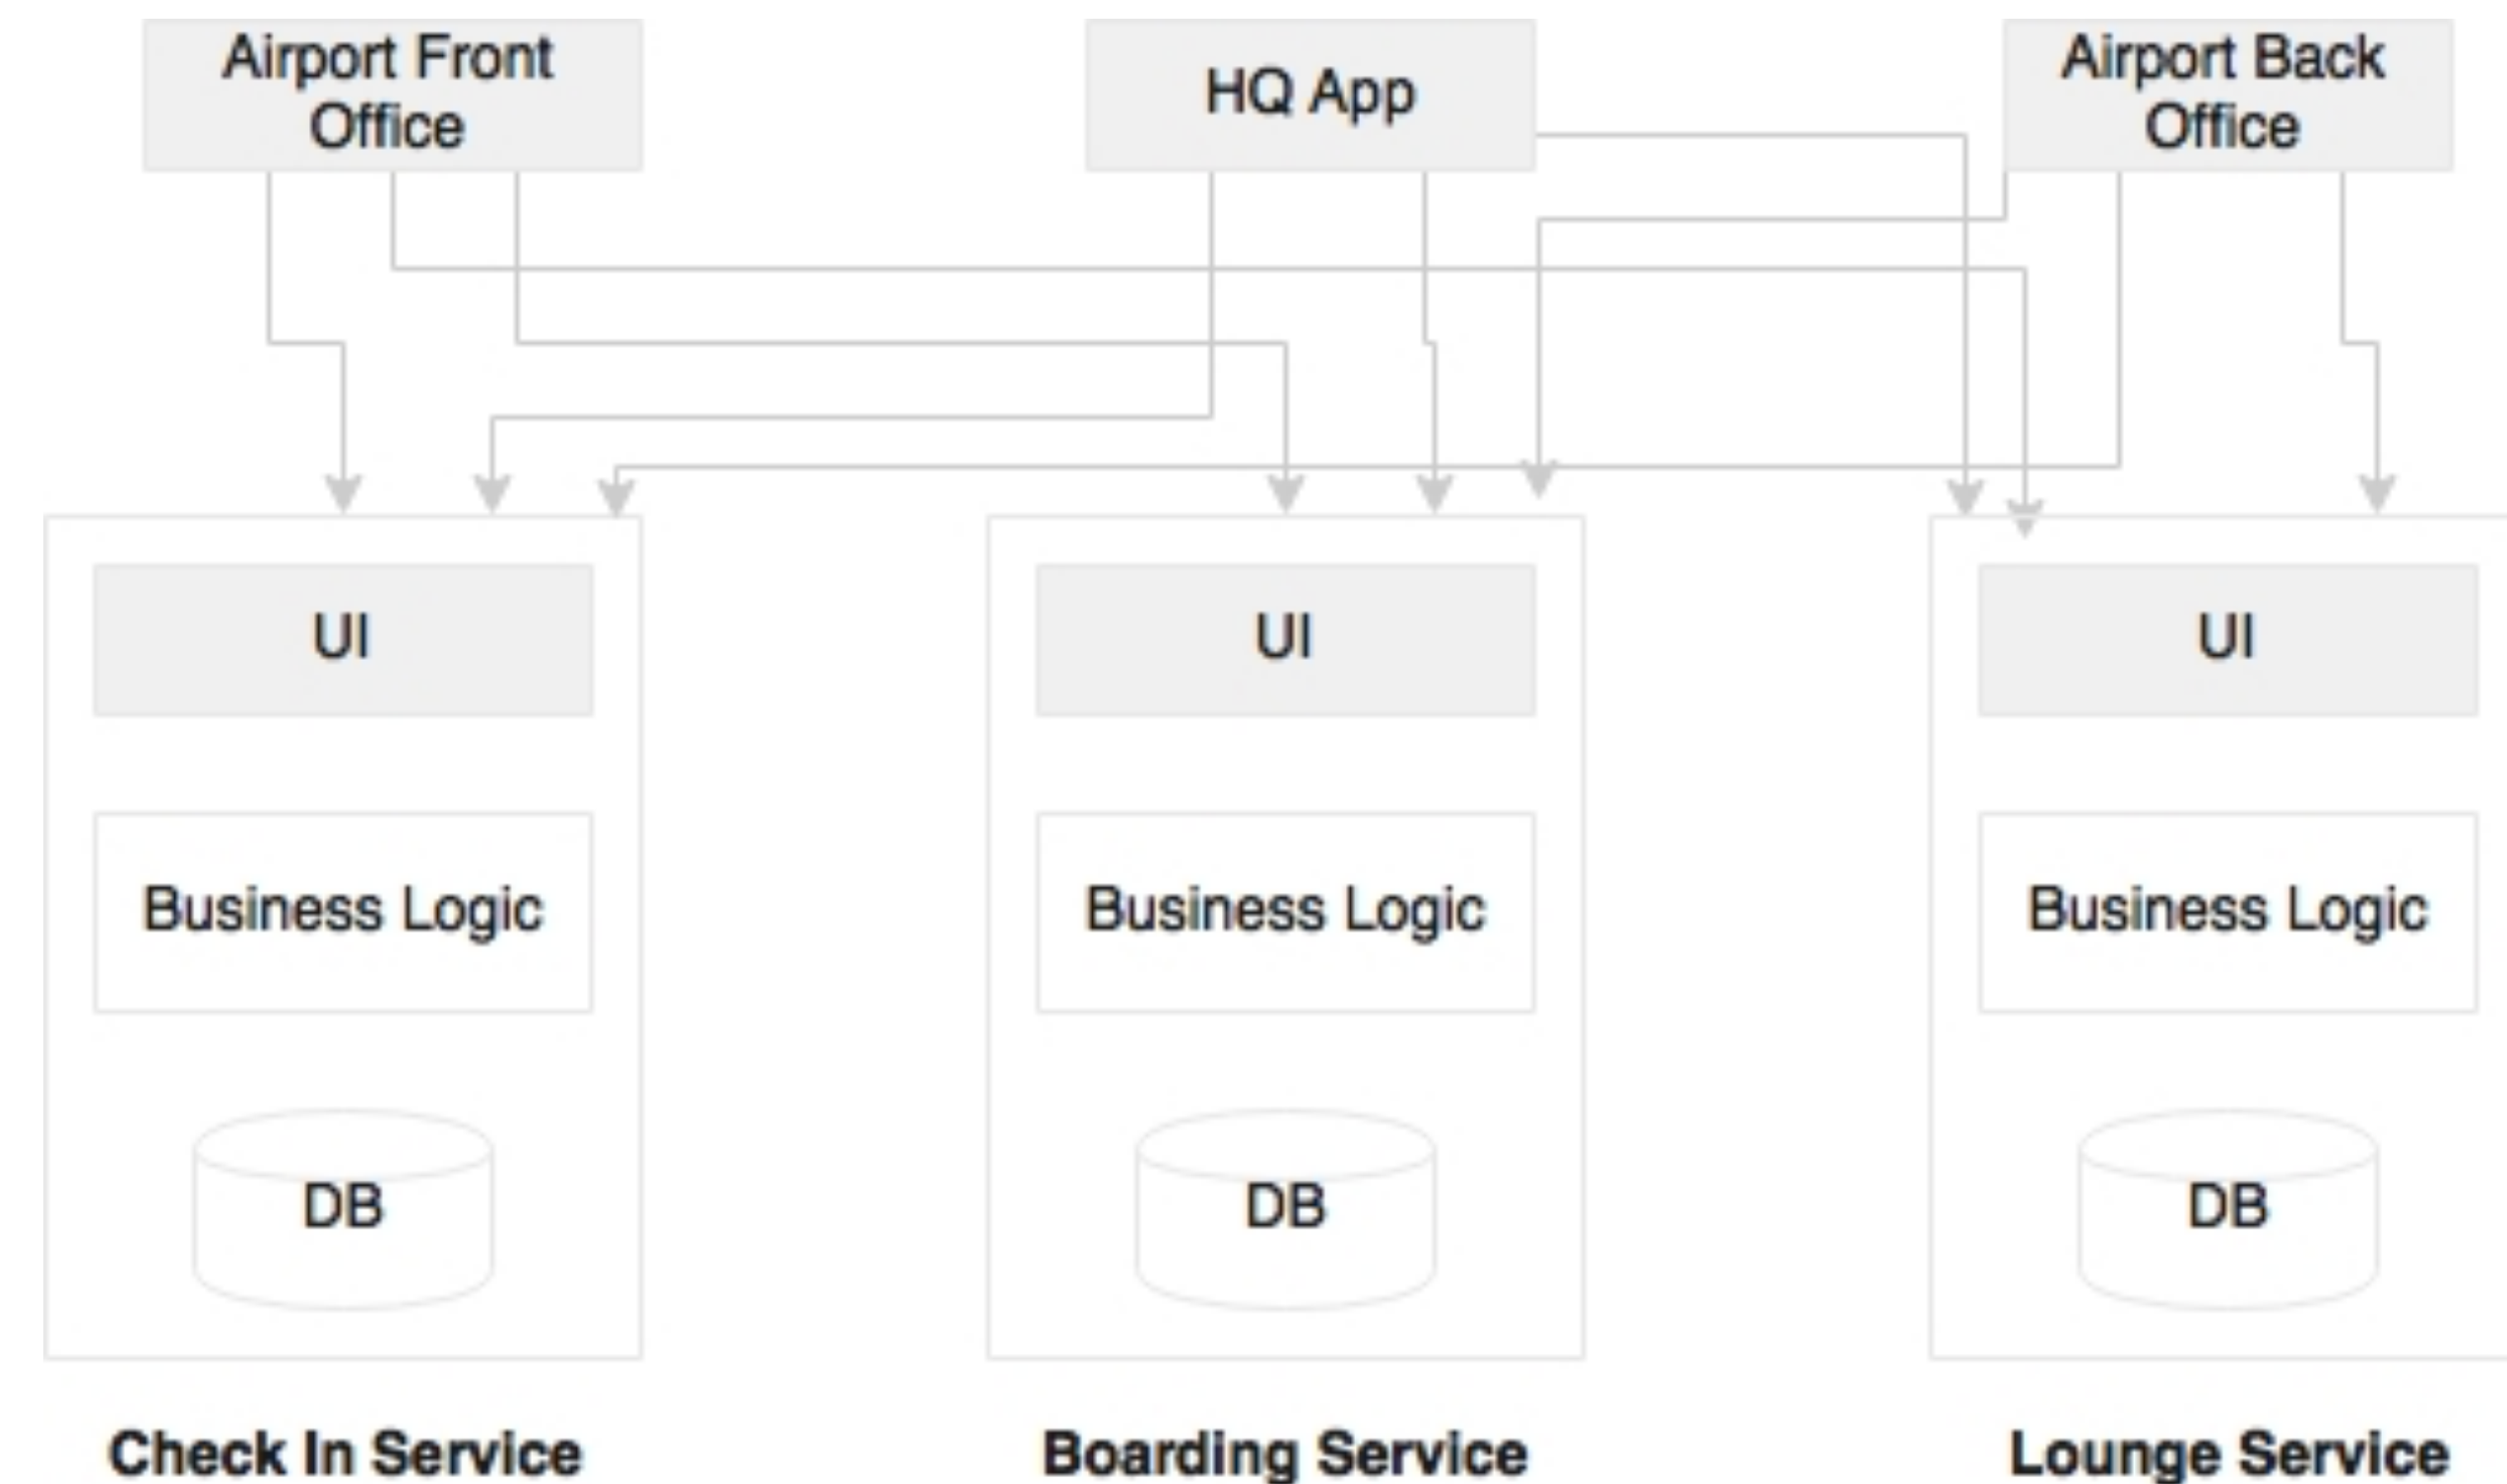
User interfaces in microservices

- The microservices principle advocates a microservice as a **vertical slice** from the database to presentation
- build quick UI and mobile applications mashing up the existing APIs



User interfaces in microservices

- Build consolidated web applications targeted to communities
- One approach is to build a **container web application** or a **placeholder web application**, which links to multiple microservices at the backend
- multiple placeholder web applications targeting different user communities



Use of API gateways in microservices

- With the advancement of client-side JavaScript frameworks, the server is expected to expose RESTful services
 - mismatch in contract expectations
 - minimal information is sent with links (HATEOAS)
- multiple calls to the server to render a page
 - used when the client makes the REST call (also sends the required fields as part of the query string)

Use of API gateways in microservices

- Introduce a level of indirection
 - A gateway component sits between the client and the server, and transforms data as per the consumer's specification
 - do not compromise on the backend service contract
 - leads to **UI services**
 - In many cases, the API gateway acts as a **proxy** to the backend, exposing a **set of consumer-specific APIs**



Use of ESB and iPaaS with microservices

- Theoretically, SOA is not all about ESBs, but the reality is that ESBs have always been at the center of many SOA implementations
- What would be the role of an ESB in the microservices world?
 - Microservices: fully cloud native systems, lightweight characteristics of microservices enable automation of deployments, scaling, and so on...
 - ESB: heavyweight in nature, not cloud friendly , protocol mediation, transformation, orchestration, and application adaptors ~> we may not need

Use of ESB and iPaaS with microservices

- Enterprises have **legacy applications**, **vendor applications**, and so on
 - Legacy services use ESBs to connect with microservices
- With the advancement of clouds, the capabilities of ESBs are not sufficient to manage integration between clouds, cloud to on-premise, and so on.
 - **Integration Platform as a Service** (iPaaS) is evolving as the next generation application integration platform, which further reduces the role of ESBs.
 - In typical deployments, iPaaS invokes API gateways to access microservices

Homework 4.2

- How can we have the same features present in EBS's with lightweight tools in the universe of microservices?

Service versioning considerations

- Versioning helps us to release **new services** without **breaking** the existing consumers
- There are three different ways in which we can version **REST services**:
 - **URI versioning**: version number is included in the URL itself

```
/api/v3/customer/1234  
/api/customer/1234 - aliased to v3.
```

```
@RestController("CustomerControllerV3")  
@RequestMapping("api/v3/customer")  
public class CustomerController {  
  
}
```

```
api/customer/100?v=1.5
```

- **Media type versioning**: version is set by the client on the HTTP Accept header

```
Accept: application/vnd.company.customer-v3+json
```


Service versioning considerations

- A less effective approach for versioning is to set the version in the **custom header**

```
@RequestMapping(value =("/{id}", method = RequestMethod.GET, headers = {"version=3"})
public Customer getCustomer(@PathVariable("id") long id,
    //other code goes here.
}
```

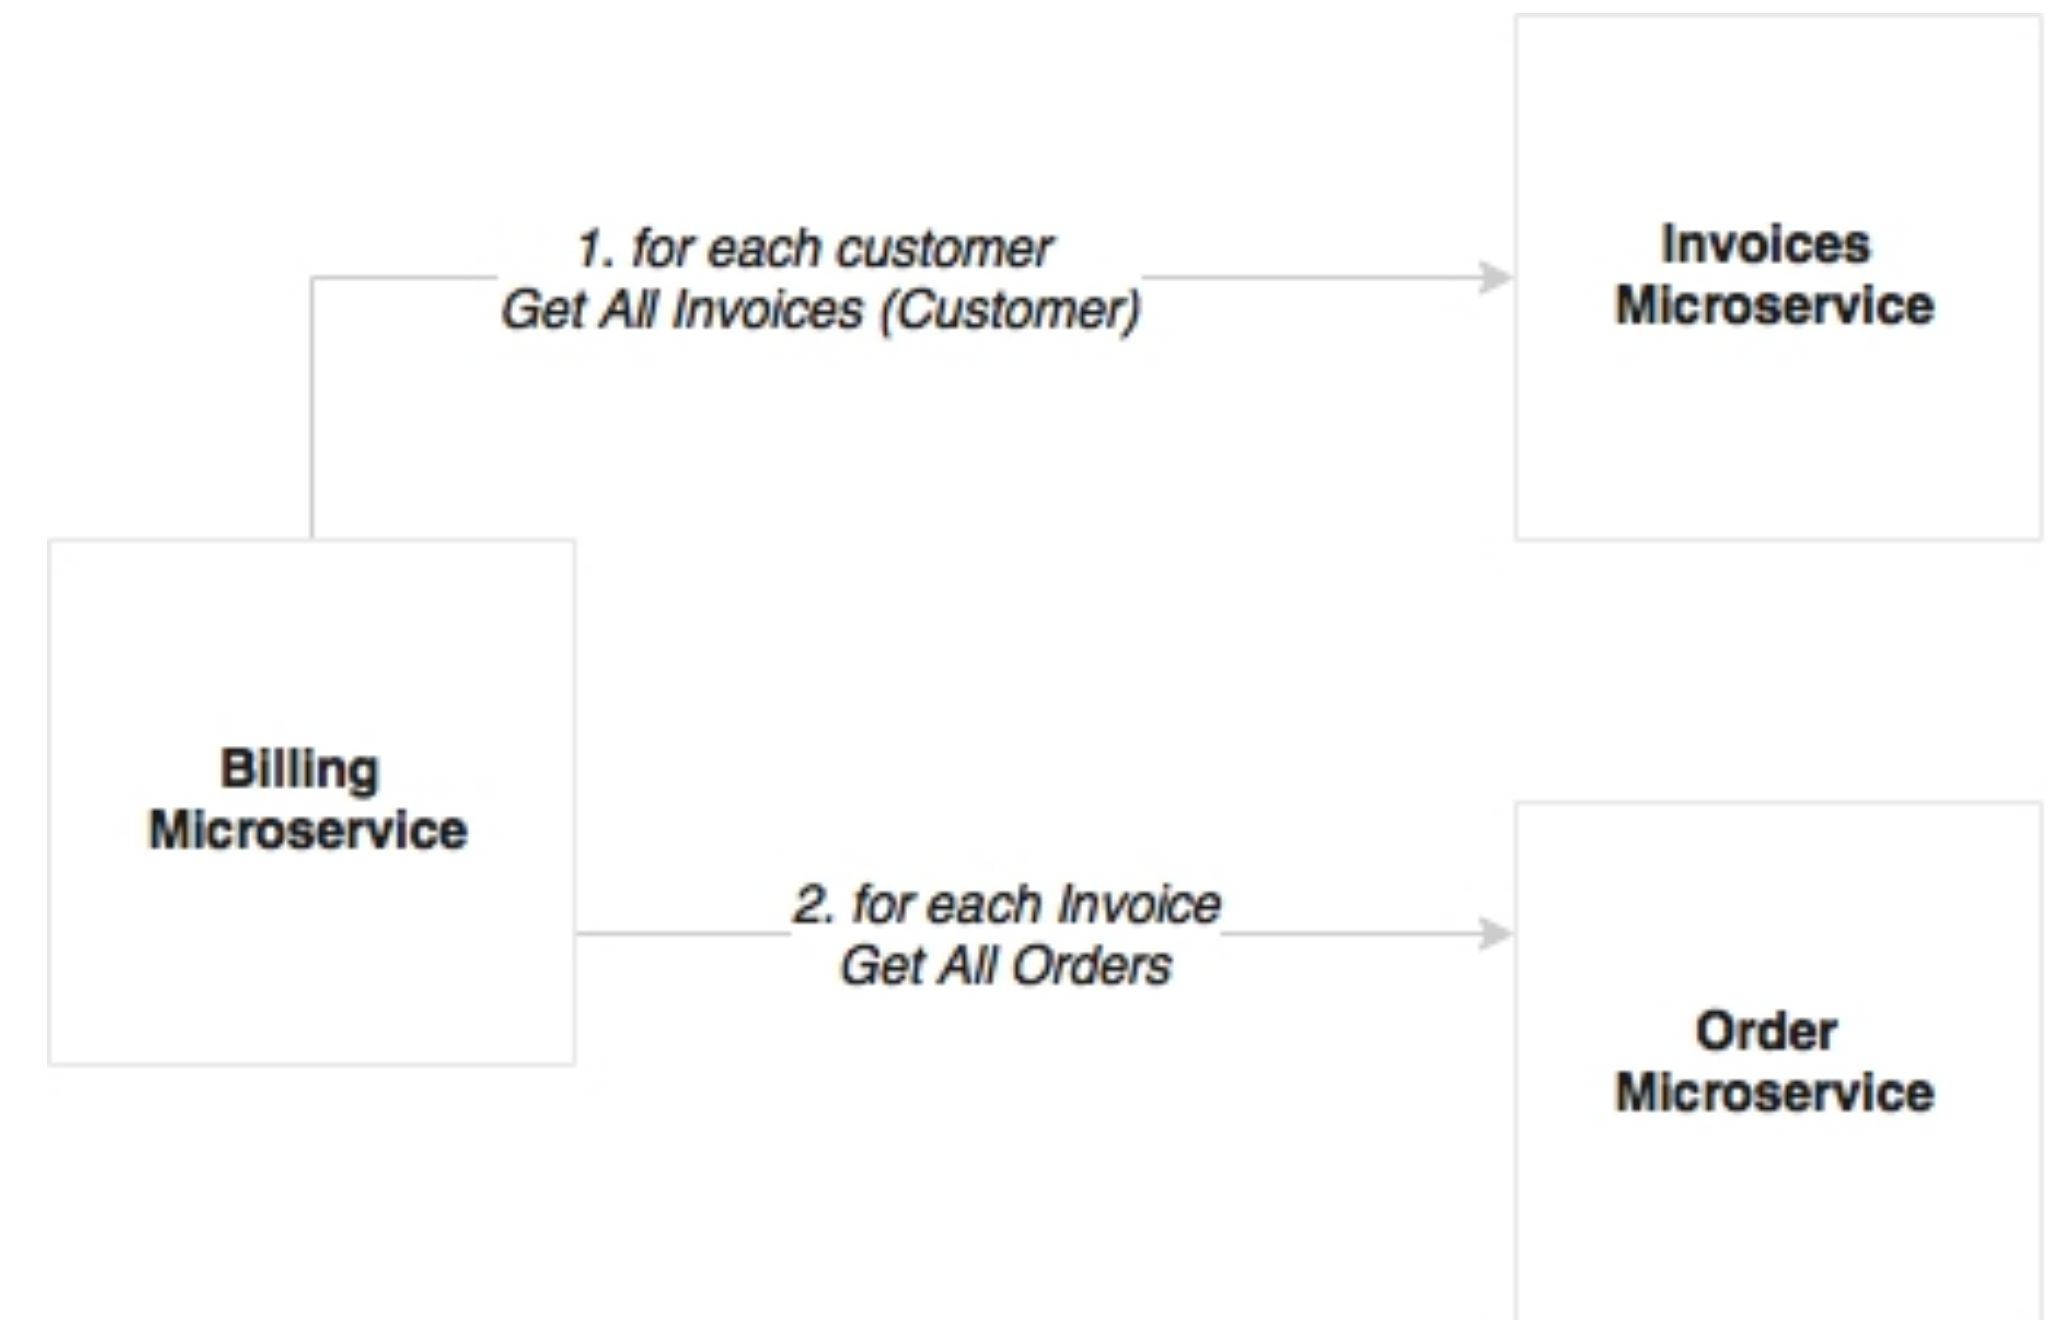
Service versioning considerations

Design for cross origin

- Composite UI web applications may call **multiple microservices** for accomplishing a task, and these could come from different domains and hosts
- **CORS** allows browser clients to send requests to services hosted on different domains
 - One approach is to enable all microservices to allow cross origin requests from other trusted domains
 - The second approach is to use an API gateway as a single trusted domain for the clients

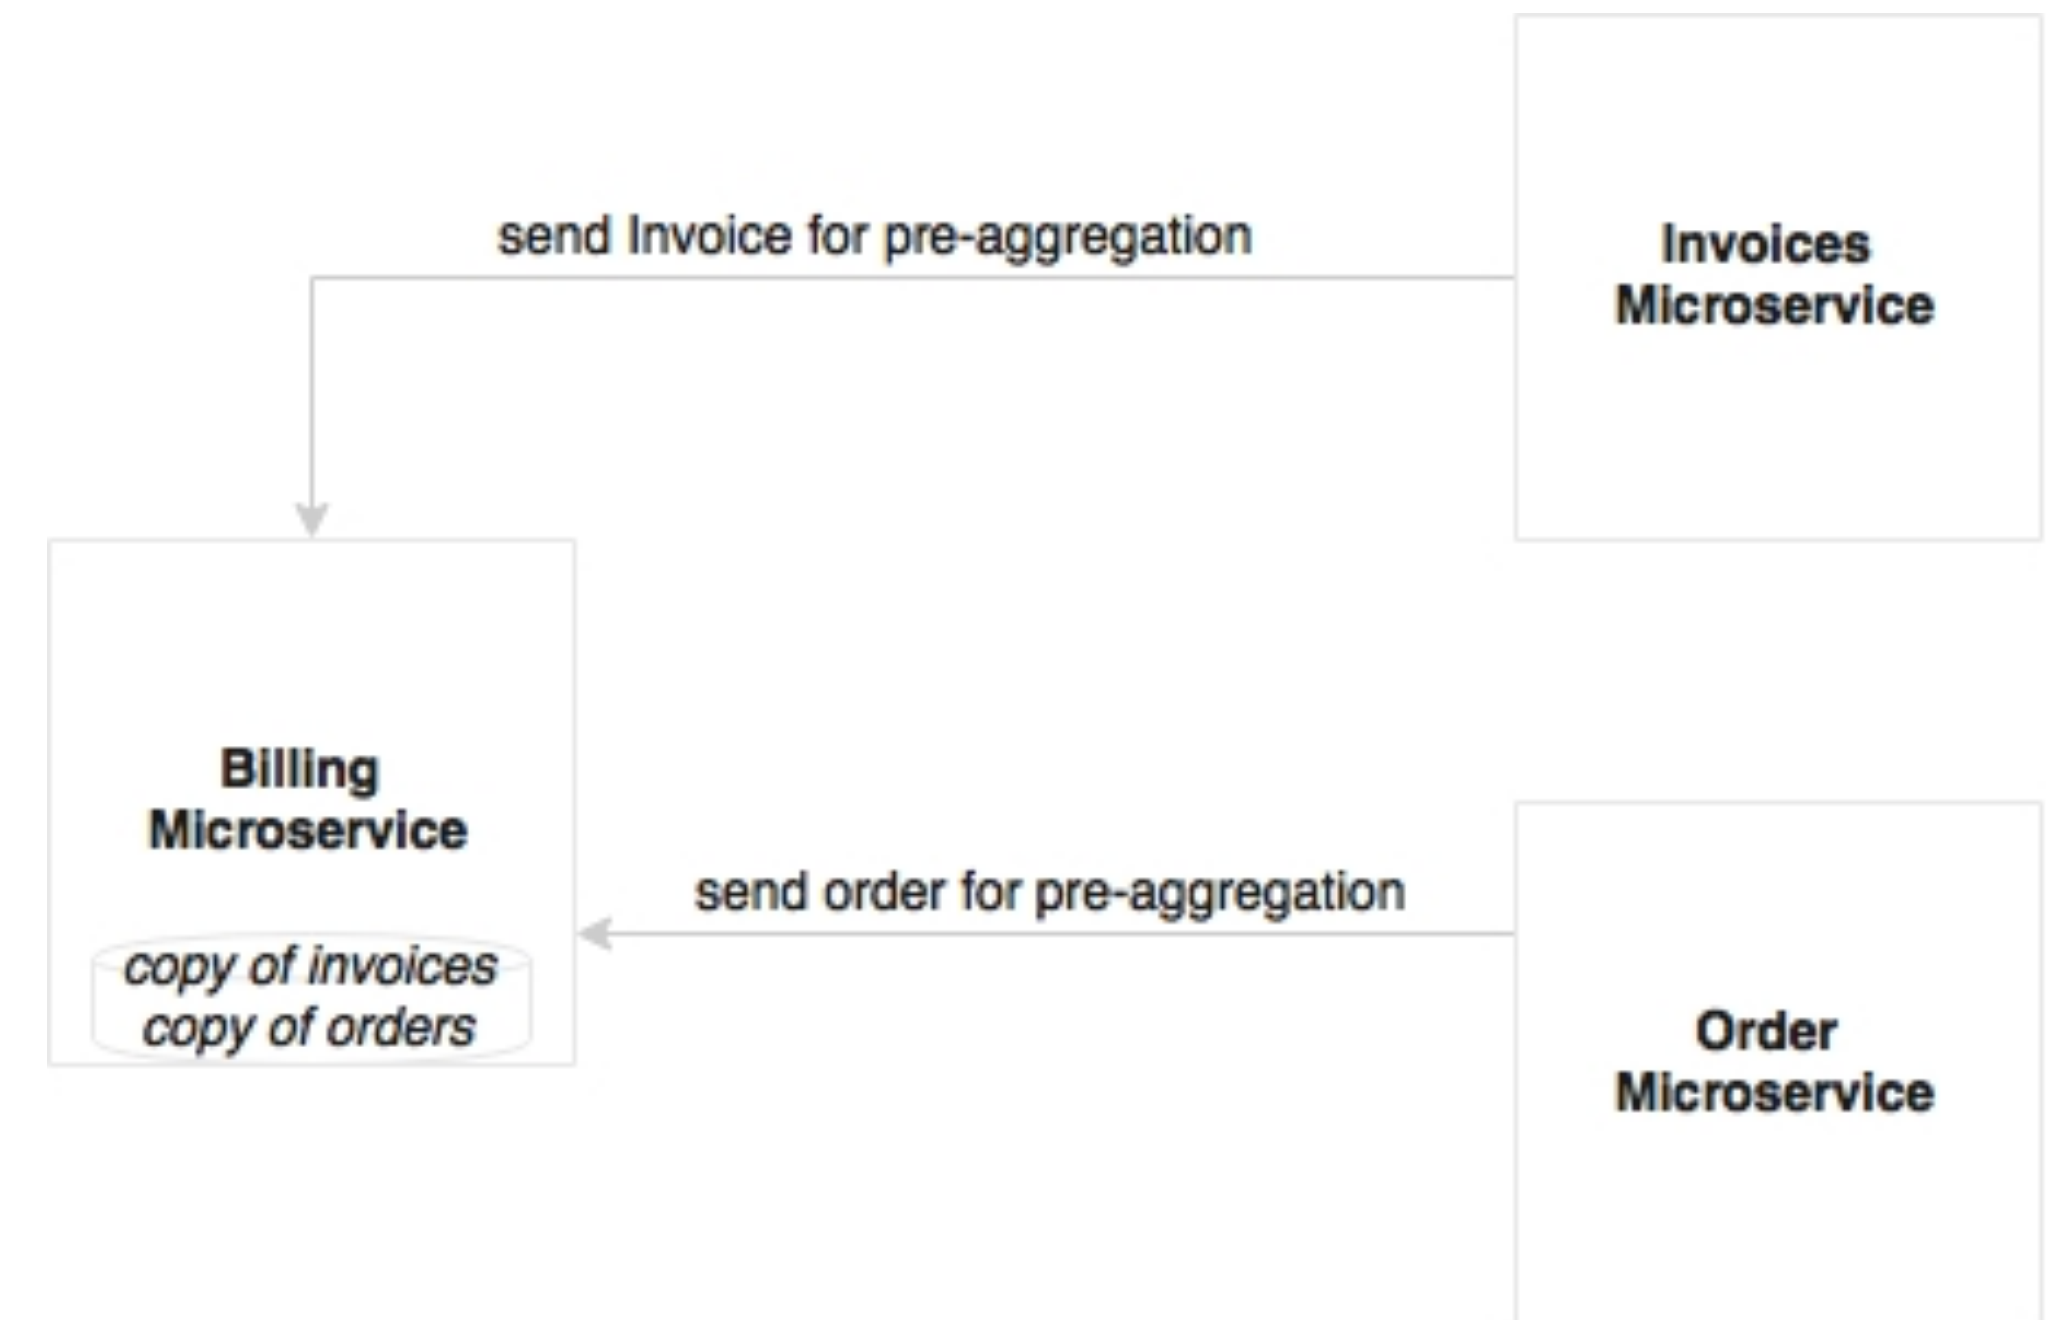
Microservices and bulk operations

- Since we have **broken** monolithic applications into **smaller**, focused services, it is no longer possible to use join queries **across** microservice data stores
- The challenge that arises is that the **Billing** service has to query the **Invoices** service for each **customer** to get all the invoices, and then for each invoice, call the **Order** service for getting the orders



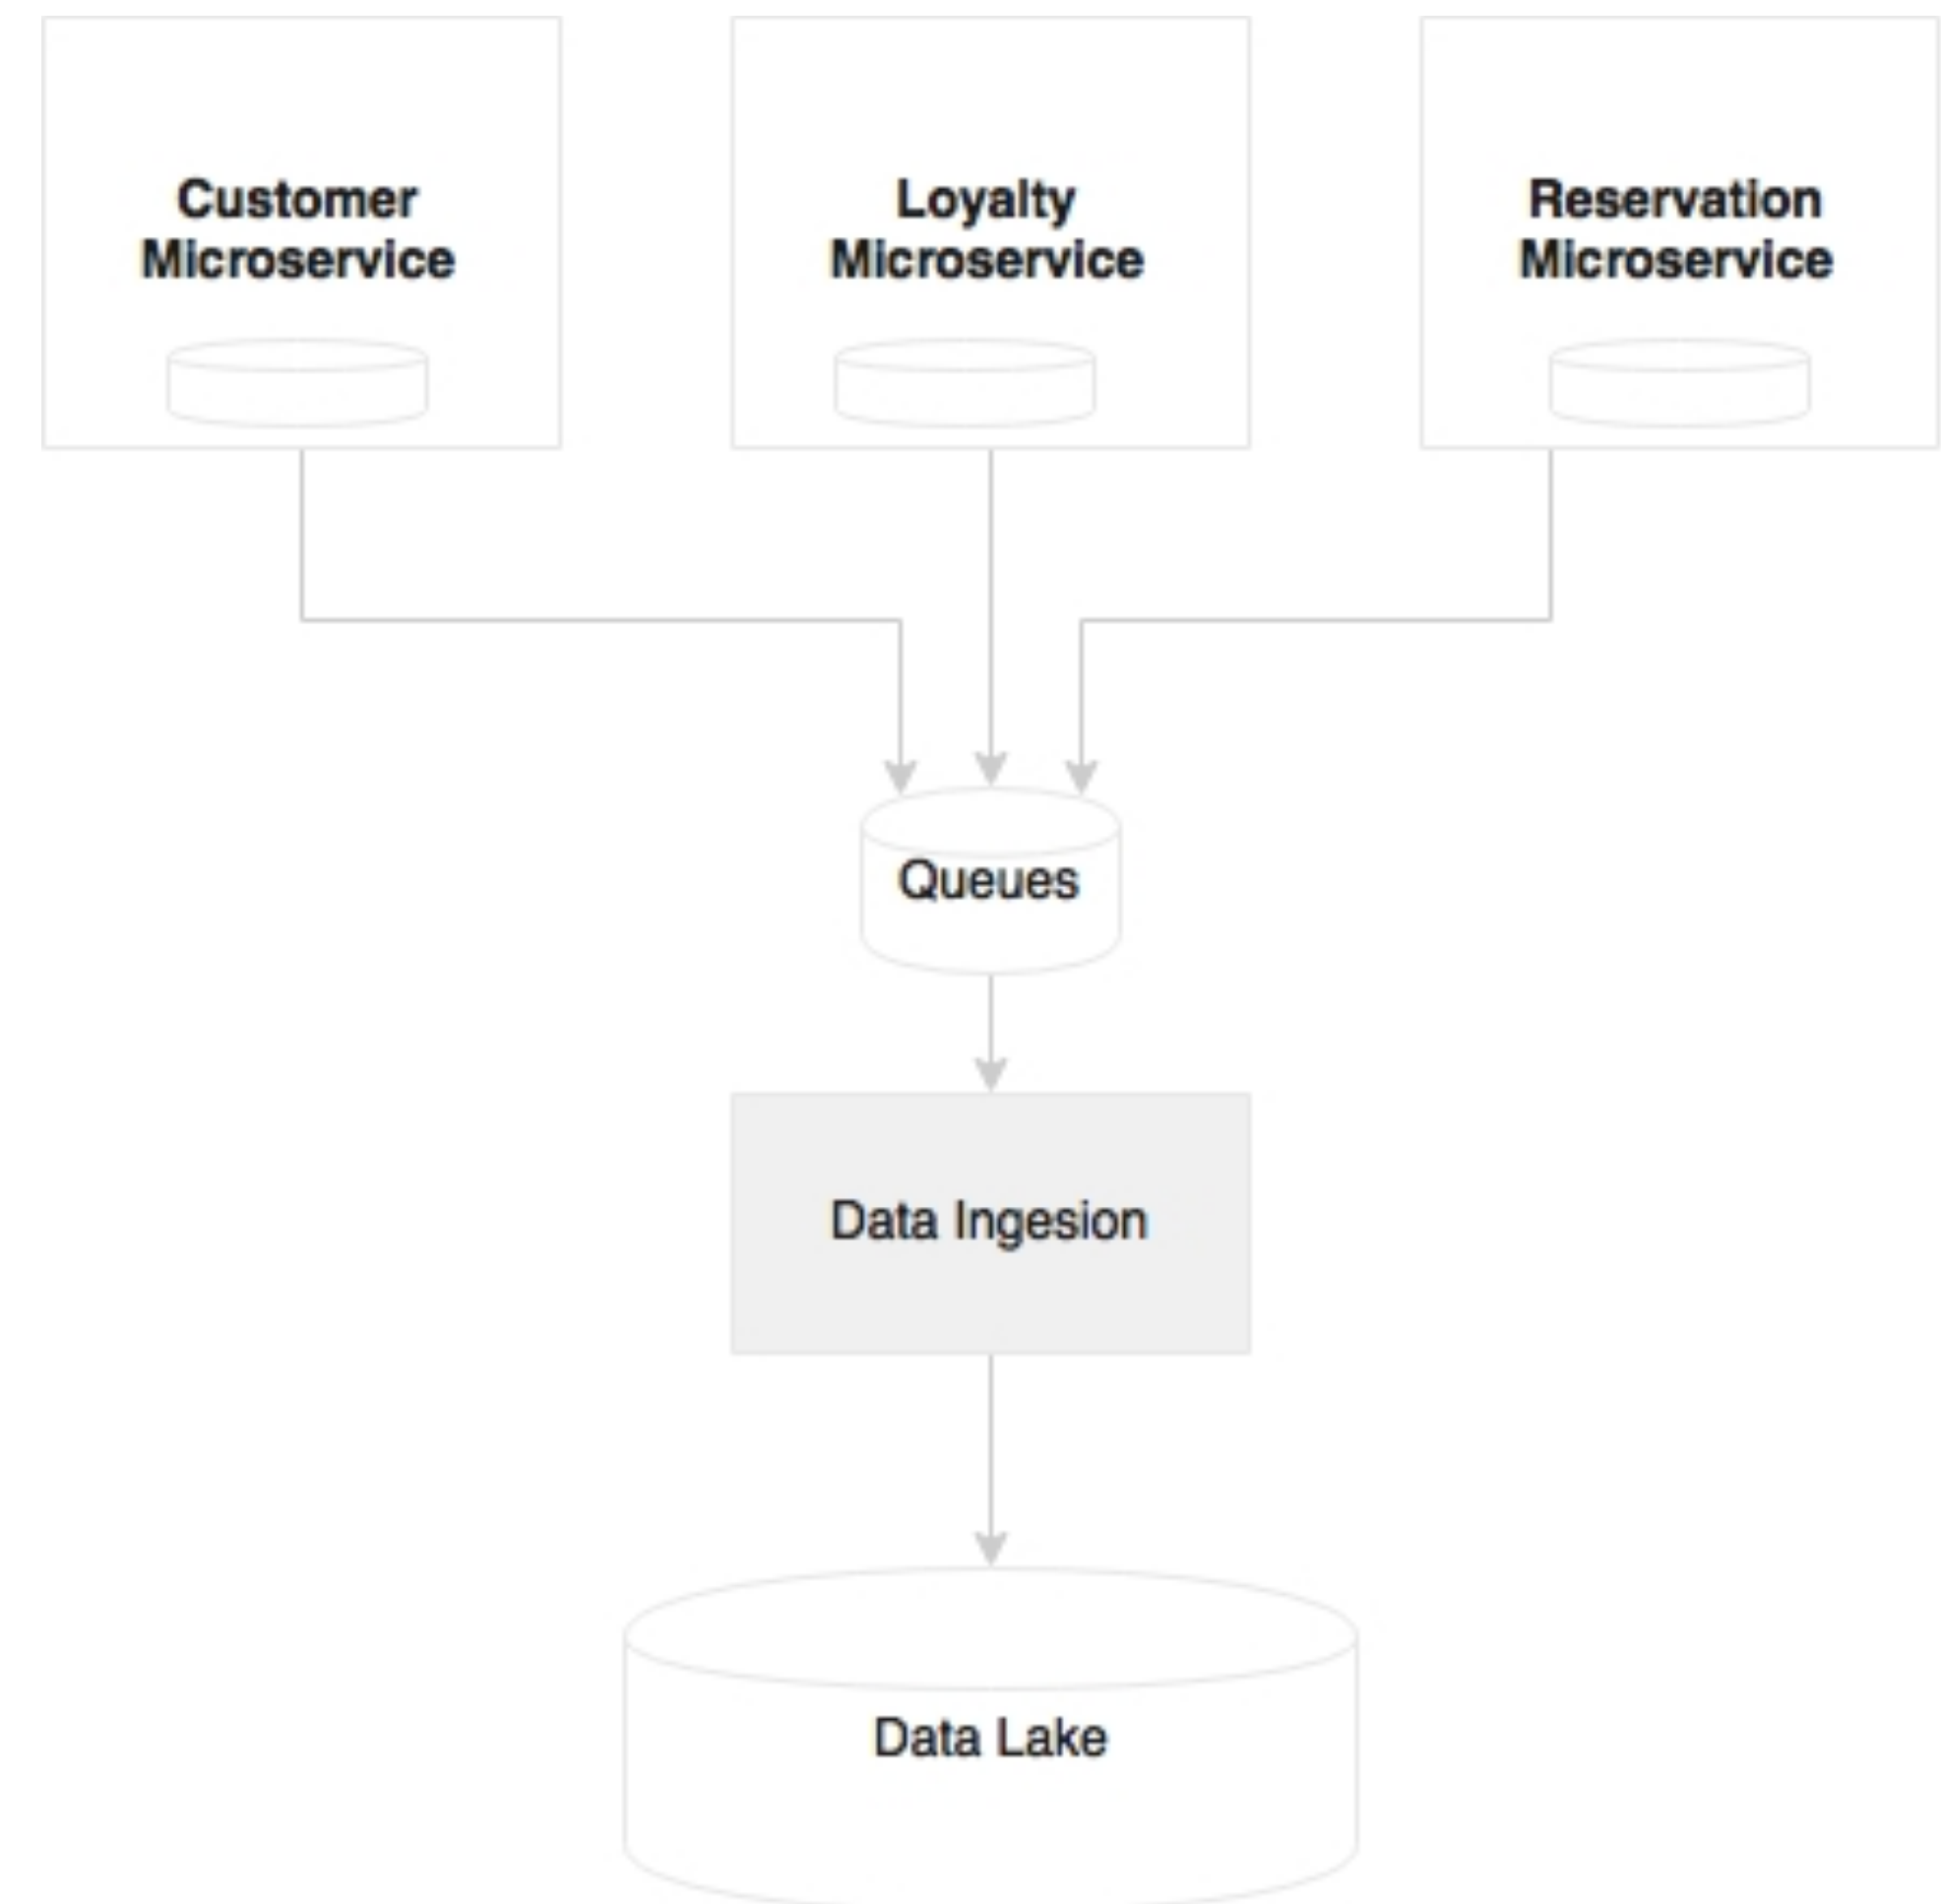
Microservices and bulk operations

- Pre-aggregate data as and when it is created
- Using of batch API
 - each batch further uses parallel threads



Microservices challenges

- Data islands
 - Fragmenting data into heterogeneous data islands
 - Traditional data warehouses like Oracle, Teradata, and others are used primarily for batch reporting
 - But with NoSQL databases (like Hadoop) and microbatching techniques, near real-time analytics is possible with the concept of **data lakes**

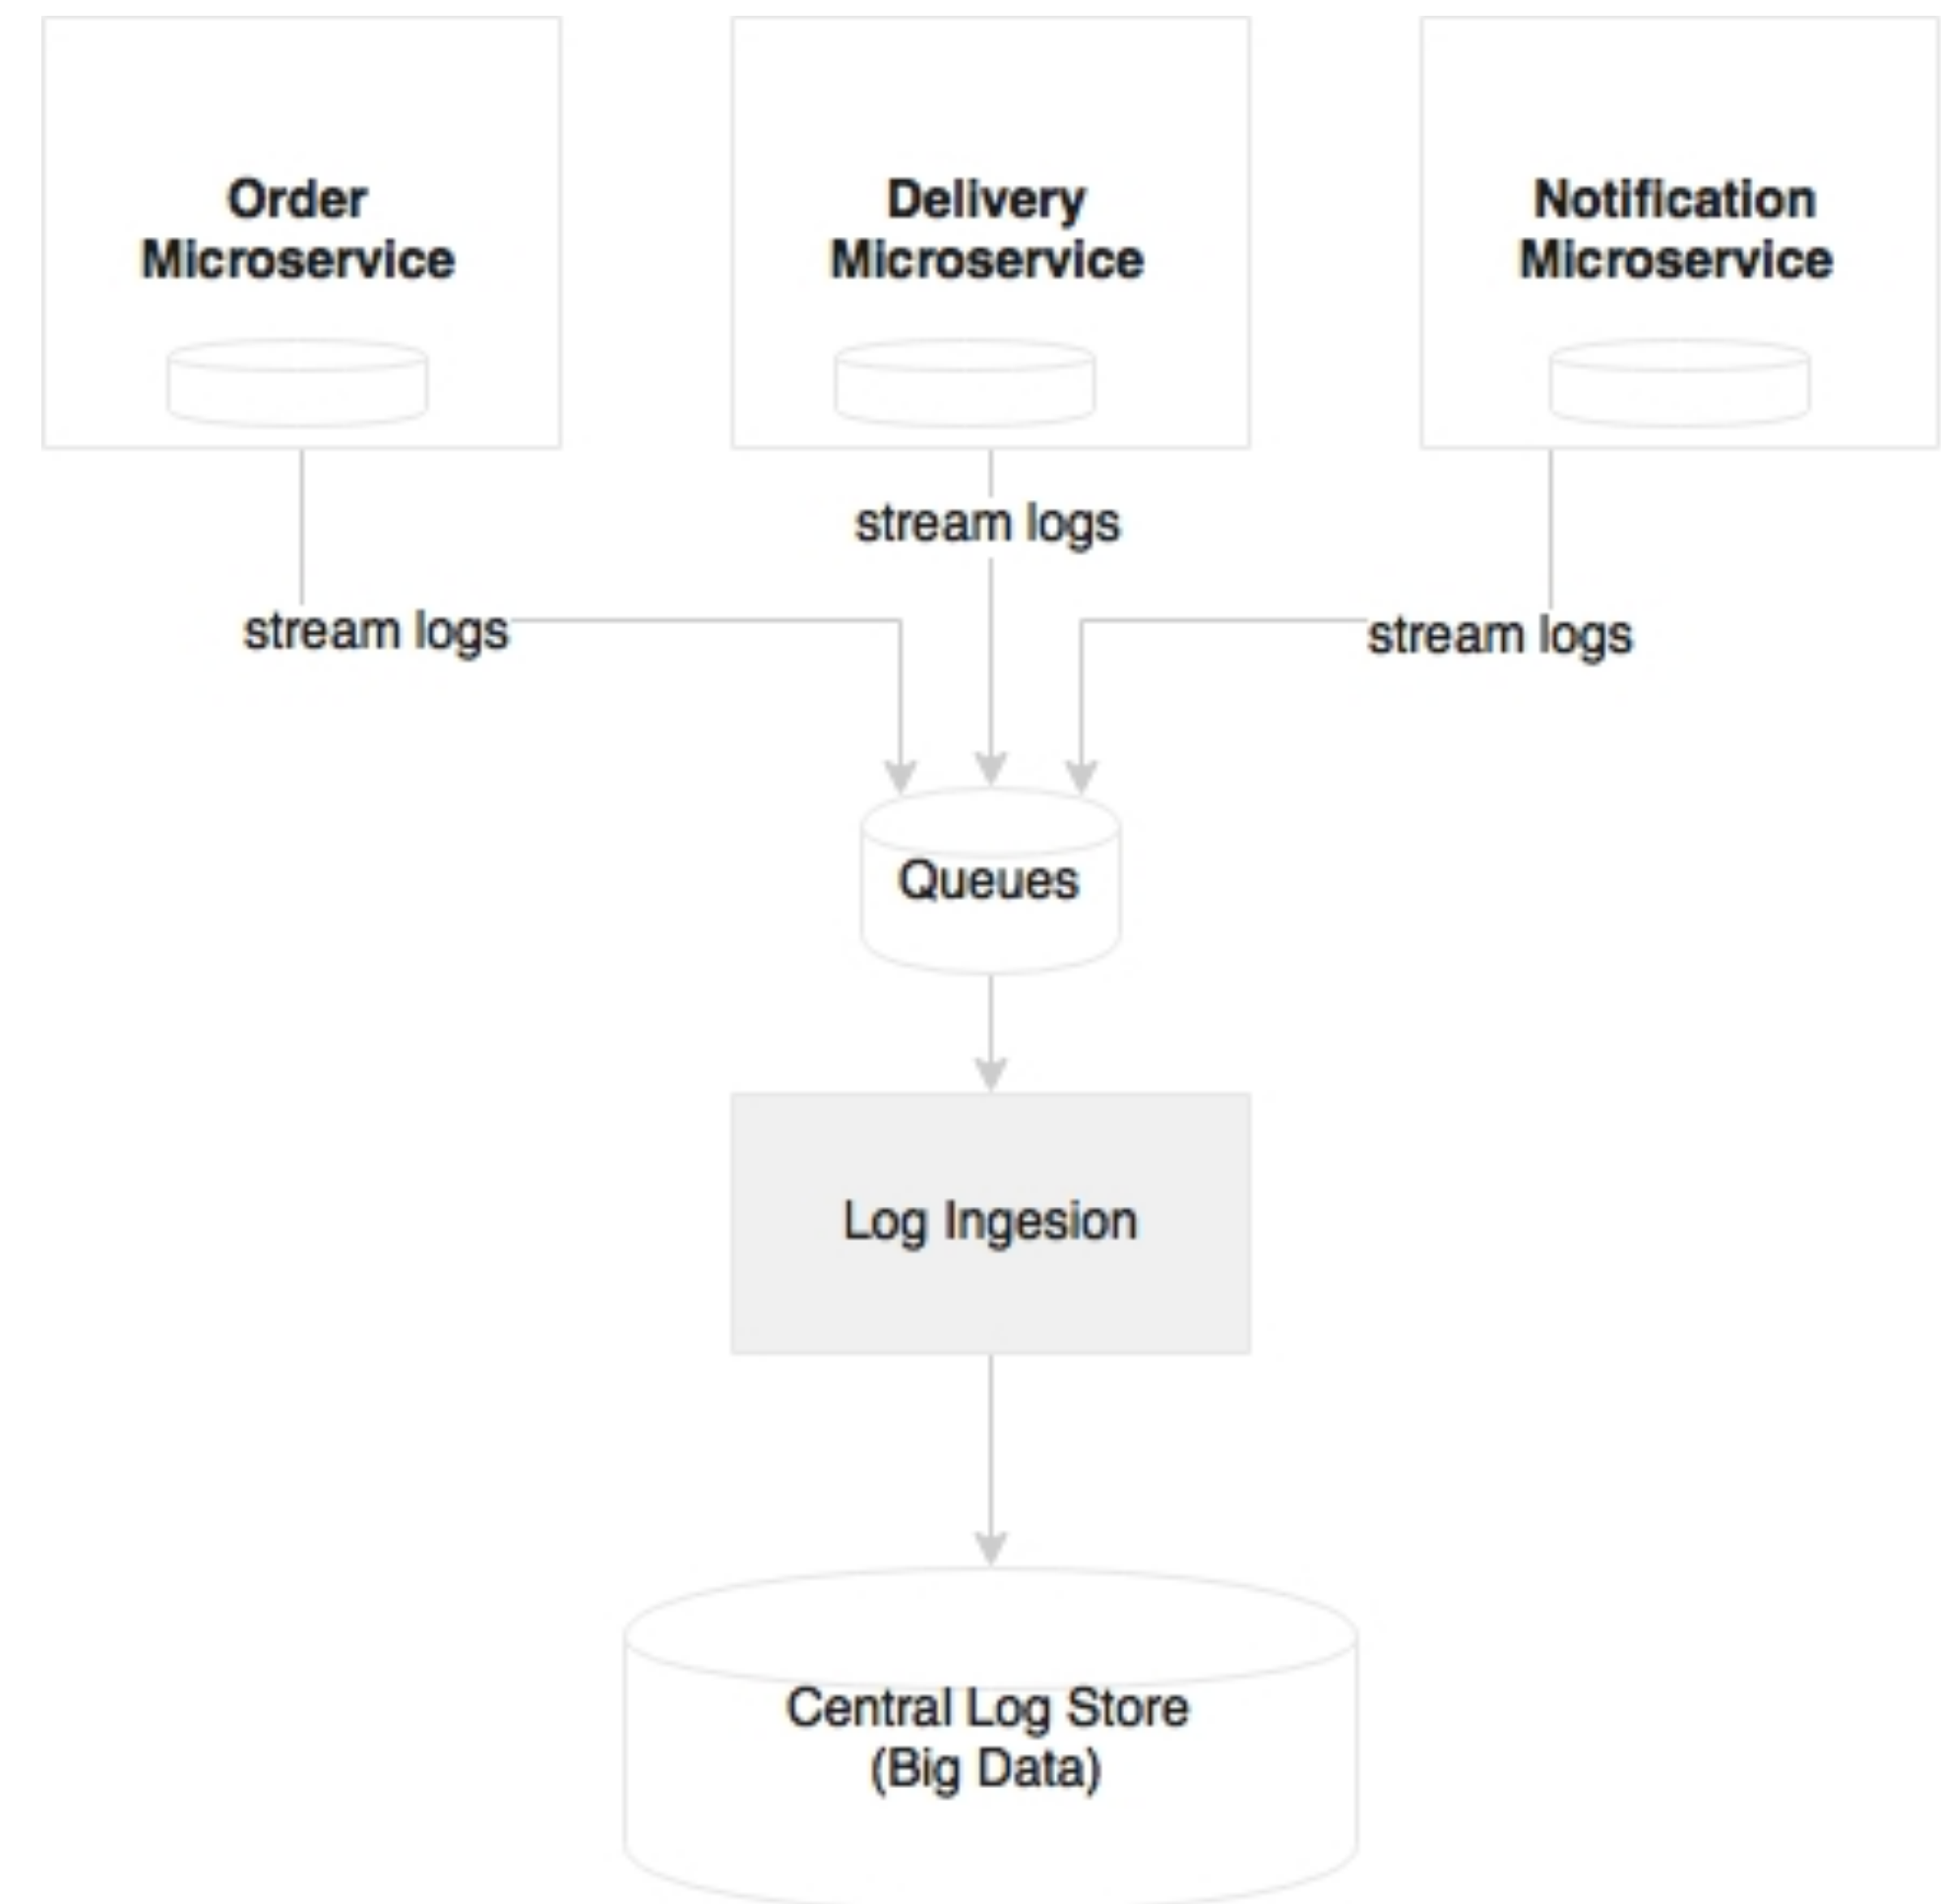


Homework 4.3

- How can be done data porting from microservices to a data lake or a data warehouse? Explain in details your answer.

Logging and monitoring

- Log files are a good piece of information for analysis and debugging
- When we scale services across multiple machines, each service instance could produce separate log files



Dependency management

- Too many dependencies could raise challenges in microservices. Four important design aspects are stated as follows:
 - Reducing dependencies by properly designing service boundaries.
 - Reducing impacts by designing dependencies as loosely coupled as possible. Also, designing service interactions through asynchronous communication styles.
 - Tackling dependency issues using patterns such as circuit breakers.
 - Monitoring dependencies using visual dependency graphs.

Organization culture

- Agile development processes, continuous integration, automated QA checks, automated delivery pipelines, automated deployments, and automatic infrastructure provisioning...

Governance challenges

- Microservices impose decentralized governance, and this is quite in contrast with the traditional SOA governance
- There are number of challenges that comes with a decentralized governance model
 - How do we understand who is consuming a service?
 - How do we ensure service reuse?
 - How do we define which services are available in the organization?
 - How do we ensure enforcement of enterprise policies?

Operation overheads

- Microservices deployment generally increases the number of deployable units and virtual machines (or containers)
- the number of **configurable items** (CIs) becomes **too high**, and the **number of servers** in which these CIs are deployed might also be **unpredictable**

Testing microservices

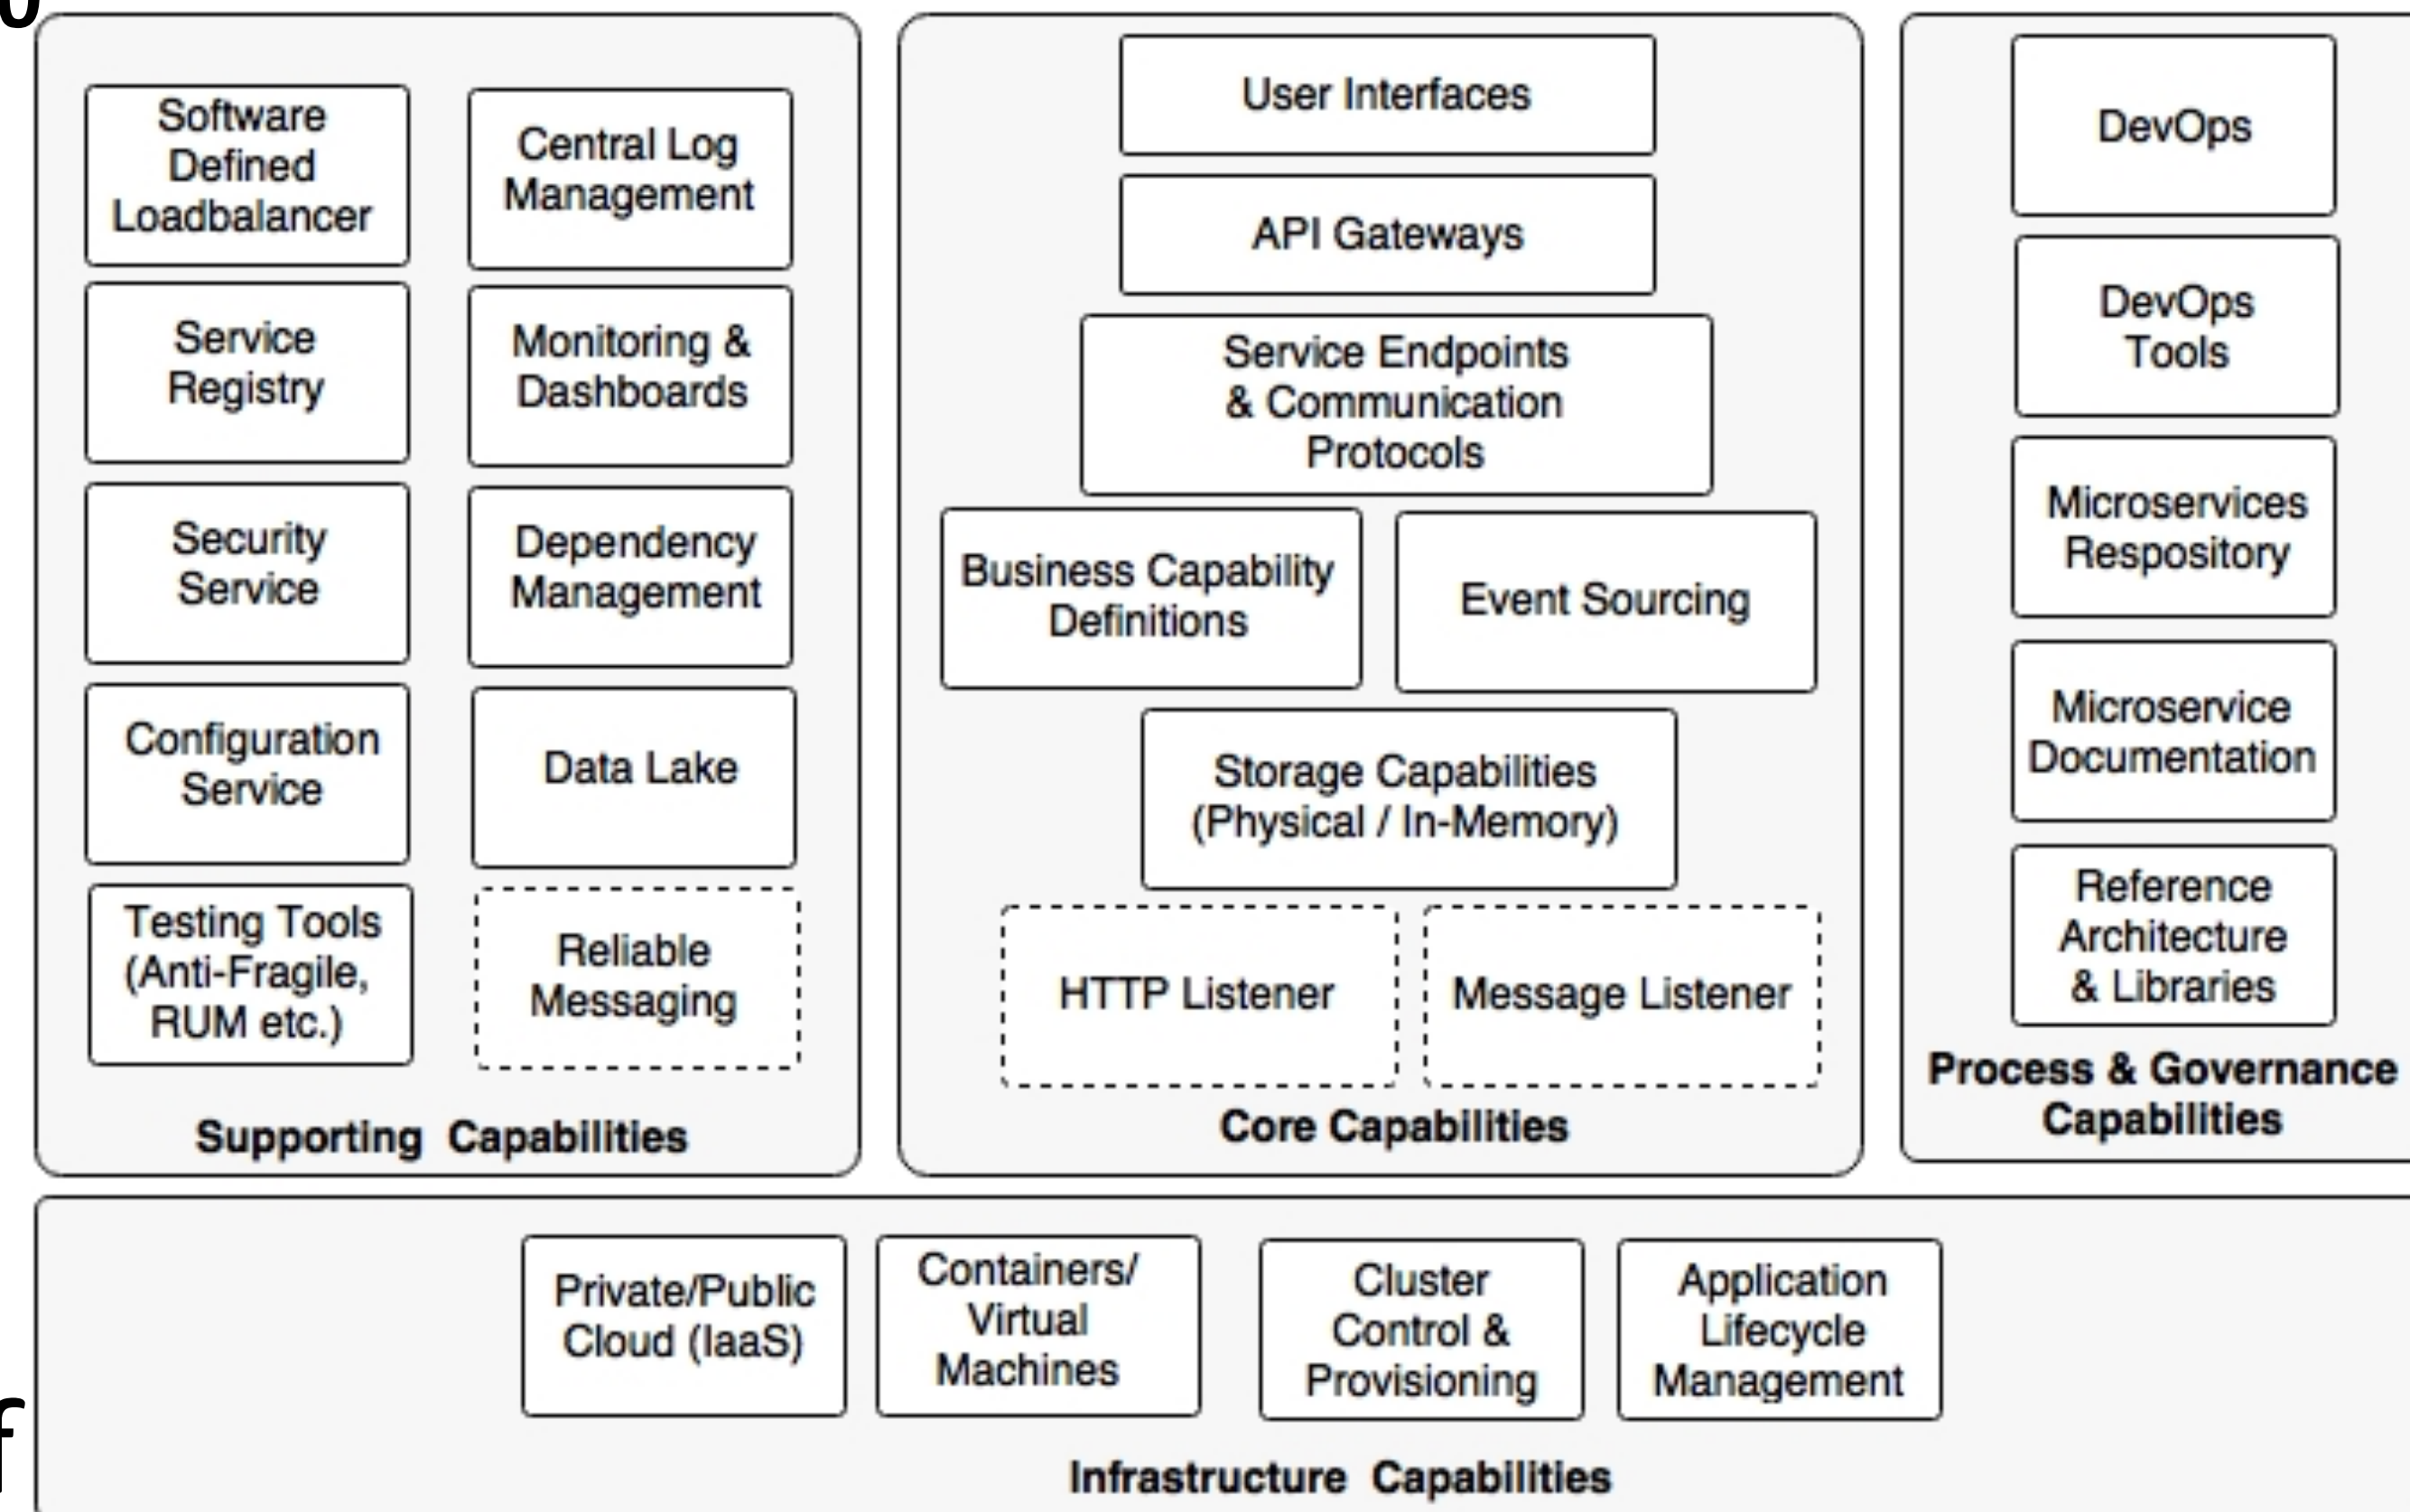
- The issue is how do we test an end-to-end service to evaluate its behavior?
 - Service virtualization or service mocking
 - Consumer driven contract
- Test automation, appropriate performance testing, and continuous delivery approaches such as A/B testing, feature flags, canary testing, blue-green deployments, and red-black deployments, all reduce the risks of production releases

Infrastructure provisioning

- Manual deployment could severely challenge the microservices rollouts
- Microservices require a supporting **elastic cloud-like infrastructure** which can **automatically provision VMs or containers, automatically deploy applications, adjust traffic flows, replicate new version to all instances, and gracefully phase out older versions**

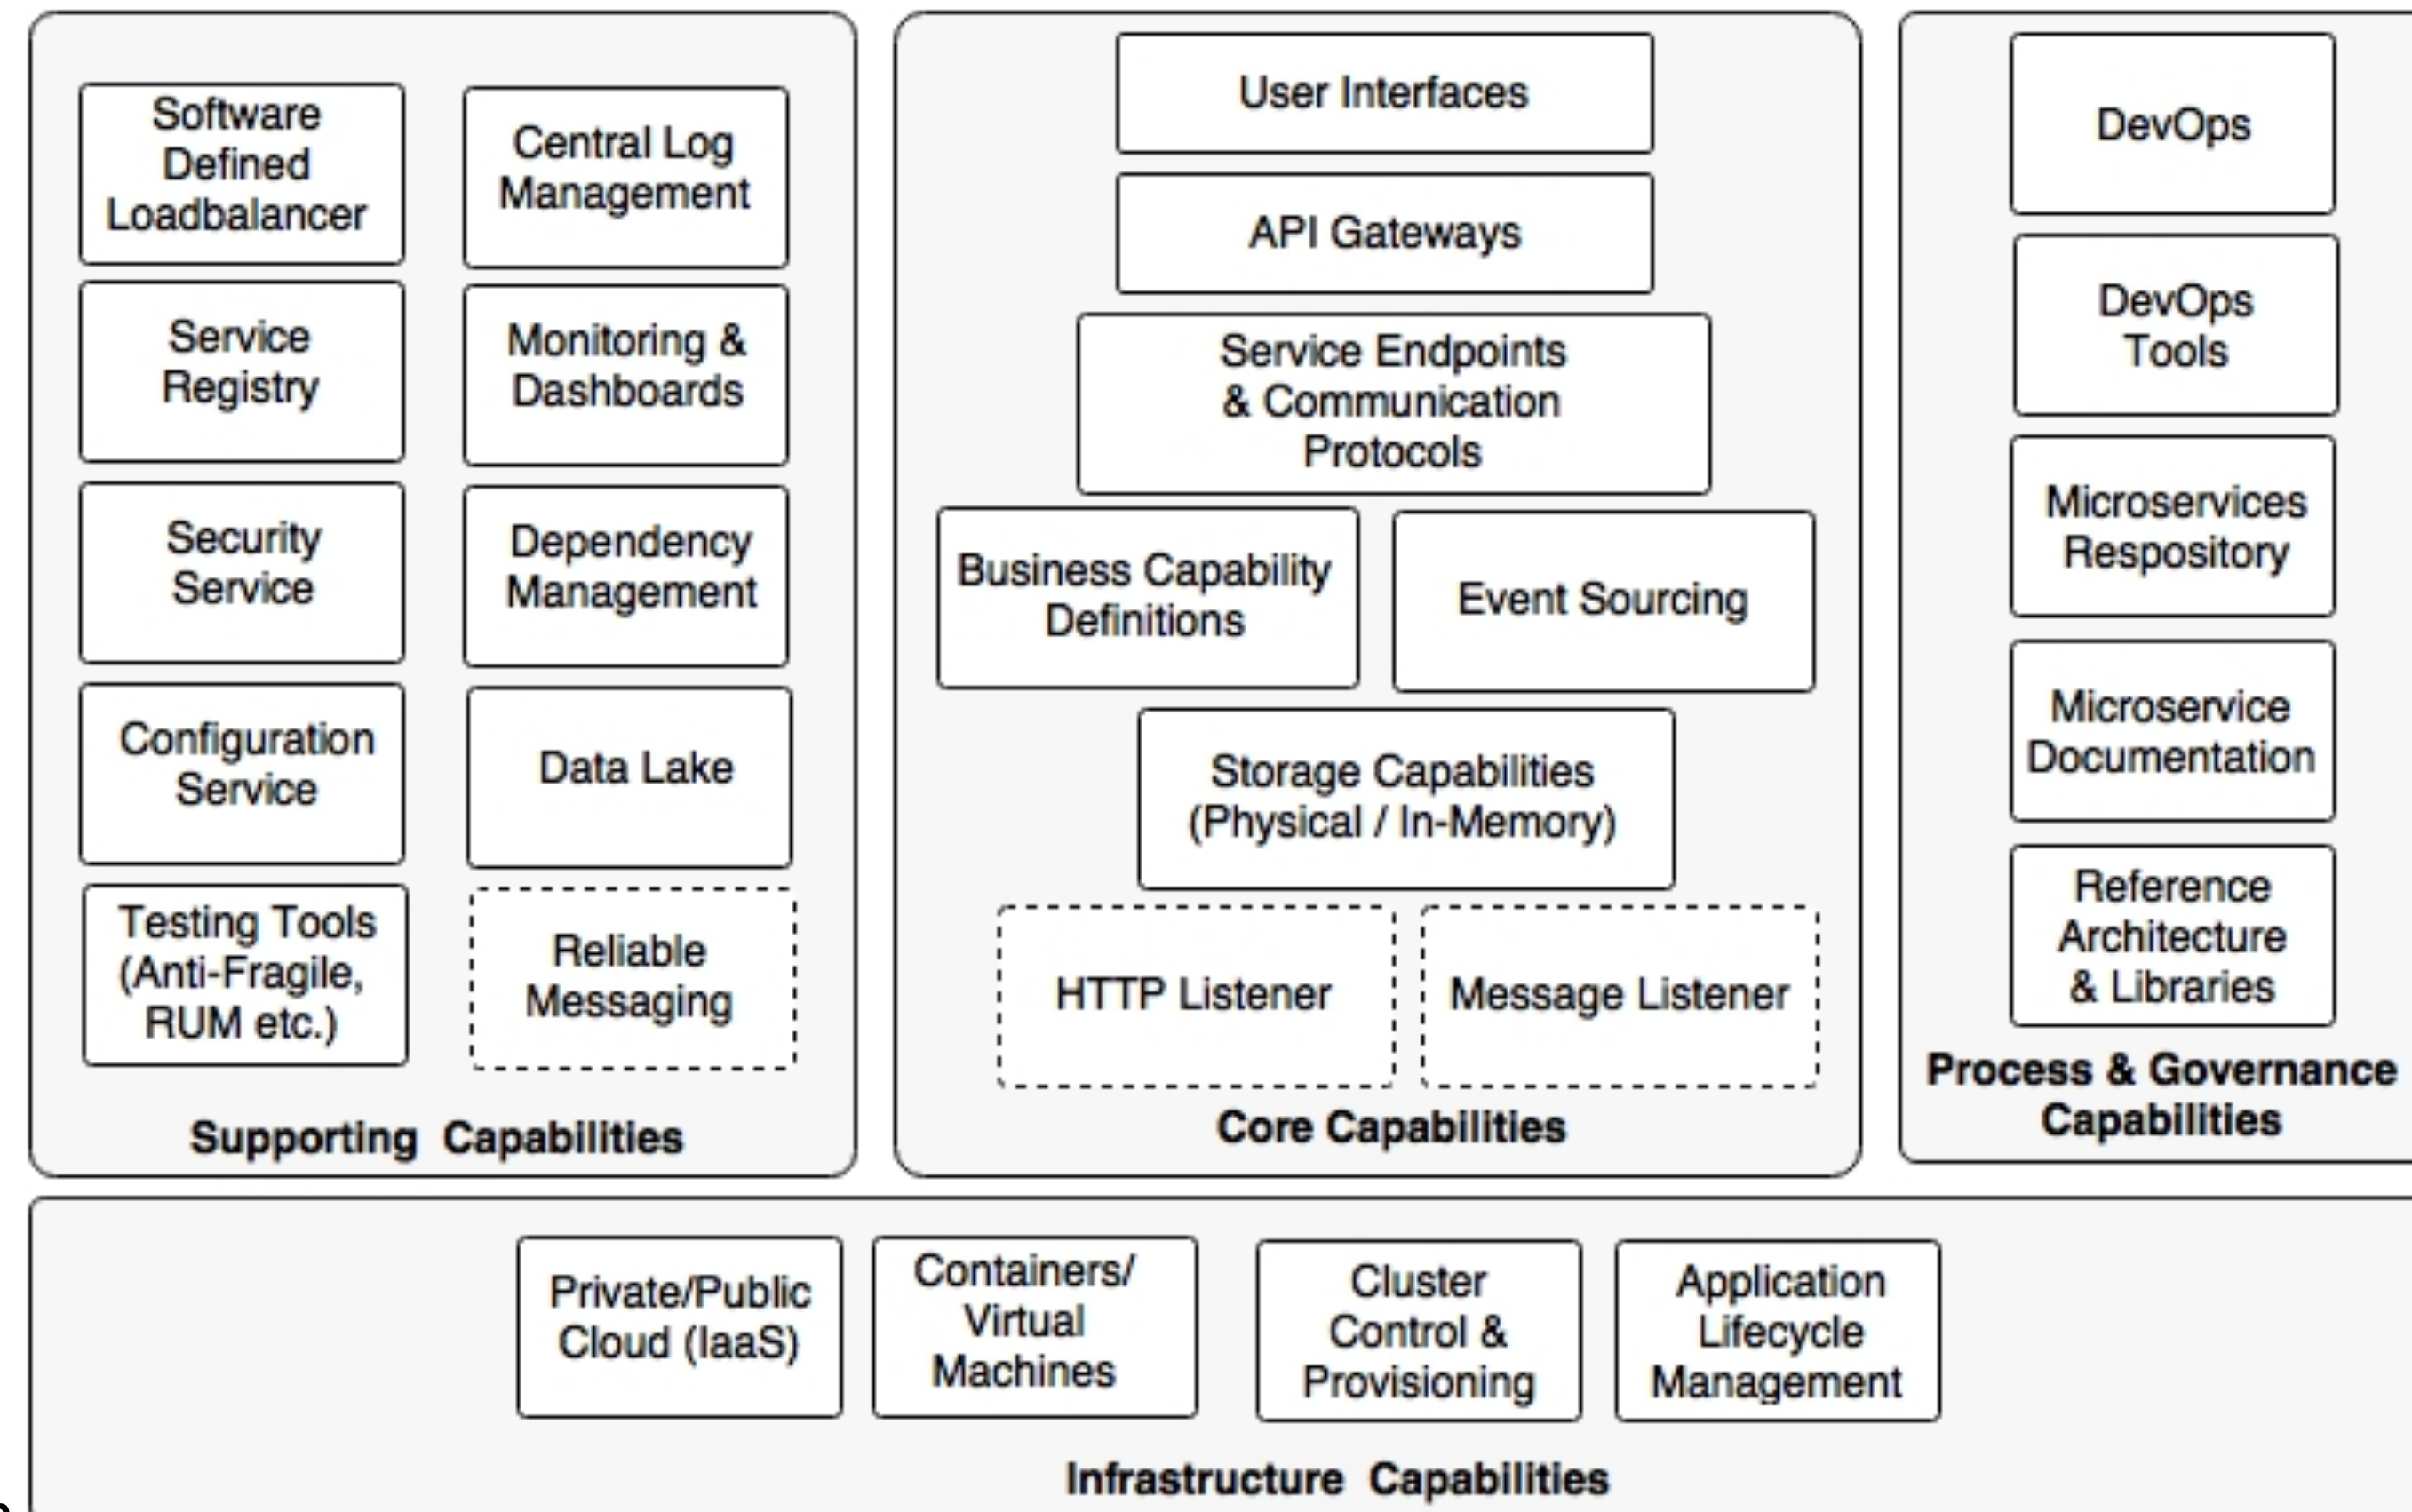
The microservices capability model

- The capability model is broadly classified in to four areas:
 - **Core capabilities:** These are part of the microservices themselves
 - **Supporting capabilities:** These are software solutions supporting core microservice implementations
 - **Infrastructure capabilities:** These are infrastructure level expectations for a successful microservices implementation
 - **Governance capabilities:** These are more of process, people, and reference information




Core capabilities

- Service listeners (HTTP/messaging)
 - HTTP listener is embedded within the microservices, thereby eliminating the need to have any external application server requirement
- Storage capability
 - could be either a physical storage, or it could be an in-memory store
- Business capability definition
 - where the business logic is implemented
- Event sourcing
 - Microservices send out state changes
- Service endpoints and communication protocols
 - Define the APIs for external consumers to consume
- API gateway
 - useful for policy enforcements
- User interfaces
 - could be implemented in any technology, and are channel- and device-agnostic.



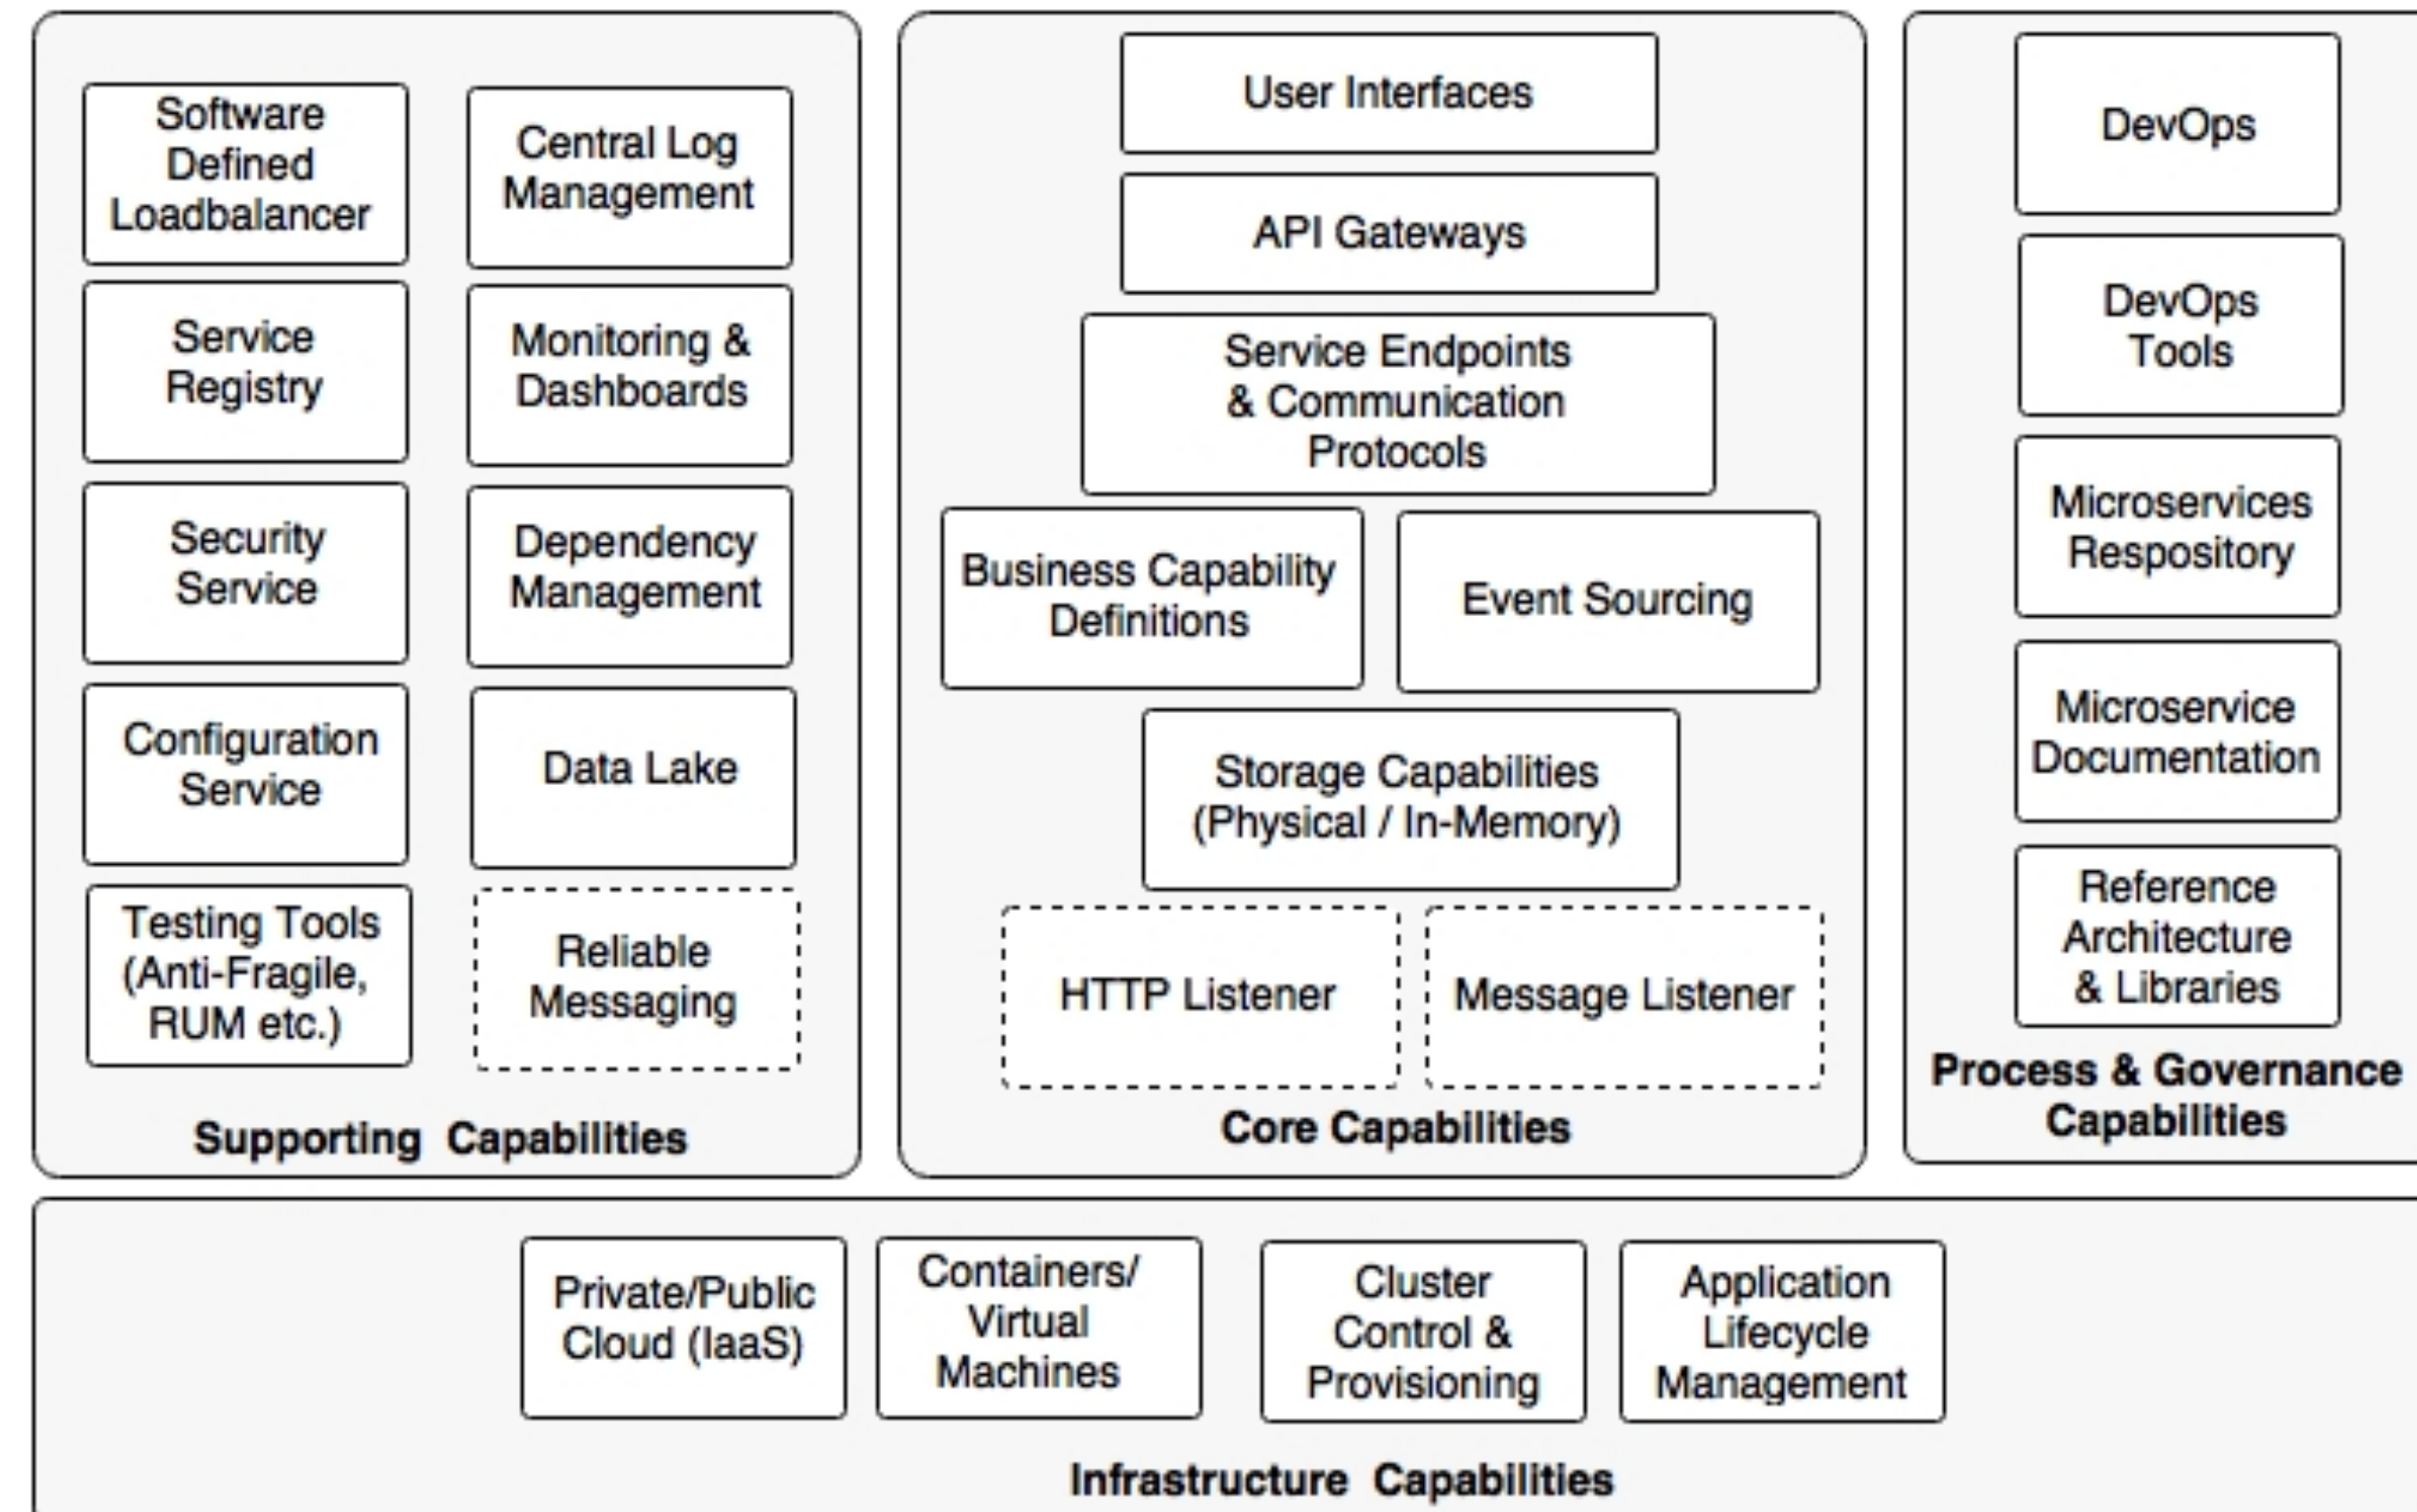
Infrastructure capabilities

- Cloud
 - traditional data center???
- Containers or virtual machines
 - Managing large physical machines is not cost effective
- Cluster control and provisioning
- Application lifecycle management



Supporting capabilities

- Software defined load balancer
- Central log management
- Service registry
- Security service
- Service configuration
- Testing tools (anti-fragile, RUM, and so on)
- Monitoring and dashboards
- Dependency and CI management
- Data lake
- Reliable messaging



Process and governance capabilities

- DevOps
- DevOps tools
- Microservices repository
- Microservices documentation
- Reference Architecture & Libraries

